

原始日志 · 主日志

原始日志 (主日志)

最终验收配置 100 步 完整训练日志

硬件: 1× 昇腾 910C 双 die (Ascend910_9382, 128GB)

框架: MindSpeed-MM v26.1.0 + triton-ascend 3.2.2

日期: 2026-09-10

C4-AI 2026 昇腾赛道复赛 · Qwen3.5-0.8B 迁移

验收主日志（最终配置 100 步）

文件: 20260910_1006_验收主日志_最终配置100步.log | 行数: 630 | 大小: 79.9 KB | 官方几何 world2/mbs4/gas1/dp2 · 全量数据 157,712 样本 · NPU Triton
采集说明: 完整日志（从程序启动开始, 未经裁剪）

```
W0910 10:06:11.862000 11326 torch/distributed/run.py:766]
W0910 10:06:11.862000 11326 torch/distributed/run.py:766] *****
W0910 10:06:11.862000 11326 torch/distributed/run.py:766] Setting OMP_NUM_THREADS environment variable for each process to be 1 in default, to avoid
your system being overloaded, please further tune the variable for optimal performance in your application as needed.
W0910 10:06:11.862000 11326 torch/distributed/run.py:766] *****
[INFO] Starting configuration loading...
[INFO] Loaded configuration file: examples/qwen3_5/e4b_final100.yaml
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdps.data.data_utils.func_utils.mod
el_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.model_name_or_path (type: <class 'str'>, value: /root/Qwen3.5-
0.8B-hf)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdps.data.data_utils.func_utils.mod
el_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.use_fast_tokenizer (type: <class 'bool'>, value: True)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdps.data.data_utils.func_utils.mod
el_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.split_special_tokens (type: <class 'bool'>, value: False)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdps.data.data_utils.func_utils.mod
el_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.image_max_pixels (type: <class 'int'>, value: 262144)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdps.data.data_utils.func_utils.mod
el_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.image_min_pixels (type: <class 'int'>, value: 1024)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdps.data.data_utils.func_utils.mod
el_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.video_max_pixels (type: <class 'int'>, value: 16384)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdps.data.data_utils.func_utils.mod
el_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.video_min_pixels (type: <class 'int'>, value: 0)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdps.data.data_utils.func_utils.mod
el_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.video_fps (type: <class 'float'>, value: 2.0)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdps.data.data_utils.func_utils.mod
el_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.video_maxlen (type: <class 'int'>, value: 64)
[INFO] Automatically registered YAML field: model.gdn_implementation (type: <class 'str'>, value: triton)
[INFO] Automatically registered YAML field: model.causal_conv1d_implementation (type: <class 'str'>, value: triton)
[INFO] Automatically registered YAML field: model.skip_gdn_recompute (type: <class 'bool'>, value: False)
[INFO] Automatically registered YAML field: model.skip_flash_attn_recompute (type: <class 'bool'>, value: False)
[INFO] Automatically registered YAML field: features.recompute (type: <class 'bool'>, value: False)
=====
BaseArguments Configuration Manager Summary
=====

Configuration class: Arguments
Configuration file: examples/qwen3_5/e4b_final100.yaml
Allow YAML extensions: True
Allow CLI override: True
Strict CLI validation: True
Allow registering YAML fields: True

Defined field count: 208
Dynamic field count: 14
Total field count: 222

===== Configuration Details =====
parallel:
  data_parallel_size: 2
  fully_shard_parallel_size: 2
  fsdp_plan:
    ignored_modules: []
    apply_modules: ['model.visual', 'model.visual.blocks.*', 'model.language_model.embed_tokens', 'model.language_model', 'model.language_model.la
yers.*', 'lm_head', 'mtp']
    param_dtype: bf16
    reduce_dtype: fp32
    output_dtype: None
    cast_forward_inputs: True
    reshard_after_forward: False
    num_to_forward_prefetch: 1
    num_to_backward_prefetch: 1
    pregather: True
    hook_modules: None
    cpu_offload: False
  tensor_parallel_size: 1
  tp_plan:
    colwise_parallel: []
    rowwise_parallel: []
    sequence_parallel: []
  ring_attention_size: 1
  ulysses_parallel_size: 1
  expert_parallel_size: 1
  expert_fully_shard_parallel_size: 2
  ep_plan:
```

```

    apply_modules: []
    use_npu_fused_ops: True
    dispatcher: alltoall
    apply_efsdp_modules: []
    recompute: False
    recompute_plan:
        apply_modules: []
        use_reentrant: False
    local_rank: 0
    global_rank: 0
    world_size: 2
model:
    model_id: qwen3_5
    model_name_or_path: /root/Qwen3.5-0.8B-hf
    trust_remote_code: True
    attn_implementation: flash_attention_2
    freeze: ['model.visual']
    mtp_num_layers: 0
    mtp_loss_scaling_factor: 0.1
    gdn_implementation: triton
    causal_conv1d_implementation: triton
    skip_gdn_recompute: False
    skip_flash_attn_recompute: False
data:
    dataset_param:
        dataset_type: huggingface
        basic_parameters:
            cache_dir: ./cache_dir/
            template: qwen3_v1_nothink
            chat_template: None
            enable_thinking: False
            dataset_dir: /root/data/coco
            dataset: [' /root/data/output_llava_coco_data.json']
            cutoff_len: 1024
            train_on_prompt: False
            mask_history: False
            streaming: False
            overwrite_cache: False
            preprocessing_batch_size: 1000
            preprocessing_num_workers: 16
            max_samples: None
            tool_format: None
            val_dataset: None
            val_max_samples: None
            val_rate: None
            packing: None
            neat_packing: False
            preprocess_on_fly: True
            async_preprocess: False
            async_preprocess_buffer_size: None
            stage: sft
        preprocess_parameters:
            model_name_or_path: /root/Qwen3.5-0.8B-hf
            use_fast_tokenizer: True
            split_special_tokens: False
            image_resolution: 512
            video_resolution: 128
            image_max_pixels: 262144
            image_min_pixels: 1024
            video_max_pixels: 16384
            video_min_pixels: 0
            video_fps: 2.0
            video_maxlen: 64
            image_do_pan_and_scan: False
            crop_to_patches: False
            use_audio_in_video: False
            audio_sampling_rate: 16000
            trust_remote_code: False
    attr:
        ranking: False
        system: None
        images: images
        videos: None
        audios: None
        prompt: None
        query: None
        response: None
        history: None
        messages: messages
        tools: None
        role_tag: role
        content_tag: content
        user_tag: user
        assistant_tag: assistant
        observation_tag: observation
        function_tag: function_call
        system_tag: system
        chosen: None
        rejected: None
        formatting: sharegpt
    val_dataset_param: None
    dataloader_param:
        dataloader_mode: sampler

```

```

sampler_type: BaseRandomBatchSampler
shuffle: False
drop_last: True
pin_memory: True
collate_param:
  collator_id: None
  model_name: qwen3vl
  ignore_pad_token_for_loss: True
  pad_to_multiple_of: 8
num_workers: 8
enable_preload: False
training:
  profile:
    enable: False
    start_step: 1
    end_step: 2
    save_path: ./profiling
    record_shapes: True
    with_memory: True
    with_stack: True
    ranks: [0]
    _train_steps: -1
    local_rank: 0
    global_rank: 0
    world_size: 2
    profile_this_rank: False
  lora:
    enable: False
    rank: 8
    alpha: 16
    target_modules: ['q_proj', 'k_proj', 'v_proj']
    dropout: 0.0
    init_lora_weights: True
    pretrained_lora_path: None
    lora_target_modules_support: None
    disable_peft_moe_conversion: True
  lr: 1e-05
  lr_min: 0.0
  lr_start: 0.0
  weight_decay: 0.0
  no_decay_modules: []
  no_decay_params: []
  optimizer: adamw
  mem_fraction_static: 0.8
  matched_adamw_rms: 0.2
  muon_momentum: 0.95
  ns_steps: 5
  adam_fused: True
  adam_beta1: 0.9
  adam_beta2: 0.999
  adam_eps: 1e-08
  clip_grad: 0.0
  clip_grad_foreach: True
  micro_batch_size: 4
  gradient_accumulation_steps: 1
  lr_warmup_ratio: 0.1
  lr_decay_style: cosine
  lr_decay_ratio: 1.0
  enable_full_determinism: False
  ckpt_manager: dcp
  save_async: False
  save_ckpt_dtype: None
  load_checkpoint_path: None
  save_steps: 0
  seed: 42
  max_steps: None
  init_model_with_meta_device: True
  train_iters: 100
  train_epochs: None
  load: /root/Qwen3.5-0.8B-dcp
  load_format: auto
  save_format: auto
  load_strict: False
  load_rank0_and_broadcast: False
  no_load_optim: True
  no_load_rng: True
  save: ./save_path
  no_save_optim: True
  no_save_rng: True
  log_interval: 1
  save_interval: 10000
  val_interval: 0
  val_micro_batch_size: None
  use_deter_comp: False
  allow_hf32: None
  plugin: ['mindspeed_mm/fsdp/models/qwen3_5', 'mindspeed_mm/fsdp/data/datasets/huggingface']
  lr_scaling_plan: None
  quantization_plan: QuantizeConfig(recipe_name=None, apply_modules=[], ignored_modules=[], quant_converters=[], extra_args={}, enable_fsdp_low_precision_all_gather=True, fsdp_low_precision_all_gather_mode='on-demand')
  _train_steps: -1
  local_rank: 0
  global_rank: 0
  world_size: 2

```

```

global_batch_size: 8
tools:
  profile:
    enable: False
    profile_type: static
    ranks: [0]
    static_param:
      level: level1
      with_stack: False
      with_memory: False
      record_shapes: False
      with_cpu: True
      save_path: ./profiling
      start_step: 10
      end_step: 11
      data_simplification: False
      aic_metrics_type: PipeUtilization
      analyse_flag: True
  memory_profile:
    enable: False
    start_step: 1
    end_step: 2
    save_path: ./memory_snapshot
    dump_ranks: [0]
    stacks: all
    max_entries: None
    mem_info: False
features:
  recompute_plan:
    apply_modules: ['model.visual.blocks.*', 'model.language_model.layers.*']
    use_reentrant: False
  loss_cfg:
    loss_type: default
    router_aux_loss_coef: 0.0
    router_aux_loss_offload: False
  enable_chunk_loss: False
  enable_dynamic_chunk_loss: False
  chunkloss_plan:
    apply_module: lm_head
    chunk_size: 1024
    total_chunk_size: 4096
  enable_activation_offload: False
  activation_offload_plan:
    apply_modules: ['model.visual.blocks.*', 'model.language_model.layers.*']
  enable_chunk_mbs: False
  chunkmbs_plan:
    apply_modules: ['model.language_model.layers.*']
    chunk_mbs: 2
    batch_dim: 0
    chunk_arg_indexes: [0]
    chunk_kwarg_names: ['position_embeddings', 'position_ids', 'rope_deltas', 'attention_mask']
  enable_ep_balance: False
  ep_balance_plan:
    max_dup_experts_num: 2
  skip_moe_pad_tokens: False
  recompute: False
data.dataset_param.preprocess_parameters.model_name_or_path: /root/Qwen3.5-0.8B-hf
data.dataset_param.preprocess_parameters.use_fast_tokenizer: True
data.dataset_param.preprocess_parameters.split_special_tokens: False
data.dataset_param.preprocess_parameters.image_max_pixels: 262144
data.dataset_param.preprocess_parameters.image_min_pixels: 1024
data.dataset_param.preprocess_parameters.video_max_pixels: 16384
data.dataset_param.preprocess_parameters.video_min_pixels: 0
data.dataset_param.preprocess_parameters.video_fps: 2.0
data.dataset_param.preprocess_parameters.video_maxlen: 64

===== Dynamic Fields =====
data.dataset_param.preprocess_parameters.image_max_pixels
data.dataset_param.preprocess_parameters.image_min_pixels
data.dataset_param.preprocess_parameters.model_name_or_path
data.dataset_param.preprocess_parameters.split_special_tokens
data.dataset_param.preprocess_parameters.use_fast_tokenizer
data.dataset_param.preprocess_parameters.video_fps
data.dataset_param.preprocess_parameters.video_max_pixels
data.dataset_param.preprocess_parameters.video_maxlen
data.dataset_param.preprocess_parameters.video_min_pixels
features.recompute
model.causal_conv1d_implementation
model.gdn_implementation
model.skip_flash_attn_recompute
model.skip_gdn_recompute
=====
Successfully imported: mindspeed_mm.fsd.models.qwen3_5
Successfully imported: mindspeed_mm.fsd.data.datasets.huggingface
[Rank 1 | Local Rank 1] 2026-09-10 10:06:19,632 INFO [datasets:114] => Disabling Tensorflow because USE_TORCH is set
[Rank 0 | Local Rank 0] 2026-09-10 10:06:19,704 INFO [datasets:114] => Disabling Tensorflow because USE_TORCH is set
[Rank 0 | Local Rank 0] 2026-09-10 10:06:20,042 INFO [mindspeed_mm.fsd.distributed.parallel_state:87] => Parallel state initialized:
  [dp] = True | Group size: 2 | device mesh:DeviceMesh('npu', [0, 1], mesh_dim_names=('dp',))
  [cp] = False | Group size: 1 | device mesh:DeviceMesh('npu', [0], mesh_dim_names=('cp',))
  [ulysses] = False | Group size: 1 | device mesh:DeviceMesh('npu', [0], mesh_dim_names=('ulysses',))
  [ring] = False | Group size: 1 | device mesh:DeviceMesh('npu', [0], mesh_dim_names=('ring',))
  [tp] = False | Group size: 1 | device mesh:DeviceMesh('npu', [0], mesh_dim_names=('tp',))

```

[illegible]

[illegible]

[illegible]

[Rank 0 | Local Rank 0] 2026-09-10 10:06:42,893 INFO [mindspeed_mm.fsd.train.train_engine:48] => [2026-09-10 10:06:42] iteration 9/ 100
| consumed samples: 72 | elapsed time per iteration (ms): 521.2 | learning rate: 8.000000E-06 | global batch size: 8 | loss: 1.267649E
+00 | grad norm: 12.815 |
[Rank 0 | Local Rank 0] 2026-09-10 10:06:43,434 INFO [mindspeed_mm.fsd.train.train_engine:48] => [2026-09-10 10:06:43] iteration 10/ 100
| consumed samples: 80 | elapsed time per iteration (ms): 540.4 | learning rate: 9.000000E-06 | global batch size: 8 | loss: 1.406928E
+00 | grad norm: 10.876 |
[Rank 0 | Local Rank 0] 2026-09-10 10:06:43,862 INFO [mindspeed_mm.fsd.train.train_engine:48] => [2026-09-10 10:06:43] iteration 11/ 100
| consumed samples: 88 | elapsed time per iteration (ms): 427.1 | learning rate: 1.000000E-05 | global batch size: 8 | loss: 1.434886E
+00 | grad norm: 12.104 |
[Rank 0 | Local Rank 0] 2026-09-10 10:06:44,397 INFO [mindspeed_mm.fsd.train.train_engine:48] => [2026-09-10 10:06:44] iteration 12/ 100
| consumed samples: 96 | elapsed time per iteration (ms): 534.5 | learning rate: 9.996954E-06 | global batch size: 8 | loss: 1.402380E
+00 | grad norm: 11.462 |
[Rank 0 | Local Rank 0] 2026-09-10 10:06:44,824 INFO [mindspeed_mm.fsd.train.train_engine:48] => [2026-09-10 10:06:44] iteration 13/ 100
| consumed samples: 104 | elapsed time per iteration (ms): 425.2 | learning rate: 9.987820E-06 | global batch size: 8 | loss: 1.466555E
+00 | grad norm: 10.055 |
[Rank 0 | Local Rank 0] 2026-09-10 10:06:45,266 INFO [mindspeed_mm.fsd.train.train_engine:48] => [2026-09-10 10:06:45] iteration 14/ 100
| consumed samples: 112 | elapsed time per iteration (ms): 441.2 | learning rate: 9.972609E-06 | global batch size: 8 | loss: 1.472093E
+00 | grad norm: 11.245 |
[Rank 0 | Local Rank 0] 2026-09-10 10:06:45,779 INFO [mindspeed_mm.fsd.train.train_engine:48] => [2026-09-10 10:06:45] iteration 15/ 100
| consumed samples: 120 | elapsed time per iteration (ms): 512.2 | learning rate: 9.951340E-06 | global batch size: 8 | loss: 1.419261E
+00 | grad norm: 9.822 |
[Rank 0 | Local Rank 0] 2026-09-10 10:06:46,183 INFO [mindspeed_mm.fsd.train.train_engine:48] => [2026-09-10 10:06:46] iteration 16/ 100
| consumed samples: 128 | elapsed time per iteration (ms): 403.6 | learning rate: 9.924039E-06 | global batch size: 8 | loss: 1.430511E
+00 | grad norm: 11.282 |
[Rank 0 | Local Rank 0] 2026-09-10 10:06:46,607 INFO [mindspeed_mm.fsd.train.train_engine:48] => [2026-09-10 10:06:46] iteration 17/ 100
| consumed samples: 136 | elapsed time per iteration (ms): 422.4 | learning rate: 9.890738E-06 | global batch size: 8 | loss: 1.410810E
+00 | grad norm: 10.318 |
[Rank 0 | Local Rank 0] 2026-09-10 10:06:47,026 INFO [mindspeed_mm.fsd.train.train_engine:48] => [2026-09-10 10:06:47] iteration 18/ 100
| consumed samples: 144 | elapsed time per iteration (ms): 418.3 | learning rate: 9.851479E-06 | global batch size: 8 | loss: 1.380086E
+00 | grad norm: 11.885 |
[Rank 0 | Local Rank 0] 2026-09-10 10:06:47,444 INFO [mindspeed_mm.fsd.train.train_engine:48] => [2026-09-10 10:06:47] iteration 19/ 100
| consumed samples: 152 | elapsed time per iteration (ms): 417.2 | learning rate: 9.806308E-06 | global batch size: 8 | loss: 1.460446E
+00 | grad norm: 10.004 |
[Rank 0 | Local Rank 0] 2026-09-10 10:06:47,867 INFO [mindspeed_mm.fsd.train.train_engine:48] => [2026-09-10 10:06:47] iteration 20/ 100
| consumed samples: 160 | elapsed time per iteration (ms): 421.6 | learning rate: 9.755283E-06 | global batch size: 8 | loss: 1.349175E
+00 | grad norm: 10.365 |
[Rank 0 | Local Rank 0] 2026-09-10 10:06:48,266 INFO [mindspeed_mm.fsd.train.train_engine:48] => [2026-09-10 10:06:48] iteration 21/ 100
| consumed samples: 168 | elapsed time per iteration (ms): 398.6 | learning rate: 9.698463E-06 | global batch size: 8 | loss: 1.255066E
+00 | grad norm: 10.719 |
[Rank 0 | Local Rank 0] 2026-09-10 10:06:48,696 INFO [mindspeed_mm.fsd.train.train_engine:48] => [2026-09-10 10:06:48] iteration 22/ 100
| consumed samples: 176 | elapsed time per iteration (ms): 429.0 | learning rate: 9.635919E-06 | global batch size: 8 | loss: 1.431060E
+00 | grad norm: 9.802 |
[Rank 0 | Local Rank 0] 2026-09-10 10:06:49,099 INFO [mindspeed_mm.fsd.train.train_engine:48] => [2026-09-10 10:06:49] iteration 23/ 100
| consumed samples: 184 | elapsed time per iteration (ms): 402.0 | learning rate: 9.567727E-06 | global batch size: 8 | loss: 1.479235E
+00 | grad norm: 11.543 |
[Rank 0 | Local Rank 0] 2026-09-10 10:06:49,517 INFO [mindspeed_mm.fsd.train.train_engine:48] => [2026-09-10 10:06:49] iteration 24/ 100
| consumed samples: 192 | elapsed time per iteration (ms): 417.0 | learning rate: 9.493970E-06 | global batch size: 8 | loss: 1.482061E
+00 | grad norm: 9.780 |
[Rank 0 | Local Rank 0] 2026-09-10 10:06:49,925 INFO [mindspeed_mm.fsd.train.train_engine:48] => [2026-09-10 10:06:49] iteration 25/ 100
| consumed samples: 200 | elapsed time per iteration (ms): 407.7 | learning rate: 9.414738E-06 | global batch size: 8 | loss: 1.420320E
+00 | grad norm: 11.049 |
[Rank 0 | Local Rank 0] 2026-09-10 10:06:50,306 INFO [mindspeed_mm.fsd.train.train_engine:48] => [2026-09-10 10:06:50] iteration 26/ 100
| consumed samples: 208 | elapsed time per iteration (ms): 379.8 | learning rate: 9.330127E-06 | global batch size: 8 | loss: 1.312449E
+00 | grad norm: 11.926 |
[Rank 0 | Local Rank 0] 2026-09-10 10:06:50,734 INFO [mindspeed_mm.fsd.train.train_engine:48] => [2026-09-10 10:06:50] iteration 27/ 100
| consumed samples: 216 | elapsed time per iteration (ms): 426.9 | learning rate: 9.240240E-06 | global batch size: 8 | loss: 1.476195E
+00 | grad norm: 10.415 |
[Rank 0 | Local Rank 0] 2026-09-10 10:06:51,158 INFO [mindspeed_mm.fsd.train.train_engine:48] => [2026-09-10 10:06:51] iteration 28/ 100
| consumed samples: 224 | elapsed time per iteration (ms): 423.8 | learning rate: 9.145188E-06 | global batch size: 8 | loss: 1.352151E
+00 | grad norm: 10.360 |
[Rank 0 | Local Rank 0] 2026-09-10 10:06:51,592 INFO [mindspeed_mm.fsd.train.train_engine:48] => [2026-09-10 10:06:51] iteration 29/ 100
| consumed samples: 232 | elapsed time per iteration (ms): 432.4 | learning rate: 9.045085E-06 | global batch size: 8 | loss: 1.451956E
+00 | grad norm: 10.195 |
[Rank 0 | Local Rank 0] 2026-09-10 10:06:52,001 INFO [mindspeed_mm.fsd.train.train_engine:48] => [2026-09-10 10:06:52] iteration 30/ 100
| consumed samples: 240 | elapsed time per iteration (ms): 408.3 | learning rate: 8.940054E-06 | global batch size: 8 | loss: 1.285332E
+00 | grad norm: 9.372 |
[Rank 0 | Local Rank 0] 2026-09-10 10:06:52,484 INFO [mindspeed_mm.fsd.train.train_engine:48] => [2026-09-10 10:06:52] iteration 31/ 100
| consumed samples: 248 | elapsed time per iteration (ms): 482.0 | learning rate: 8.830222E-06 | global batch size: 8 | loss: 1.320294E
+00 | grad norm: 9.492 |
[Rank 0 | Local Rank 0] 2026-09-10 10:06:52,902 INFO [mindspeed_mm.fsd.train.train_engine:48] => [2026-09-10 10:06:52] iteration 32/ 100
| consumed samples: 256 | elapsed time per iteration (ms): 417.5 | learning rate: 8.715724E-06 | global batch size: 8 | loss: 1.437102E
+00 | grad norm: 9.379 |
[Rank 0 | Local Rank 0] 2026-09-10 10:06:53,338 INFO [mindspeed_mm.fsd.train.train_engine:48] => [2026-09-10 10:06:53] iteration 33/ 100
| consumed samples: 264 | elapsed time per iteration (ms): 435.2 | learning rate: 8.596699E-06 | global batch size: 8 | loss: 1.276882E
+00 | grad norm: 9.736 |
[Rank 0 | Local Rank 0] 2026-09-10 10:06:53,759 INFO [mindspeed_mm.fsd.train.train_engine:48] => [2026-09-10 10:06:53] iteration 34/ 100
| consumed samples: 272 | elapsed time per iteration (ms): 420.3 | learning rate: 8.473292E-06 | global batch size: 8 | loss: 1.292732E
+00 | grad norm: 10.341 |
[Rank 0 | Local Rank 0] 2026-09-10 10:06:54,259 INFO [mindspeed_mm.fsd.train.train_engine:48] => [2026-09-10 10:06:54] iteration 35/ 100
| consumed samples: 280 | elapsed time per iteration (ms): 498.7 | learning rate: 8.345653E-06 | global batch size: 8 | loss: 1.410749E
+00 | grad norm: 9.229 |
[Rank 0 | Local Rank 0] 2026-09-10 10:06:54,690 INFO [mindspeed_mm.fsd.train.train_engine:48] => [2026-09-10 10:06:54] iteration 36/ 100
| consumed samples: 288 | elapsed time per iteration (ms): 429.6 | learning rate: 8.213938E-06 | global batch size: 8 | loss: 1.407417E
+00 | grad norm: 9.931 |
[Rank 0 | Local Rank 0] 2026-09-10 10:06:55,118 INFO [mindspeed_mm.fsd.train.train_engine:48] => [2026-09-10 10:06:55] iteration 37/ 100
| consumed samples: 296 | elapsed time per iteration (ms): 427.4 | learning rate: 8.078307E-06 | global batch size: 8 | loss: 1.294433E
+00 | grad norm: 11.828 |
[Rank 0 | Local Rank 0] 2026-09-10 10:06:55,550 INFO [mindspeed_mm.fsd.train.train_engine:48] => [2026-09-10 10:06:55] iteration 38/ 100
| consumed samples: 304 | elapsed time per iteration (ms): 431.1 | learning rate: 7.938926E-06 | global batch size: 8 | loss: 1.470158E
+00 | grad norm: 10.302 |
[Rank 0 | Local Rank 0] 2026-09-10 10:06:55,985 INFO [mindspeed_mm.fsd.train.train_engine:48] => [2026-09-10 10:06:55] iteration 39/ 100
| consumed samples: 312 | elapsed time per iteration (ms): 434.8 | learning rate: 7.795965E-06 | global batch size: 8 | loss: 1.354276E
+00 | grad norm: 9.555 |

[Rank 0 | Local Rank 0] 2026-09-10 10:06:56,413 INFO [mindspeed_mm.fsd.dp.train.train_engine:48] => [2026-09-10 10:06:56] iteration 40/ 100
| consumed samples: 320 | elapsed time per iteration (ms): 426.7 | learning rate: 7.649596E-06 | global batch size: 8 | loss: 1.514748E
+00 | grad norm: 8.950 |
[Rank 0 | Local Rank 0] 2026-09-10 10:06:56,841 INFO [mindspeed_mm.fsd.dp.train.train_engine:48] => [2026-09-10 10:06:56] iteration 41/ 100
| consumed samples: 328 | elapsed time per iteration (ms): 427.5 | learning rate: 7.500000E-06 | global batch size: 8 | loss: 1.401769E
+00 | grad norm: 10.474 |
[Rank 0 | Local Rank 0] 2026-09-10 10:06:57,268 INFO [mindspeed_mm.fsd.dp.train.train_engine:48] => [2026-09-10 10:06:57] iteration 42/ 100
| consumed samples: 336 | elapsed time per iteration (ms): 425.5 | learning rate: 7.347358E-06 | global batch size: 8 | loss: 1.288610E
+00 | grad norm: 9.200 |
[Rank 0 | Local Rank 0] 2026-09-10 10:06:57,956 INFO [mindspeed_mm.fsd.dp.train.train_engine:48] => [2026-09-10 10:06:57] iteration 43/ 100
| consumed samples: 344 | elapsed time per iteration (ms): 687.5 | learning rate: 7.191856E-06 | global batch size: 8 | loss: 1.322055E
+00 | grad norm: 10.310 |
[Rank 0 | Local Rank 0] 2026-09-10 10:06:58,392 INFO [mindspeed_mm.fsd.dp.train.train_engine:48] => [2026-09-10 10:06:58] iteration 44/ 100
| consumed samples: 352 | elapsed time per iteration (ms): 434.8 | learning rate: 7.033683E-06 | global batch size: 8 | loss: 1.358778E
+00 | grad norm: 11.078 |
[Rank 0 | Local Rank 0] 2026-09-10 10:06:58,804 INFO [mindspeed_mm.fsd.dp.train.train_engine:48] => [2026-09-10 10:06:58] iteration 45/ 100
| consumed samples: 360 | elapsed time per iteration (ms): 411.2 | learning rate: 6.873033E-06 | global batch size: 8 | loss: 1.289471E
+00 | grad norm: 9.466 |
[Rank 0 | Local Rank 0] 2026-09-10 10:06:59,199 INFO [mindspeed_mm.fsd.dp.train.train_engine:48] => [2026-09-10 10:06:59] iteration 46/ 100
| consumed samples: 368 | elapsed time per iteration (ms): 394.1 | learning rate: 6.710101E-06 | global batch size: 8 | loss: 1.330112E
+00 | grad norm: 10.449 |
[Rank 0 | Local Rank 0] 2026-09-10 10:06:59,618 INFO [mindspeed_mm.fsd.dp.train.train_engine:48] => [2026-09-10 10:06:59] iteration 47/ 100
| consumed samples: 376 | elapsed time per iteration (ms): 418.0 | learning rate: 6.545085E-06 | global batch size: 8 | loss: 1.389744E
+00 | grad norm: 10.828 |
[Rank 0 | Local Rank 0] 2026-09-10 10:07:00,036 INFO [mindspeed_mm.fsd.dp.train.train_engine:48] => [2026-09-10 10:07:00] iteration 48/ 100
| consumed samples: 384 | elapsed time per iteration (ms): 417.1 | learning rate: 6.378187E-06 | global batch size: 8 | loss: 1.470739E
+00 | grad norm: 9.599 |
[Rank 0 | Local Rank 0] 2026-09-10 10:07:00,450 INFO [mindspeed_mm.fsd.dp.train.train_engine:48] => [2026-09-10 10:07:00] iteration 49/ 100
| consumed samples: 392 | elapsed time per iteration (ms): 413.2 | learning rate: 6.209609E-06 | global batch size: 8 | loss: 1.272147E
+00 | grad norm: 10.066 |
[Rank 0 | Local Rank 0] 2026-09-10 10:07:01,093 INFO [mindspeed_mm.fsd.dp.train.train_engine:48] => [2026-09-10 10:07:01] iteration 50/ 100
| consumed samples: 400 | elapsed time per iteration (ms): 643.0 | learning rate: 6.039558E-06 | global batch size: 8 | loss: 1.286628E
+00 | grad norm: 9.138 |
[Rank 0 | Local Rank 0] 2026-09-10 10:07:01,513 INFO [mindspeed_mm.fsd.dp.train.train_engine:48] => [2026-09-10 10:07:01] iteration 51/ 100
| consumed samples: 408 | elapsed time per iteration (ms): 418.2 | learning rate: 5.868241E-06 | global batch size: 8 | loss: 1.360852E
+00 | grad norm: 9.252 |
[Rank 0 | Local Rank 0] 2026-09-10 10:07:01,945 INFO [mindspeed_mm.fsd.dp.train.train_engine:48] => [2026-09-10 10:07:01] iteration 52/ 100
| consumed samples: 416 | elapsed time per iteration (ms): 431.3 | learning rate: 5.695866E-06 | global batch size: 8 | loss: 1.557105E
+00 | grad norm: 9.377 |
[Rank 0 | Local Rank 0] 2026-09-10 10:07:02,359 INFO [mindspeed_mm.fsd.dp.train.train_engine:48] => [2026-09-10 10:07:02] iteration 53/ 100
| consumed samples: 424 | elapsed time per iteration (ms): 413.2 | learning rate: 5.522642E-06 | global batch size: 8 | loss: 1.303965E
+00 | grad norm: 10.076 |
[Rank 0 | Local Rank 0] 2026-09-10 10:07:02,779 INFO [mindspeed_mm.fsd.dp.train.train_engine:48] => [2026-09-10 10:07:02] iteration 54/ 100
| consumed samples: 432 | elapsed time per iteration (ms): 419.3 | learning rate: 5.348782E-06 | global batch size: 8 | loss: 1.278784E
+00 | grad norm: 9.573 |
[Rank 0 | Local Rank 0] 2026-09-10 10:07:03,192 INFO [mindspeed_mm.fsd.dp.train.train_engine:48] => [2026-09-10 10:07:03] iteration 55/ 100
| consumed samples: 440 | elapsed time per iteration (ms): 411.9 | learning rate: 5.174497E-06 | global batch size: 8 | loss: 1.300990E
+00 | grad norm: 10.604 |
[Rank 0 | Local Rank 0] 2026-09-10 10:07:03,626 INFO [mindspeed_mm.fsd.dp.train.train_engine:48] => [2026-09-10 10:07:03] iteration 56/ 100
| consumed samples: 448 | elapsed time per iteration (ms): 433.1 | learning rate: 5.000000E-06 | global batch size: 8 | loss: 1.402128E
+00 | grad norm: 9.481 |
[Rank 0 | Local Rank 0] 2026-09-10 10:07:04,064 INFO [mindspeed_mm.fsd.dp.train.train_engine:48] => [2026-09-10 10:07:04] iteration 57/ 100
| consumed samples: 456 | elapsed time per iteration (ms): 437.7 | learning rate: 4.825503E-06 | global batch size: 8 | loss: 1.393018E
+00 | grad norm: 9.644 |
[Rank 0 | Local Rank 0] 2026-09-10 10:07:04,482 INFO [mindspeed_mm.fsd.dp.train.train_engine:48] => [2026-09-10 10:07:04] iteration 58/ 100
| consumed samples: 464 | elapsed time per iteration (ms): 416.9 | learning rate: 4.651218E-06 | global batch size: 8 | loss: 1.377542E
+00 | grad norm: 8.861 |
[Rank 0 | Local Rank 0] 2026-09-10 10:07:04,879 INFO [mindspeed_mm.fsd.dp.train.train_engine:48] => [2026-09-10 10:07:04] iteration 59/ 100
| consumed samples: 472 | elapsed time per iteration (ms): 395.5 | learning rate: 4.477358E-06 | global batch size: 8 | loss: 1.211915E
+00 | grad norm: 10.237 |
[Rank 0 | Local Rank 0] 2026-09-10 10:07:05,272 INFO [mindspeed_mm.fsd.dp.train.train_engine:48] => [2026-09-10 10:07:05] iteration 60/ 100
| consumed samples: 480 | elapsed time per iteration (ms): 392.6 | learning rate: 4.304134E-06 | global batch size: 8 | loss: 1.269879E
+00 | grad norm: 10.532 |
[Rank 0 | Local Rank 0] 2026-09-10 10:07:05,668 INFO [mindspeed_mm.fsd.dp.train.train_engine:48] => [2026-09-10 10:07:05] iteration 61/ 100
| consumed samples: 488 | elapsed time per iteration (ms): 394.9 | learning rate: 4.131759E-06 | global batch size: 8 | loss: 1.301453E
+00 | grad norm: 9.548 |
[Rank 0 | Local Rank 0] 2026-09-10 10:07:06,086 INFO [mindspeed_mm.fsd.dp.train.train_engine:48] => [2026-09-10 10:07:06] iteration 62/ 100
| consumed samples: 496 | elapsed time per iteration (ms): 416.9 | learning rate: 3.960442E-06 | global batch size: 8 | loss: 1.175144E
+00 | grad norm: 11.458 |
[Rank 0 | Local Rank 0] 2026-09-10 10:07:06,497 INFO [mindspeed_mm.fsd.dp.train.train_engine:48] => [2026-09-10 10:07:06] iteration 63/ 100
| consumed samples: 504 | elapsed time per iteration (ms): 410.7 | learning rate: 3.790391E-06 | global batch size: 8 | loss: 1.407577E
+00 | grad norm: 9.719 |
[Rank 0 | Local Rank 0] 2026-09-10 10:07:06,917 INFO [mindspeed_mm.fsd.dp.train.train_engine:48] => [2026-09-10 10:07:06] iteration 64/ 100
| consumed samples: 512 | elapsed time per iteration (ms): 418.5 | learning rate: 3.621813E-06 | global batch size: 8 | loss: 1.333787E
+00 | grad norm: 10.392 |
[Rank 0 | Local Rank 0] 2026-09-10 10:07:07,310 INFO [mindspeed_mm.fsd.dp.train.train_engine:48] => [2026-09-10 10:07:07] iteration 65/ 100
| consumed samples: 520 | elapsed time per iteration (ms): 392.3 | learning rate: 3.454915E-06 | global batch size: 8 | loss: 1.297108E
+00 | grad norm: 10.470 |
[Rank 0 | Local Rank 0] 2026-09-10 10:07:07,735 INFO [mindspeed_mm.fsd.dp.train.train_engine:48] => [2026-09-10 10:07:07] iteration 66/ 100
| consumed samples: 528 | elapsed time per iteration (ms): 424.3 | learning rate: 3.289899E-06 | global batch size: 8 | loss: 1.416344E
+00 | grad norm: 9.587 |
[Rank 0 | Local Rank 0] 2026-09-10 10:07:08,163 INFO [mindspeed_mm.fsd.dp.train.train_engine:48] => [2026-09-10 10:07:08] iteration 67/ 100
| consumed samples: 536 | elapsed time per iteration (ms): 427.4 | learning rate: 3.126967E-06 | global batch size: 8 | loss: 1.294969E
+00 | grad norm: 9.848 |
[Rank 0 | Local Rank 0] 2026-09-10 10:07:08,589 INFO [mindspeed_mm.fsd.dp.train.train_engine:48] => [2026-09-10 10:07:08] iteration 68/ 100
| consumed samples: 544 | elapsed time per iteration (ms): 425.1 | learning rate: 2.966317E-06 | global batch size: 8 | loss: 1.509635E
+00 | grad norm: 10.319 |
[Rank 0 | Local Rank 0] 2026-09-10 10:07:09,033 INFO [mindspeed_mm.fsd.dp.train.train_engine:48] => [2026-09-10 10:07:09] iteration 69/ 100
| consumed samples: 552 | elapsed time per iteration (ms): 443.3 | learning rate: 2.808144E-06 | global batch size: 8 | loss: 1.424917E
+00 | grad norm: 9.366 |
[Rank 0 | Local Rank 0] 2026-09-10 10:07:09,436 INFO [mindspeed_mm.fsd.dp.train.train_engine:48] => [2026-09-10 10:07:09] iteration 70/ 100
| consumed samples: 560 | elapsed time per iteration (ms): 401.3 | learning rate: 2.652642E-06 | global batch size: 8 | loss: 1.333512E
+00 | grad norm: 10.356 |

```
[Rank 0 | Local Rank 0] 2026-09-10 10:07:09,871 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 10:07:09] iteration 71/ 100
| consumed samples: 568 | elapsed time per iteration (ms): 434.3 | learning rate: 2.500000E-06 | global batch size: 8 | loss: 1.277492E
+00 | grad norm: 10.732 |
[Rank 0 | Local Rank 0] 2026-09-10 10:07:10,295 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 10:07:10] iteration 72/ 100
| consumed samples: 576 | elapsed time per iteration (ms): 423.7 | learning rate: 2.350404E-06 | global batch size: 8 | loss: 1.400283E
+00 | grad norm: 9.731 |
[Rank 0 | Local Rank 0] 2026-09-10 10:07:10,712 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 10:07:10] iteration 73/ 100
| consumed samples: 584 | elapsed time per iteration (ms): 415.9 | learning rate: 2.204035E-06 | global batch size: 8 | loss: 1.378706E
+00 | grad norm: 9.602 |
[Rank 0 | Local Rank 0] 2026-09-10 10:07:11,141 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 10:07:11] iteration 74/ 100
| consumed samples: 592 | elapsed time per iteration (ms): 428.0 | learning rate: 2.061074E-06 | global batch size: 8 | loss: 1.327326E
+00 | grad norm: 10.257 |
[Rank 0 | Local Rank 0] 2026-09-10 10:07:11,535 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 10:07:11] iteration 75/ 100
| consumed samples: 600 | elapsed time per iteration (ms): 392.8 | learning rate: 1.921693E-06 | global batch size: 8 | loss: 1.365415E
+00 | grad norm: 9.564 |
[Rank 0 | Local Rank 0] 2026-09-10 10:07:11,963 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 10:07:11] iteration 76/ 100
| consumed samples: 608 | elapsed time per iteration (ms): 427.7 | learning rate: 1.786062E-06 | global batch size: 8 | loss: 1.337496E
+00 | grad norm: 9.749 |
[Rank 0 | Local Rank 0] 2026-09-10 10:07:12,374 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 10:07:12] iteration 77/ 100
| consumed samples: 616 | elapsed time per iteration (ms): 409.2 | learning rate: 1.654347E-06 | global batch size: 8 | loss: 1.400981E
+00 | grad norm: 9.563 |
[Rank 0 | Local Rank 0] 2026-09-10 10:07:12,774 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 10:07:12] iteration 78/ 100
| consumed samples: 624 | elapsed time per iteration (ms): 399.2 | learning rate: 1.526708E-06 | global batch size: 8 | loss: 1.417662E
+00 | grad norm: 9.865 |
[Rank 0 | Local Rank 0] 2026-09-10 10:07:13,209 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 10:07:13] iteration 79/ 100
| consumed samples: 632 | elapsed time per iteration (ms): 434.1 | learning rate: 1.403301E-06 | global batch size: 8 | loss: 1.448659E
+00 | grad norm: 9.673 |
[Rank 0 | Local Rank 0] 2026-09-10 10:07:13,631 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 10:07:13] iteration 80/ 100
| consumed samples: 640 | elapsed time per iteration (ms): 421.1 | learning rate: 1.284276E-06 | global batch size: 8 | loss: 1.275453E
+00 | grad norm: 8.807 |
[Rank 0 | Local Rank 0] 2026-09-10 10:07:14,023 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 10:07:14] iteration 81/ 100
| consumed samples: 648 | elapsed time per iteration (ms): 391.3 | learning rate: 1.169778E-06 | global batch size: 8 | loss: 1.296216E
+00 | grad norm: 10.195 |
[Rank 0 | Local Rank 0] 2026-09-10 10:07:14,455 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 10:07:14] iteration 82/ 100
| consumed samples: 656 | elapsed time per iteration (ms): 431.3 | learning rate: 1.059946E-06 | global batch size: 8 | loss: 1.274168E
+00 | grad norm: 9.961 |
[Rank 0 | Local Rank 0] 2026-09-10 10:07:14,863 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 10:07:14] iteration 83/ 100
| consumed samples: 664 | elapsed time per iteration (ms): 407.2 | learning rate: 9.549150E-07 | global batch size: 8 | loss: 1.264691E
+00 | grad norm: 9.447 |
[Rank 0 | Local Rank 0] 2026-09-10 10:07:15,290 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 10:07:15] iteration 84/ 100
| consumed samples: 672 | elapsed time per iteration (ms): 425.7 | learning rate: 8.548121E-07 | global batch size: 8 | loss: 1.233577E
+00 | grad norm: 9.187 |
[Rank 0 | Local Rank 0] 2026-09-10 10:07:15,716 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 10:07:15] iteration 85/ 100
| consumed samples: 680 | elapsed time per iteration (ms): 425.1 | learning rate: 7.597595E-07 | global batch size: 8 | loss: 1.319390E
+00 | grad norm: 11.679 |
[Rank 0 | Local Rank 0] 2026-09-10 10:07:16,112 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 10:07:16] iteration 86/ 100
| consumed samples: 688 | elapsed time per iteration (ms): 395.5 | learning rate: 6.698730E-07 | global batch size: 8 | loss: 1.360353E
+00 | grad norm: 10.701 |
[Rank 0 | Local Rank 0] 2026-09-10 10:07:16,842 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 10:07:16] iteration 87/ 100
| consumed samples: 696 | elapsed time per iteration (ms): 729.6 | learning rate: 5.852620E-07 | global batch size: 8 | loss: 1.367474E
+00 | grad norm: 8.787 |
[Rank 0 | Local Rank 0] 2026-09-10 10:07:17,271 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 10:07:17] iteration 88/ 100
| consumed samples: 704 | elapsed time per iteration (ms): 427.5 | learning rate: 5.060298E-07 | global batch size: 8 | loss: 1.280551E
+00 | grad norm: 10.096 |
[Rank 0 | Local Rank 0] 2026-09-10 10:07:17,684 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 10:07:17] iteration 89/ 100
| consumed samples: 712 | elapsed time per iteration (ms): 412.6 | learning rate: 4.322727E-07 | global batch size: 8 | loss: 1.111414E
+00 | grad norm: 9.697 |
[Rank 0 | Local Rank 0] 2026-09-10 10:07:18,105 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 10:07:18] iteration 90/ 100
| consumed samples: 720 | elapsed time per iteration (ms): 420.0 | learning rate: 3.640807E-07 | global batch size: 8 | loss: 1.302263E
+00 | grad norm: 9.603 |
[Rank 0 | Local Rank 0] 2026-09-10 10:07:18,561 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 10:07:18] iteration 91/ 100
| consumed samples: 728 | elapsed time per iteration (ms): 455.2 | learning rate: 3.015369E-07 | global batch size: 8 | loss: 1.352852E
+00 | grad norm: 8.144 |
[Rank 0 | Local Rank 0] 2026-09-10 10:07:18,947 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 10:07:18] iteration 92/ 100
| consumed samples: 736 | elapsed time per iteration (ms): 385.1 | learning rate: 2.447174E-07 | global batch size: 8 | loss: 1.289150E
+00 | grad norm: 10.371 |
[Rank 0 | Local Rank 0] 2026-09-10 10:07:19,363 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 10:07:19] iteration 93/ 100
| consumed samples: 744 | elapsed time per iteration (ms): 414.8 | learning rate: 1.936915E-07 | global batch size: 8 | loss: 1.200442E
+00 | grad norm: 10.067 |
[Rank 0 | Local Rank 0] 2026-09-10 10:07:19,806 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 10:07:19] iteration 94/ 100
| consumed samples: 752 | elapsed time per iteration (ms): 441.7 | learning rate: 1.485214E-07 | global batch size: 8 | loss: 1.441499E
+00 | grad norm: 9.690 |
[Rank 0 | Local Rank 0] 2026-09-10 10:07:20,225 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 10:07:20] iteration 95/ 100
| consumed samples: 760 | elapsed time per iteration (ms): 418.2 | learning rate: 1.092620E-07 | global batch size: 8 | loss: 1.313565E
+00 | grad norm: 9.769 |
[Rank 0 | Local Rank 0] 2026-09-10 10:07:20,791 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 10:07:20] iteration 96/ 100
| consumed samples: 768 | elapsed time per iteration (ms): 565.8 | learning rate: 7.596123E-08 | global batch size: 8 | loss: 1.322018E
+00 | grad norm: 8.598 |
[Rank 0 | Local Rank 0] 2026-09-10 10:07:21,207 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 10:07:21] iteration 97/ 100
| consumed samples: 776 | elapsed time per iteration (ms): 414.3 | learning rate: 4.865966E-08 | global batch size: 8 | loss: 1.330229E
+00 | grad norm: 9.995 |
[Rank 0 | Local Rank 0] 2026-09-10 10:07:21,581 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 10:07:21] iteration 98/ 100
| consumed samples: 784 | elapsed time per iteration (ms): 373.3 | learning rate: 2.739052E-08 | global batch size: 8 | loss: 1.306947E
+00 | grad norm: 11.287 |
[Rank 0 | Local Rank 0] 2026-09-10 10:07:22,012 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 10:07:22] iteration 99/ 100
| consumed samples: 792 | elapsed time per iteration (ms): 430.4 | learning rate: 1.217975E-08 | global batch size: 8 | loss: 1.364842E
+00 | grad norm: 10.147 |
[Rank 0 | Local Rank 0] 2026-09-10 10:07:22,432 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 10:07:22] iteration 100/ 100
| consumed samples: 800 | elapsed time per iteration (ms): 419.0 | learning rate: 3.045865E-09 | global batch size: 8 | loss: 1.447562E
+00 | grad norm: 10.073 |
/usr/local/lib/python3.10/dist-packages/torch/distributed/checkpoint/filesystem.py:592: UserWarning: Detected an existing checkpoint in save_path/it
er_0000100/.metadata, overwriting since self.overwrite=True. Past version 2.5 of PyTorch, `overwrite` will default to False. Set this variable to Tr
ue to maintain this functionality or False to raise when an existing checkpoint is found.
```

```
warnings.warn(
/usr/local/lib/python3.10/dist-packages/torch/distributed/checkpoint/filesystem.py:592: UserWarning: Detected an existing checkpoint in save_path/iter_0000100/.metadata, overwriting since self.overwrite=True. Past version 2.5 of PyTorch, `overwrite` will default to False. Set this variable to True to maintain this functionality or False to raise when an existing checkpoint is found.
warnings.warn(
/usr/local/lib/python3.10/dist-packages/torch/distributed/checkpoint/filesystem.py:111: UserWarning: TypedStorage is deprecated. It will be removed in the future and UntypedStorage will be the only storage class. This should only matter to you if you are using storages directly. To access UntypedStorage directly, use tensor.untyped_storage() instead of tensor.storage()
if tensor.storage().size() != tensor.numel():
/usr/local/lib/python3.10/dist-packages/torch/distributed/checkpoint/filesystem.py:111: UserWarning: TypedStorage is deprecated. It will be removed in the future and UntypedStorage will be the only storage class. This should only matter to you if you are using storages directly. To access UntypedStorage directly, use tensor.untyped_storage() instead of tensor.storage()
if tensor.storage().size() != tensor.numel():
[Rank 0 | Local Rank 0] 2026-09-10 10:07:28,756 INFO [mindspeed_mm.fsdcp.checkpoint.dcp_checkpoint:48] => Saved checkpoint to ./save_path/iter_0000100
```