

原始日志 · 补充日志集

原始日志 (补充日志集)

全流程阶段日志汇编 (7 份)

含: 精度锚点 / 几何探索 / 跨卡验证 / 性能优化 / A2 对照 / 官方基线

日期: 2026-09-11

C4-AI 2026 昇腾赛道复赛 · Qwen3.5-0.8B 迁移

20260907_1440_全量数据dp2_100步

文件: 20260907_1440_全量数据dp2_100步.log | 行数: 717 | 大小: 94.5 KB | 全量数据 (157,712 样本) 官方几何 100 步 —— 精度锚点
采集说明: 完整日志 (从程序启动开始, 未经裁剪)

```
W0907 14:40:21.190000 38166 torch/distributed/run.py:766]
W0907 14:40:21.190000 38166 torch/distributed/run.py:766] *****
W0907 14:40:21.190000 38166 torch/distributed/run.py:766] Setting OMP_NUM_THREADS environment variable for each process to be 1 in default, to avoid
your system being overloaded, please further tune the variable for optimal performance in your application as needed.
W0907 14:40:21.190000 38166 torch/distributed/run.py:766] *****
[INFO] Starting configuration loading...
[INFO] Loaded configuration file: examples/qwen3_5/qwen3_5_0_8B_geo_m4g1_dp2_full100.yaml
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdps.data.data_utils.func_utils.mod
el_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.model_name_or_path (type: <class 'str'>, value: /root/Qwen3.5-
0.8B-hf)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdps.data.data_utils.func_utils.mod
el_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.use_fast_tokenizer (type: <class 'bool'>, value: True)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdps.data.data_utils.func_utils.mod
el_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.split_special_tokens (type: <class 'bool'>, value: False)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdps.data.data_utils.func_utils.mod
el_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.image_max_pixels (type: <class 'int'>, value: 262144)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdps.data.data_utils.func_utils.mod
el_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.image_min_pixels (type: <class 'int'>, value: 1024)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdps.data.data_utils.func_utils.mod
el_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.video_max_pixels (type: <class 'int'>, value: 16384)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdps.data.data_utils.func_utils.mod
el_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.video_min_pixels (type: <class 'int'>, value: 0)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdps.data.data_utils.func_utils.mod
el_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.video_fps (type: <class 'float'>, value: 2.0)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdps.data.data_utils.func_utils.mod
el_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.video_maxlen (type: <class 'int'>, value: 64)
[INFO] Automatically registered YAML field: model.gdn_implementation (type: <class 'str'>, value: triton)
[INFO] Automatically registered YAML field: model.causal_conv1d_implementation (type: <class 'str'>, value: triton)
[INFO] Automatically registered YAML field: model.skip_gdn_recompute (type: <class 'bool'>, value: False)
[INFO] Automatically registered YAML field: model.skip_flash_attn_recompute (type: <class 'bool'>, value: False)
[INFO] Automatically registered YAML field: features.recompute (type: <class 'bool'>, value: True)
=====
BaseArguments Configuration Manager Summary
=====

Configuration class: Arguments
Configuration file: examples/qwen3_5/qwen3_5_0_8B_geo_m4g1_dp2_full100.yaml
Allow YAML extensions: True
Allow CLI override: True
Strict CLI validation: True
Allow registering YAML fields: True

Defined field count: 208
Dynamic field count: 14
Total field count: 222

===== Configuration Details =====
parallel:
  data_parallel_size: 2
  fully_shard_parallel_size: 2
  fsdp_plan:
    ignored_modules: []
    apply_modules: ['model.visual', 'model.visual.blocks.*', 'model.language_model.embed_tokens', 'model.language_model', 'model.language_model.la
yers.*', 'lm_head', 'mtp']
    param_dtype: bf16
    reduce_dtype: fp32
    output_dtype: None
    cast_forward_inputs: True
    reshard_after_forward: False
    num_to_forward_prefetch: 1
    num_to_backward_prefetch: 1
    pregather: False
    hook_modules: None
    cpu_offload: False
  tensor_parallel_size: 1
  tp_plan:
    colwise_parallel: []
    rowwise_parallel: []
    sequence_parallel: []
  ring_attention_size: 1
  ulysses_parallel_size: 1
  expert_parallel_size: 1
  expert_fully_shard_parallel_size: 2
  ep_plan:
```

```
    apply_modules: []
    use_npu_fused_ops: True
    dispatcher: alltoall
    apply_efsdp_modules: []
recompute: False
recompute_plan:
    apply_modules: []
    use_reentrant: False
local_rank: 0
global_rank: 0
world_size: 2
model:
    model_id: qwen3_5
    model_name_or_path: /root/Qwen3.5-0.8B-hf
    trust_remote_code: True
    attn_implementation: flash_attention_2
    freeze: ['model.visual']
    mtp_num_layers: 0
    mtp_loss_scaling_factor: 0.1
    gdn_implementation: triton
    causal_conv1d_implementation: triton
    skip_gdn_recompute: False
    skip_flash_attn_recompute: False
data:
    dataset_param:
        dataset_type: huggingface
        basic_parameters:
            cache_dir: ./cache_dir/
            template: qwen3_vl_nothink
            chat_template: None
            enable_thinking: False
            dataset_dir: /root/data/coco
            dataset: [' /root/data/output_llava_coco_data.json']
            cutoff_len: 1024
            train_on_prompt: False
            mask_history: False
            streaming: False
            overwrite_cache: False
            preprocessing_batch_size: 1000
            preprocessing_num_workers: 16
            max_samples: None
            tool_format: None
            val_dataset: None
            val_max_samples: None
            val_rate: None
            packing: None
            neat_packing: False
            preprocess_on_fly: True
            async_preprocess: False
            async_preprocess_buffer_size: None
            stage: sft
        preprocess_parameters:
            model_name_or_path: /root/Qwen3.5-0.8B-hf
            use_fast_tokenizer: True
            split_special_tokens: False
            image_resolution: 512
            video_resolution: 128
            image_max_pixels: 262144
            image_min_pixels: 1024
            video_max_pixels: 16384
            video_min_pixels: 0
            video_fps: 2.0
            video_maxlen: 64
            image_do_pan_and_scan: False
            crop_to_patches: False
            use_audio_in_video: False
            audio_sampling_rate: 16000
            trust_remote_code: False
    attr:
        ranking: False
        system: None
        images: images
        videos: None
        audios: None
        prompt: None
        query: None
        response: None
        history: None
        messages: messages
        tools: None
        role_tag: role
        content_tag: content
        user_tag: user
        assistant_tag: assistant
        observation_tag: observation
        function_tag: function_call
        system_tag: system
        chosen: None
        rejected: None
        formatting: sharegpt
    val_dataset_param: None
    dataloader_param:
        dataloader_mode: sampler
```

```

sampler_type: BaseRandomBatchSampler
shuffle: False
drop_last: True
pin_memory: True
collate_param:
  collator_id: None
  model_name: qwen3vl
  ignore_pad_token_for_loss: True
  pad_to_multiple_of: 8
num_workers: 8
enable_preload: False
training:
  profile:
    enable: False
    start_step: 1
    end_step: 2
    save_path: ./profiling
    record_shapes: True
    with_memory: True
    with_stack: True
    ranks: [0]
    _train_steps: -1
    local_rank: 0
    global_rank: 0
    world_size: 2
    profile_this_rank: False
  lora:
    enable: False
    rank: 8
    alpha: 16
    target_modules: ['q_proj', 'k_proj', 'v_proj']
    dropout: 0.0
    init_lora_weights: True
    pretrained_lora_path: None
    lora_target_modules_support: None
    disable_peft_moe_conversion: True
  lr: 1e-05
  lr_min: 0.0
  lr_start: 0.0
  weight_decay: 0.0
  no_decay_modules: []
  no_decay_params: []
  optimizer: adamw
  mem_fraction_static: 0.8
  matched_adamw_rms: 0.2
  muon_momentum: 0.95
  ns_steps: 5
  adam_fused: True
  adam_beta1: 0.9
  adam_beta2: 0.999
  adam_eps: 1e-08
  clip_grad: 0.0
  clip_grad_foreach: True
  micro_batch_size: 4
  gradient_accumulation_steps: 1
  lr_warmup_ratio: 0.1
  lr_decay_style: cosine
  lr_decay_ratio: 1.0
  enable_full_determinism: False
  ckpt_manager: dcp
  save_async: False
  save_ckpt_dtype: None
  load_checkpoint_path: None
  save_steps: 0
  seed: 42
  max_steps: None
  init_model_with_meta_device: True
  train_iters: 100
  train_epochs: None
  load: /root/Qwen3.5-0.8B-dcp
  load_format: auto
  save_format: auto
  load_strict: False
  load_rank0_and_broadcast: False
  no_load_optim: True
  no_load_rng: True
  save: ./save_path
  no_save_optim: True
  no_save_rng: True
  log_interval: 1
  save_interval: 10000
  val_interval: 0
  val_micro_batch_size: None
  use_deter_comp: False
  allow_hf32: None
  plugin: ['mindspeed_mm/fsdp/models/qwen3_5', 'mindspeed_mm/fsdp/data/datasets/huggingface']
  lr_scaling_plan: None
  quantization_plan: QuantizeConfig(recipe_name=None, apply_modules=[], ignored_modules=[], quant_converters=[], extra_args={}, enable_fsdp_low_precision_all_gather=True, fsdp_low_precision_all_gather_mode='on-demand')
  _train_steps: -1
  local_rank: 0
  global_rank: 0
  world_size: 2

```

```

global_batch_size: 8
tools:
  profile:
    enable: False
    profile_type: static
    ranks: [0]
    static_param:
      level: level1
      with_stack: False
      with_memory: False
      record_shapes: False
      with_cpu: True
      save_path: ./profiling
      start_step: 10
      end_step: 11
      data_simplification: False
      aic_metrics_type: PipeUtilization
      analyse_flag: True
  memory_profile:
    enable: False
    start_step: 1
    end_step: 2
    save_path: ./memory_snapshot
    dump_ranks: [0]
    stacks: all
    max_entries: None
    mem_info: False
features:
  recompute_plan:
    apply_modules: ['model.visual.blocks.*', 'model.language_model.layers.*']
    use_reentrant: False
  loss_cfg:
    loss_type: default
    router_aux_loss_coef: 0.0
    router_aux_loss_offload: False
  enable_chunk_loss: True
  enable_dynamic_chunk_loss: False
  chunkloss_plan:
    apply_module: lm_head
    chunk_size: 1024
    total_chunk_size: 4096
  enable_activation_offload: True
  activation_offload_plan:
    apply_modules: ['model.visual.blocks.*', 'model.language_model.layers.*']
  enable_chunk_mbs: False
  chunkmbs_plan:
    apply_modules: ['model.language_model.layers.*']
    chunk_mbs: 2
    batch_dim: 0
    chunk_arg_indexes: [0]
    chunk_kward_names: ['position_embeddings', 'position_ids', 'rope_deltas', 'attention_mask']
  enable_ep_balance: False
  ep_balance_plan:
    max_dup_experts_num: 2
  skip_moe_pad_tokens: False
  recompute: True
data.dataset_param.preprocess_parameters.model_name_or_path: /root/Qwen3.5-0.8B-hf
data.dataset_param.preprocess_parameters.use_fast_tokenizer: True
data.dataset_param.preprocess_parameters.split_special_tokens: False
data.dataset_param.preprocess_parameters.image_max_pixels: 262144
data.dataset_param.preprocess_parameters.image_min_pixels: 1024
data.dataset_param.preprocess_parameters.video_max_pixels: 16384
data.dataset_param.preprocess_parameters.video_min_pixels: 0
data.dataset_param.preprocess_parameters.video_fps: 2.0
data.dataset_param.preprocess_parameters.video_maxlen: 64

===== Dynamic Fields =====
data.dataset_param.preprocess_parameters.image_max_pixels
data.dataset_param.preprocess_parameters.image_min_pixels
data.dataset_param.preprocess_parameters.model_name_or_path
data.dataset_param.preprocess_parameters.split_special_tokens
data.dataset_param.preprocess_parameters.use_fast_tokenizer
data.dataset_param.preprocess_parameters.video_fps
data.dataset_param.preprocess_parameters.video_max_pixels
data.dataset_param.preprocess_parameters.video_maxlen
data.dataset_param.preprocess_parameters.video_min_pixels
features.recompute
model.causal_conv1d_implementation
model.gdn_implementation
model.skip_flash_attn_recompute
model.skip_gdn_recompute
=====
Successfully imported: mindspeed_mm.fsd.models.qwen3_5
Successfully imported: mindspeed_mm.fsd.data.datasets.huggingface
[Rank 1 | Local Rank 1] 2026-09-07 14:40:28,790 INFO [datasets:114] => Disabling Tensorflow because USE_TORCH is set
[Rank 0 | Local Rank 0] 2026-09-07 14:40:28,843 INFO [datasets:114] => Disabling Tensorflow because USE_TORCH is set
[Rank 0 | Local Rank 0] 2026-09-07 14:40:29,175 INFO [mindspeed_mm.fsd.distributed.parallel_state:87] => Parallel state initialized:
  [dp] = True | Group size: 2 | device mesh:DeviceMesh('npu', [0, 1], mesh_dim_names=('dp',))
  [cp] = False | Group size: 1 | device mesh:DeviceMesh('npu', [0], mesh_dim_names=('cp',))
  [ulysses] = False | Group size: 1 | device mesh:DeviceMesh('npu', [0], mesh_dim_names=('ulysses',))
  [ring] = False | Group size: 1 | device mesh:DeviceMesh('npu', [0], mesh_dim_names=('ring',))
  [tp] = False | Group size: 1 | device mesh:DeviceMesh('npu', [0], mesh_dim_names=('tp',))

```

```
[edp] = False | Group size: 1 | device mesh:DeviceMesh('npu', [0], mesh_dim_names=('edp',))
[efsdp] = True | Group size: 2 | device mesh:DeviceMesh('npu', [0, 1], mesh_dim_names=('efsdp',))
[ep] = False | Group size: 1 | device mesh:DeviceMesh('npu', [0], mesh_dim_names=('ep',))
[efsdp] = True | Group size: 2 | device mesh:DeviceMesh('npu', [0, 1], mesh_dim_names=('dp_shard_cp',))
```

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

| consumed samples: 120 | elapsed time per iteration (ms): 3868.4 | learning rate: 9.951340E-06 | global batch size: 8 | loss: 1.418323
E+00 | grad norm: 9.817 |
[Rank 0 | Local Rank 0] 2026-09-07 14:41:49,793 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:41:49] iteration 16/ 100
| consumed samples: 128 | elapsed time per iteration (ms): 524.1 | learning rate: 9.924039E-06 | global batch size: 8 | loss: 1.430387E
+00 | grad norm: 11.306 |
[Rank 0 | Local Rank 0] 2026-09-07 14:41:50,326 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:41:50] iteration 17/ 100
| consumed samples: 136 | elapsed time per iteration (ms): 532.4 | learning rate: 9.890738E-06 | global batch size: 8 | loss: 1.410449E
+00 | grad norm: 10.325 |
[Rank 0 | Local Rank 0] 2026-09-07 14:41:50,847 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:41:50] iteration 18/ 100
| consumed samples: 144 | elapsed time per iteration (ms): 520.3 | learning rate: 9.851479E-06 | global batch size: 8 | loss: 1.379003E
+00 | grad norm: 11.854 |
[Rank 0 | Local Rank 0] 2026-09-07 14:41:51,654 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:41:51] iteration 19/ 100
| consumed samples: 152 | elapsed time per iteration (ms): 805.7 | learning rate: 9.806308E-06 | global batch size: 8 | loss: 1.462117E
+00 | grad norm: 9.997 |
[Rank 0 | Local Rank 0] 2026-09-07 14:41:52,181 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:41:52] iteration 20/ 100
| consumed samples: 160 | elapsed time per iteration (ms): 526.8 | learning rate: 9.755283E-06 | global batch size: 8 | loss: 1.349811E
+00 | grad norm: 10.371 |
[Rank 0 | Local Rank 0] 2026-09-07 14:41:52,693 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:41:52] iteration 21/ 100
| consumed samples: 168 | elapsed time per iteration (ms): 510.3 | learning rate: 9.698463E-06 | global batch size: 8 | loss: 1.255451E
+00 | grad norm: 10.737 |
[Rank 0 | Local Rank 0] 2026-09-07 14:41:53,235 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:41:53] iteration 22/ 100
| consumed samples: 176 | elapsed time per iteration (ms): 541.8 | learning rate: 9.635919E-06 | global batch size: 8 | loss: 1.431602E
+00 | grad norm: 9.805 |
[Rank 0 | Local Rank 0] 2026-09-07 14:41:53,765 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:41:53] iteration 23/ 100
| consumed samples: 184 | elapsed time per iteration (ms): 528.7 | learning rate: 9.567727E-06 | global batch size: 8 | loss: 1.478101E
+00 | grad norm: 11.522 |
[Rank 0 | Local Rank 0] 2026-09-07 14:41:54,287 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:41:54] iteration 24/ 100
| consumed samples: 192 | elapsed time per iteration (ms): 521.4 | learning rate: 9.493970E-06 | global batch size: 8 | loss: 1.481449E
+00 | grad norm: 9.781 |
[Rank 0 | Local Rank 0] 2026-09-07 14:41:54,786 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:41:54] iteration 25/ 100
| consumed samples: 200 | elapsed time per iteration (ms): 498.0 | learning rate: 9.414738E-06 | global batch size: 8 | loss: 1.420968E
+00 | grad norm: 11.041 |
[Rank 0 | Local Rank 0] 2026-09-07 14:41:55,263 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:41:55] iteration 26/ 100
| consumed samples: 208 | elapsed time per iteration (ms): 475.5 | learning rate: 9.330127E-06 | global batch size: 8 | loss: 1.313488E
+00 | grad norm: 11.950 |
[Rank 0 | Local Rank 0] 2026-09-07 14:41:55,835 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:41:55] iteration 27/ 100
| consumed samples: 216 | elapsed time per iteration (ms): 571.4 | learning rate: 9.240240E-06 | global batch size: 8 | loss: 1.476822E
+00 | grad norm: 10.422 |
[Rank 0 | Local Rank 0] 2026-09-07 14:41:56,400 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:41:56] iteration 28/ 100
| consumed samples: 224 | elapsed time per iteration (ms): 563.1 | learning rate: 9.145188E-06 | global batch size: 8 | loss: 1.353210E
+00 | grad norm: 10.365 |
[Rank 0 | Local Rank 0] 2026-09-07 14:41:56,952 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:41:56] iteration 29/ 100
| consumed samples: 232 | elapsed time per iteration (ms): 551.2 | learning rate: 9.045085E-06 | global batch size: 8 | loss: 1.453027E
+00 | grad norm: 10.195 |
[Rank 0 | Local Rank 0] 2026-09-07 14:41:57,471 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:41:57] iteration 30/ 100
| consumed samples: 240 | elapsed time per iteration (ms): 518.4 | learning rate: 8.940054E-06 | global batch size: 8 | loss: 1.285901E
+00 | grad norm: 9.373 |
[Rank 0 | Local Rank 0] 2026-09-07 14:41:58,071 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:41:58] iteration 31/ 100
| consumed samples: 248 | elapsed time per iteration (ms): 598.9 | learning rate: 8.830222E-06 | global batch size: 8 | loss: 1.319767E
+00 | grad norm: 9.474 |
[Rank 0 | Local Rank 0] 2026-09-07 14:41:58,599 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:41:58] iteration 32/ 100
| consumed samples: 256 | elapsed time per iteration (ms): 527.4 | learning rate: 8.715724E-06 | global batch size: 8 | loss: 1.436119E
+00 | grad norm: 9.365 |
[Rank 0 | Local Rank 0] 2026-09-07 14:41:59,143 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:41:59] iteration 33/ 100
| consumed samples: 264 | elapsed time per iteration (ms): 542.5 | learning rate: 8.596699E-06 | global batch size: 8 | loss: 1.276496E
+00 | grad norm: 9.753 |
[Rank 0 | Local Rank 0] 2026-09-07 14:41:59,664 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:41:59] iteration 34/ 100
| consumed samples: 272 | elapsed time per iteration (ms): 520.1 | learning rate: 8.473292E-06 | global batch size: 8 | loss: 1.293541E
+00 | grad norm: 10.339 |
[Rank 0 | Local Rank 0] 2026-09-07 14:42:00,246 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:00] iteration 35/ 100
| consumed samples: 280 | elapsed time per iteration (ms): 580.8 | learning rate: 8.345653E-06 | global batch size: 8 | loss: 1.412303E
+00 | grad norm: 9.274 |
[Rank 0 | Local Rank 0] 2026-09-07 14:42:00,808 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:00] iteration 36/ 100
| consumed samples: 288 | elapsed time per iteration (ms): 561.4 | learning rate: 8.213938E-06 | global batch size: 8 | loss: 1.407448E
+00 | grad norm: 9.934 |
[Rank 0 | Local Rank 0] 2026-09-07 14:42:01,354 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:01] iteration 37/ 100
| consumed samples: 296 | elapsed time per iteration (ms): 544.5 | learning rate: 8.078307E-06 | global batch size: 8 | loss: 1.294207E
+00 | grad norm: 11.831 |
[Rank 0 | Local Rank 0] 2026-09-07 14:42:01,906 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:01] iteration 38/ 100
| consumed samples: 304 | elapsed time per iteration (ms): 550.9 | learning rate: 7.938926E-06 | global batch size: 8 | loss: 1.470340E
+00 | grad norm: 10.311 |
[Rank 0 | Local Rank 0] 2026-09-07 14:42:02,468 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:02] iteration 39/ 100
| consumed samples: 312 | elapsed time per iteration (ms): 560.3 | learning rate: 7.795965E-06 | global batch size: 8 | loss: 1.353412E
+00 | grad norm: 9.558 |
[Rank 0 | Local Rank 0] 2026-09-07 14:42:03,013 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:03] iteration 40/ 100
| consumed samples: 320 | elapsed time per iteration (ms): 543.7 | learning rate: 7.649596E-06 | global batch size: 8 | loss: 1.515226E
+00 | grad norm: 8.964 |
[Rank 0 | Local Rank 0] 2026-09-07 14:42:03,906 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:03] iteration 41/ 100
| consumed samples: 328 | elapsed time per iteration (ms): 892.9 | learning rate: 7.500000E-06 | global batch size: 8 | loss: 1.402188E
+00 | grad norm: 10.483 |
[Rank 0 | Local Rank 0] 2026-09-07 14:42:04,446 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:04] iteration 42/ 100
| consumed samples: 336 | elapsed time per iteration (ms): 538.7 | learning rate: 7.347358E-06 | global batch size: 8 | loss: 1.288017E
+00 | grad norm: 9.202 |
[Rank 0 | Local Rank 0] 2026-09-07 14:42:04,942 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:04] iteration 43/ 100
| consumed samples: 344 | elapsed time per iteration (ms): 494.5 | learning rate: 7.191856E-06 | global batch size: 8 | loss: 1.321906E
+00 | grad norm: 10.298 |
[Rank 0 | Local Rank 0] 2026-09-07 14:42:05,511 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:05] iteration 44/ 100
| consumed samples: 352 | elapsed time per iteration (ms): 568.4 | learning rate: 7.033683E-06 | global batch size: 8 | loss: 1.358801E
+00 | grad norm: 11.099 |
[Rank 0 | Local Rank 0] 2026-09-07 14:42:06,058 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:06] iteration 45/ 100
| consumed samples: 360 | elapsed time per iteration (ms): 545.7 | learning rate: 6.873033E-06 | global batch size: 8 | loss: 1.288986E
+00 | grad norm: 9.460 |
[Rank 0 | Local Rank 0] 2026-09-07 14:42:06,575 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:06] iteration 46/ 100

| consumed samples: 368 | elapsed time per iteration (ms): 515.8 | learning rate: 6.710101E-06 | global batch size: 8 | loss: 1.329785E+00 | grad norm: 10.437 |

[Rank 0 | Local Rank 0] 2026-09-07 14:42:07,112 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:07] iteration 47/ 100
| consumed samples: 376 | elapsed time per iteration (ms): 535.3 | learning rate: 6.545085E-06 | global batch size: 8 | loss: 1.390461E+00 | grad norm: 10.827 |

[Rank 0 | Local Rank 0] 2026-09-07 14:42:07,660 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:07] iteration 48/ 100
| consumed samples: 384 | elapsed time per iteration (ms): 546.9 | learning rate: 6.378187E-06 | global batch size: 8 | loss: 1.470727E+00 | grad norm: 9.590 |

[Rank 0 | Local Rank 0] 2026-09-07 14:42:08,189 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:08] iteration 49/ 100
| consumed samples: 392 | elapsed time per iteration (ms): 527.6 | learning rate: 6.209609E-06 | global batch size: 8 | loss: 1.271199E+00 | grad norm: 10.058 |

[Rank 0 | Local Rank 0] 2026-09-07 14:42:09,207 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:09] iteration 50/ 100
| consumed samples: 400 | elapsed time per iteration (ms): 1017.0 | learning rate: 6.039558E-06 | global batch size: 8 | loss: 1.284779E+00 | grad norm: 9.134 |

[Rank 0 | Local Rank 0] 2026-09-07 14:42:09,768 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:09] iteration 51/ 100
| consumed samples: 408 | elapsed time per iteration (ms): 560.2 | learning rate: 5.868241E-06 | global batch size: 8 | loss: 1.360786E+00 | grad norm: 9.235 |

[Rank 0 | Local Rank 0] 2026-09-07 14:42:10,329 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:10] iteration 52/ 100
| consumed samples: 416 | elapsed time per iteration (ms): 559.5 | learning rate: 5.695866E-06 | global batch size: 8 | loss: 1.557341E+00 | grad norm: 9.367 |

[Rank 0 | Local Rank 0] 2026-09-07 14:42:10,854 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:10] iteration 53/ 100
| consumed samples: 424 | elapsed time per iteration (ms): 524.4 | learning rate: 5.522642E-06 | global batch size: 8 | loss: 1.302375E+00 | grad norm: 10.062 |

[Rank 0 | Local Rank 0] 2026-09-07 14:42:11,378 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:11] iteration 54/ 100
| consumed samples: 432 | elapsed time per iteration (ms): 523.4 | learning rate: 5.348782E-06 | global batch size: 8 | loss: 1.279331E+00 | grad norm: 9.579 |

[Rank 0 | Local Rank 0] 2026-09-07 14:42:11,896 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:11] iteration 55/ 100
| consumed samples: 440 | elapsed time per iteration (ms): 516.4 | learning rate: 5.174497E-06 | global batch size: 8 | loss: 1.301835E+00 | grad norm: 10.614 |

[Rank 0 | Local Rank 0] 2026-09-07 14:42:12,463 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:12] iteration 56/ 100
| consumed samples: 448 | elapsed time per iteration (ms): 565.7 | learning rate: 5.000000E-06 | global batch size: 8 | loss: 1.403436E+00 | grad norm: 9.496 |

[Rank 0 | Local Rank 0] 2026-09-07 14:42:13,008 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:13] iteration 57/ 100
| consumed samples: 456 | elapsed time per iteration (ms): 543.6 | learning rate: 4.825503E-06 | global batch size: 8 | loss: 1.392962E+00 | grad norm: 9.645 |

[Rank 0 | Local Rank 0] 2026-09-07 14:42:13,522 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:13] iteration 58/ 100
| consumed samples: 464 | elapsed time per iteration (ms): 513.6 | learning rate: 4.651218E-06 | global batch size: 8 | loss: 1.378616E+00 | grad norm: 8.869 |

[Rank 0 | Local Rank 0] 2026-09-07 14:42:14,019 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:14] iteration 59/ 100
| consumed samples: 472 | elapsed time per iteration (ms): 496.1 | learning rate: 4.477358E-06 | global batch size: 8 | loss: 1.214342E+00 | grad norm: 10.254 |

[Rank 0 | Local Rank 0] 2026-09-07 14:42:14,515 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:14] iteration 60/ 100
| consumed samples: 480 | elapsed time per iteration (ms): 495.1 | learning rate: 4.304134E-06 | global batch size: 8 | loss: 1.269960E+00 | grad norm: 10.553 |

[Rank 0 | Local Rank 0] 2026-09-07 14:42:15,029 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:15] iteration 61/ 100
| consumed samples: 488 | elapsed time per iteration (ms): 512.6 | learning rate: 4.131759E-06 | global batch size: 8 | loss: 1.301221E+00 | grad norm: 9.550 |

[Rank 0 | Local Rank 0] 2026-09-07 14:42:15,571 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:15] iteration 62/ 100
| consumed samples: 496 | elapsed time per iteration (ms): 541.3 | learning rate: 3.960442E-06 | global batch size: 8 | loss: 1.175182E+00 | grad norm: 11.453 |

[Rank 0 | Local Rank 0] 2026-09-07 14:42:16,111 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:16] iteration 63/ 100
| consumed samples: 504 | elapsed time per iteration (ms): 538.5 | learning rate: 3.790391E-06 | global batch size: 8 | loss: 1.407070E+00 | grad norm: 9.722 |

[Rank 0 | Local Rank 0] 2026-09-07 14:42:16,673 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:16] iteration 64/ 100
| consumed samples: 512 | elapsed time per iteration (ms): 560.7 | learning rate: 3.621813E-06 | global batch size: 8 | loss: 1.333373E+00 | grad norm: 10.390 |

[Rank 0 | Local Rank 0] 2026-09-07 14:42:17,247 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:17] iteration 65/ 100
| consumed samples: 520 | elapsed time per iteration (ms): 572.8 | learning rate: 3.454915E-06 | global batch size: 8 | loss: 1.296349E+00 | grad norm: 10.481 |

[Rank 0 | Local Rank 0] 2026-09-07 14:42:17,806 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:17] iteration 66/ 100
| consumed samples: 528 | elapsed time per iteration (ms): 557.7 | learning rate: 3.289899E-06 | global batch size: 8 | loss: 1.415460E+00 | grad norm: 9.579 |

[Rank 0 | Local Rank 0] 2026-09-07 14:42:18,395 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:18] iteration 67/ 100
| consumed samples: 536 | elapsed time per iteration (ms): 587.3 | learning rate: 3.126967E-06 | global batch size: 8 | loss: 1.293443E+00 | grad norm: 9.833 |

[Rank 0 | Local Rank 0] 2026-09-07 14:42:18,981 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:18] iteration 68/ 100
| consumed samples: 544 | elapsed time per iteration (ms): 584.9 | learning rate: 2.966317E-06 | global batch size: 8 | loss: 1.509499E+00 | grad norm: 10.320 |

[Rank 0 | Local Rank 0] 2026-09-07 14:42:19,550 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:19] iteration 69/ 100
| consumed samples: 552 | elapsed time per iteration (ms): 567.9 | learning rate: 2.808144E-06 | global batch size: 8 | loss: 1.425953E+00 | grad norm: 9.366 |

[Rank 0 | Local Rank 0] 2026-09-07 14:42:20,065 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:20] iteration 70/ 100
| consumed samples: 560 | elapsed time per iteration (ms): 513.8 | learning rate: 2.652642E-06 | global batch size: 8 | loss: 1.333466E+00 | grad norm: 10.345 |

[Rank 0 | Local Rank 0] 2026-09-07 14:42:20,612 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:20] iteration 71/ 100
| consumed samples: 568 | elapsed time per iteration (ms): 546.4 | learning rate: 2.500000E-06 | global batch size: 8 | loss: 1.277310E+00 | grad norm: 10.730 |

[Rank 0 | Local Rank 0] 2026-09-07 14:42:21,131 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:21] iteration 72/ 100
| consumed samples: 576 | elapsed time per iteration (ms): 517.7 | learning rate: 2.350404E-06 | global batch size: 8 | loss: 1.399917E+00 | grad norm: 9.736 |

[Rank 0 | Local Rank 0] 2026-09-07 14:42:22,019 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:22] iteration 73/ 100
| consumed samples: 584 | elapsed time per iteration (ms): 887.1 | learning rate: 2.204035E-06 | global batch size: 8 | loss: 1.377935E+00 | grad norm: 9.605 |

[Rank 0 | Local Rank 0] 2026-09-07 14:42:22,558 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:22] iteration 74/ 100
| consumed samples: 592 | elapsed time per iteration (ms): 537.6 | learning rate: 2.061074E-06 | global batch size: 8 | loss: 1.327169E+00 | grad norm: 10.243 |

[Rank 0 | Local Rank 0] 2026-09-07 14:42:23,082 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:23] iteration 75/ 100
| consumed samples: 600 | elapsed time per iteration (ms): 523.2 | learning rate: 1.921693E-06 | global batch size: 8 | loss: 1.365602E+00 | grad norm: 9.581 |

[Rank 0 | Local Rank 0] 2026-09-07 14:42:23,631 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:23] iteration 76/ 100
| consumed samples: 608 | elapsed time per iteration (ms): 548.5 | learning rate: 1.786062E-06 | global batch size: 8 | loss: 1.338079E+00 | grad norm: 9.835 |

[Rank 0 | Local Rank 0] 2026-09-07 14:42:24,172 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:24] iteration 77/ 100

```
| consumed samples:          616 | elapsed time per iteration (ms): 539.0 | learning rate: 1.654347E-06 | global batch size:      8 | loss: 1.401612E
+00 | grad norm: 9.564 |
[Rank 0 | Local Rank 0] 2026-09-07 14:42:24,739 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:24] iteration      78/     100
| consumed samples:          624 | elapsed time per iteration (ms): 565.9 | learning rate: 1.526708E-06 | global batch size:      8 | loss: 1.417971E
+00 | grad norm: 9.873 |
[Rank 0 | Local Rank 0] 2026-09-07 14:42:25,321 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:25] iteration      79/     100
| consumed samples:          632 | elapsed time per iteration (ms): 581.2 | learning rate: 1.403301E-06 | global batch size:      8 | loss: 1.448473E
+00 | grad norm: 9.666 |
[Rank 0 | Local Rank 0] 2026-09-07 14:42:25,897 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:25] iteration      80/     100
| consumed samples:          640 | elapsed time per iteration (ms): 575.2 | learning rate: 1.284276E-06 | global batch size:      8 | loss: 1.276526E
+00 | grad norm: 8.820 |
[Rank 0 | Local Rank 0] 2026-09-07 14:42:26,677 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:26] iteration      81/     100
| consumed samples:          648 | elapsed time per iteration (ms): 778.9 | learning rate: 1.169778E-06 | global batch size:      8 | loss: 1.296517E
+00 | grad norm: 10.200 |
[Rank 0 | Local Rank 0] 2026-09-07 14:42:27,245 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:27] iteration      82/     100
| consumed samples:          656 | elapsed time per iteration (ms): 565.9 | learning rate: 1.059946E-06 | global batch size:      8 | loss: 1.275604E
+00 | grad norm: 9.979 |
[Rank 0 | Local Rank 0] 2026-09-07 14:42:27,790 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:27] iteration      83/     100
| consumed samples:          664 | elapsed time per iteration (ms): 544.2 | learning rate: 9.549150E-07 | global batch size:      8 | loss: 1.264835E
+00 | grad norm: 9.460 |
[Rank 0 | Local Rank 0] 2026-09-07 14:42:28,350 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:28] iteration      84/     100
| consumed samples:          672 | elapsed time per iteration (ms): 558.6 | learning rate: 8.548121E-07 | global batch size:      8 | loss: 1.234508E
+00 | grad norm: 9.197 |
[Rank 0 | Local Rank 0] 2026-09-07 14:42:28,901 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:28] iteration      85/     100
| consumed samples:          680 | elapsed time per iteration (ms): 549.8 | learning rate: 7.597595E-07 | global batch size:      8 | loss: 1.319488E
+00 | grad norm: 11.696 |
[Rank 0 | Local Rank 0] 2026-09-07 14:42:29,423 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:29] iteration      86/     100
| consumed samples:          688 | elapsed time per iteration (ms): 521.5 | learning rate: 6.698730E-07 | global batch size:      8 | loss: 1.361060E
+00 | grad norm: 10.703 |
[Rank 0 | Local Rank 0] 2026-09-07 14:42:29,982 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:29] iteration      87/     100
| consumed samples:          696 | elapsed time per iteration (ms): 557.1 | learning rate: 5.852620E-07 | global batch size:      8 | loss: 1.366600E
+00 | grad norm: 8.777 |
[Rank 0 | Local Rank 0] 2026-09-07 14:42:30,514 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:30] iteration      88/     100
| consumed samples:          704 | elapsed time per iteration (ms): 530.9 | learning rate: 5.060298E-07 | global batch size:      8 | loss: 1.280035E
+00 | grad norm: 10.083 |
[Rank 0 | Local Rank 0] 2026-09-07 14:42:31,038 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:31] iteration      89/     100
| consumed samples:          712 | elapsed time per iteration (ms): 523.7 | learning rate: 4.322727E-07 | global batch size:      8 | loss: 1.111551E
+00 | grad norm: 9.693 |
[Rank 0 | Local Rank 0] 2026-09-07 14:42:31,575 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:31] iteration      90/     100
| consumed samples:          720 | elapsed time per iteration (ms): 535.4 | learning rate: 3.640807E-07 | global batch size:      8 | loss: 1.302250E
+00 | grad norm: 9.608 |
[Rank 0 | Local Rank 0] 2026-09-07 14:42:32,127 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:32] iteration      91/     100
| consumed samples:          728 | elapsed time per iteration (ms): 551.3 | learning rate: 3.015369E-07 | global batch size:      8 | loss: 1.352673E
+00 | grad norm: 8.151 |
[Rank 0 | Local Rank 0] 2026-09-07 14:42:32,611 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:32] iteration      92/     100
| consumed samples:          736 | elapsed time per iteration (ms): 483.1 | learning rate: 2.447174E-07 | global batch size:      8 | loss: 1.289345E
+00 | grad norm: 10.370 |
[Rank 0 | Local Rank 0] 2026-09-07 14:42:33,131 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:33] iteration      93/     100
| consumed samples:          744 | elapsed time per iteration (ms): 518.7 | learning rate: 1.936915E-07 | global batch size:      8 | loss: 1.199757E
+00 | grad norm: 10.068 |
[Rank 0 | Local Rank 0] 2026-09-07 14:42:33,723 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:33] iteration      94/     100
| consumed samples:          752 | elapsed time per iteration (ms): 590.9 | learning rate: 1.485214E-07 | global batch size:      8 | loss: 1.442541E
+00 | grad norm: 9.706 |
[Rank 0 | Local Rank 0] 2026-09-07 14:42:34,285 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:34] iteration      95/     100
| consumed samples:          760 | elapsed time per iteration (ms): 561.5 | learning rate: 1.092620E-07 | global batch size:      8 | loss: 1.314139E
+00 | grad norm: 9.781 |
[Rank 0 | Local Rank 0] 2026-09-07 14:42:38,136 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:38] iteration      96/     100
| consumed samples:          768 | elapsed time per iteration (ms): 3849.2 | learning rate: 7.596123E-08 | global batch size:      8 | loss: 1.322539
E+00 | grad norm: 8.599 |
[Rank 0 | Local Rank 0] 2026-09-07 14:42:38,691 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:38] iteration      97/     100
| consumed samples:          776 | elapsed time per iteration (ms): 551.5 | learning rate: 4.865966E-08 | global batch size:      8 | loss: 1.328328E
+00 | grad norm: 9.990 |
[Rank 0 | Local Rank 0] 2026-09-07 14:42:39,213 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:39] iteration      98/     100
| consumed samples:          784 | elapsed time per iteration (ms): 521.2 | learning rate: 2.739052E-08 | global batch size:      8 | loss: 1.307037E
+00 | grad norm: 11.282 |
[Rank 0 | Local Rank 0] 2026-09-07 14:42:39,798 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:39] iteration      99/     100
| consumed samples:          792 | elapsed time per iteration (ms): 583.2 | learning rate: 1.217975E-08 | global batch size:      8 | loss: 1.363741E
+00 | grad norm: 10.155 |
[Rank 0 | Local Rank 0] 2026-09-07 14:42:40,349 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 14:42:40] iteration     100/     100
| consumed samples:          800 | elapsed time per iteration (ms): 550.6 | learning rate: 3.045865E-09 | global batch size:      8 | loss: 1.447508E
+00 | grad norm: 10.078 |
/usr/local/lib/python3.10/dist-packages/torch/distributed/checkpoint/filesystem.py:592: UserWarning: Detected an existing checkpoint in save_path/it
er_0000100/.metadata, overwriting since self.overwrite=True. Past version 2.5 of PyTorch, `overwrite` will default to False. Set this variable to Tr
ue to maintain this functionality or False to raise when an existing checkpoint is found.
  warnings.warn(
/usr/local/lib/python3.10/dist-packages/torch/distributed/checkpoint/filesystem.py:592: UserWarning: Detected an existing checkpoint in save_path/it
er_0000100/.metadata, overwriting since self.overwrite=True. Past version 2.5 of PyTorch, `overwrite` will default to False. Set this variable to Tr
ue to maintain this functionality or False to raise when an existing checkpoint is found.
  warnings.warn(
[Rank 0 | Local Rank 0] 2026-09-07 14:42:46,603 INFO [mindspeed_mm.fsdsp.checkpoint.dcp_checkpoint:48] => Saved checkpoint to ./save_path/iter_0000
100
```

20260907_1334_官方几何dp2_100步

文件: 20260907_1334_官方几何dp2_100步.log | 行数: 712 | 大小: 94.2 KB | 首次复刻官方并行几何 (双 die) 100 步
采集说明: 完整日志 (从程序启动开始, 未经裁剪)

```
W0907 13:34:28.479000 32506 torch/distributed/run.py:766]
W0907 13:34:28.479000 32506 torch/distributed/run.py:766] *****
W0907 13:34:28.479000 32506 torch/distributed/run.py:766] Setting OMP_NUM_THREADS environment variable for each process to be 1 in default, to avoid
your system being overloaded, please further tune the variable for optimal performance in your application as needed.
W0907 13:34:28.479000 32506 torch/distributed/run.py:766] *****
[INFO] Starting configuration loading...
[INFO] Loaded configuration file: examples/qwen3_5/qwen3_5_0_8B_geo_m4g1_dp2_100.yaml
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdps.data.data_utils.func_utils.mod
el_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.model_name_or_path (type: <class 'str'>, value: /root/Qwen3.5-
0.8B-hf)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdps.data.data_utils.func_utils.mod
el_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.use_fast_tokenizer (type: <class 'bool'>, value: True)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdps.data.data_utils.func_utils.mod
el_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.split_special_tokens (type: <class 'bool'>, value: False)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdps.data.data_utils.func_utils.mod
el_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.image_max_pixels (type: <class 'int'>, value: 262144)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdps.data.data_utils.func_utils.mod
el_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.image_min_pixels (type: <class 'int'>, value: 1024)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdps.data.data_utils.func_utils.mod
el_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.video_max_pixels (type: <class 'int'>, value: 16384)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdps.data.data_utils.func_utils.mod
el_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.video_min_pixels (type: <class 'int'>, value: 0)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdps.data.data_utils.func_utils.mod
el_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.video_fps (type: <class 'float'>, value: 2.0)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdps.data.data_utils.func_utils.mod
el_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.video_maxlen (type: <class 'int'>, value: 64)
[INFO] Automatically registered YAML field: model.gdn_implementation (type: <class 'str'>, value: triton)
[INFO] Automatically registered YAML field: model.causal_conv1d_implementation (type: <class 'str'>, value: triton)
[INFO] Automatically registered YAML field: model.skip_gdn_recompute (type: <class 'bool'>, value: False)
[INFO] Automatically registered YAML field: model.skip_flash_attn_recompute (type: <class 'bool'>, value: False)
[INFO] Automatically registered YAML field: features.recompute (type: <class 'bool'>, value: True)
=====
BaseArguments Configuration Manager Summary
=====

Configuration class: Arguments
Configuration file: examples/qwen3_5/qwen3_5_0_8B_geo_m4g1_dp2_100.yaml
Allow YAML extensions: True
Allow CLI override: True
Strict CLI validation: True
Allow registering YAML fields: True

Defined field count: 208
Dynamic field count: 14
Total field count: 222

===== Configuration Details =====
parallel:
  data_parallel_size: 2
  fully_shard_parallel_size: 2
  fsdp_plan:
    ignored_modules: []
    apply_modules: ['model.visual', 'model.visual.blocks.*', 'model.language_model.embed_tokens', 'model.language_model', 'model.language_model.la
yers.*'], 'lm_head', 'mtp']
    param_dtype: bf16
    reduce_dtype: fp32
    output_dtype: None
    cast_forward_inputs: True
    reshard_after_forward: False
    num_to_forward_prefetch: 1
    num_to_backward_prefetch: 1
    pregather: False
    hook_modules: None
    cpu_offload: False
  tensor_parallel_size: 1
  tp_plan:
    colwise_parallel: []
    rowwise_parallel: []
    sequence_parallel: []
  ring_attention_size: 1
  ulysses_parallel_size: 1
  expert_parallel_size: 1
  expert_fully_shard_parallel_size: 2
  ep_plan:
    apply_modules: []
    use_npu_fused_ops: True
```

```
    dispatcher: alltoall
    apply_efsdp_modules: []
recompute: False
recompute_plan:
    apply_modules: []
    use_reentrant: False
local_rank: 0
global_rank: 0
world_size: 2
model:
    model_id: qwen3_5
    model_name_or_path: /root/Qwen3.5-0.8B-hf
    trust_remote_code: True
    attn_implementation: flash_attention_2
    freeze: ['model.visual']
    mtp_num_layers: 0
    mtp_loss_scaling_factor: 0.1
    gdn_implementation: triton
    causal_conv1d_implementation: triton
    skip_gdn_recompute: False
    skip_flash_attn_recompute: False
data:
    dataset_param:
        dataset_type: huggingface
        basic_parameters:
            cache_dir: ./cache_dir/
            template: qwen3_vl_nothink
            chat_template: None
            enable_thinking: False
            dataset_dir: /root/coco_d1/extracted/dataset
            dataset: ['/root/coco_d1/extracted/dataset/annotations_slim.json']
            cutoff_len: 1024
            train_on_prompt: False
            mask_history: False
            streaming: False
            overwrite_cache: False
            preprocessing_batch_size: 1000
            preprocessing_num_workers: 16
            max_samples: None
            tool_format: None
            val_dataset: None
            val_max_samples: None
            val_rate: None
            packing: None
            neat_packing: False
            preprocess_on_fly: True
            async_preprocess: False
            async_preprocess_buffer_size: None
            stage: sft
        preprocess_parameters:
            model_name_or_path: /root/Qwen3.5-0.8B-hf
            use_fast_tokenizer: True
            split_special_tokens: False
            image_resolution: 512
            video_resolution: 128
            image_max_pixels: 262144
            image_min_pixels: 1024
            video_max_pixels: 16384
            video_min_pixels: 0
            video_fps: 2.0
            video_maxlen: 64
            image_do_pan_and_scan: False
            crop_to_patches: False
            use_audio_in_video: False
            audio_sampling_rate: 16000
            trust_remote_code: False
    attr:
        ranking: False
        system: None
        images: images
        videos: None
        audios: None
        prompt: None
        query: None
        response: None
        history: None
        messages: messages
        tools: None
        role_tag: role
        content_tag: content
        user_tag: user
        assistant_tag: assistant
        observation_tag: observation
        function_tag: function_call
        system_tag: system
        chosen: None
        rejected: None
        formatting: sharegpt
val_dataset_param: None
dataloading_param:
    dataloader_mode: sampler
    sampler_type: BaseRandomBatchSampler
    shuffle: False
```

```

drop_last: True
pin_memory: True
collate_param:
  collator_id: None
  model_name: qwen3v1
  ignore_pad_token_for_loss: True
  pad_to_multiple_of: 8
num_workers: 8
enable_preload: False
training:
  profile:
    enable: False
    start_step: 1
    end_step: 2
    save_path: ./profiling
    record_shapes: True
    with_memory: True
    with_stack: True
    ranks: [0]
    _train_steps: -1
    local_rank: 0
    global_rank: 0
    world_size: 2
    profile_this_rank: False
lora:
  enable: False
  rank: 8
  alpha: 16
  target_modules: ['q_proj', 'k_proj', 'v_proj']
  dropout: 0.0
  init_lora_weights: True
  pretrained_lora_path: None
  lora_target_modules_support: None
  disable_peft_moe_conversion: True
lr: 1e-05
lr_min: 0.0
lr_start: 0.0
weight_decay: 0.0
no_decay_modules: []
no_decay_params: []
optimizer: adamw
mem_fraction_static: 0.8
matched_adamw_rms: 0.2
muon_momentum: 0.95
ns_steps: 5
adam_fused: True
adam_beta1: 0.9
adam_beta2: 0.999
adam_eps: 1e-08
clip_grad: 0.0
clip_grad_foreach: True
micro_batch_size: 4
gradient_accumulation_steps: 1
lr_warmup_ratio: 0.1
lr_decay_style: cosine
lr_decay_ratio: 1.0
enable_full_determinism: False
ckpt_manager: dcp
save_async: False
save_ckpt_dtype: None
load_checkpoint_path: None
save_steps: 0
seed: 42
max_steps: None
init_model_with_meta_device: True
train_iters: 100
train_epochs: None
load: /root/Qwen3.5-0.8B-dcp
load_format: auto
save_format: auto
load_strict: False
load_rank0_and_broadcast: False
no_load_optim: True
no_load_rng: True
save: ./save_path
no_save_optim: True
no_save_rng: True
log_interval: 1
save_interval: 10000
val_interval: 0
val_micro_batch_size: None
use_deter_comp: False
allow_hf32: None
plugin: ['mindspeed_mm/fsdp/models/qwen3_5', 'mindspeed_mm/fsdp/data/datasets/huggingface']
lr_scaling_plan: None
quantization_plan: QuantizeConfig(recipe_name=None, apply_modules=[], ignored_modules=[], quant_converters=[], extra_args={}, enable_fsdp_low_precision_all_gather=True, fsdp_low_precision_all_gather_mode='on-demand')
_train_steps: -1
local_rank: 0
global_rank: 0
world_size: 2
global_batch_size: 8
tools:

```

```

profile:
  enable: False
  profile_type: static
  ranks: [0]
  static_param:
    level: level1
    with_stack: False
    with_memory: False
    record_shapes: False
    with_cpu: True
    save_path: ./profiling
    start_step: 10
    end_step: 11
    data_simplification: False
    aic_metrics_type: PipeUtilization
    analyse_flag: True
memory_profile:
  enable: False
  start_step: 1
  end_step: 2
  save_path: ./memory_snapshot
  dump_ranks: [0]
  stacks: all
  max_entries: None
  mem_info: False
features:
  recompute_plan:
    apply_modules: ['model.visual.blocks.*', 'model.language_model.layers.*']
    use_reentrant: False
  loss_cfg:
    loss_type: default
    router_aux_loss_coef: 0.0
    router_aux_loss_offload: False
  enable_chunk_loss: True
  enable_dynamic_chunk_loss: False
  chunkloss_plan:
    apply_module: lm_head
    chunk_size: 1024
    total_chunk_size: 4096
  enable_activation_offload: True
  activation_offload_plan:
    apply_modules: ['model.visual.blocks.*', 'model.language_model.layers.*']
  enable_chunk_mbs: False
  chunkmbs_plan:
    apply_modules: ['model.language_model.layers.*']
    chunk_mbs: 2
    batch_dim: 0
    chunk_arg_indexes: [0]
    chunk_kwang_names: ['position_embeddings', 'position_ids', 'rope_deltas', 'attention_mask']
  enable_ep_balance: False
  ep_balance_plan:
    max_dup_experts_num: 2
  skip_moe_pad_tokens: False
  recompute: True
data.dataset_param.preprocess_parameters.model_name_or_path: /root/Qwen3.5-0.8B-hf
data.dataset_param.preprocess_parameters.use_fast_tokenizer: True
data.dataset_param.preprocess_parameters.split_special_tokens: False
data.dataset_param.preprocess_parameters.image_max_pixels: 262144
data.dataset_param.preprocess_parameters.image_min_pixels: 1024
data.dataset_param.preprocess_parameters.video_max_pixels: 16384
data.dataset_param.preprocess_parameters.video_min_pixels: 0
data.dataset_param.preprocess_parameters.video_fps: 2.0
data.dataset_param.preprocess_parameters.video_maxlen: 64

```

===== Dynamic Fields =====

```

data.dataset_param.preprocess_parameters.image_max_pixels
data.dataset_param.preprocess_parameters.image_min_pixels
data.dataset_param.preprocess_parameters.model_name_or_path
data.dataset_param.preprocess_parameters.split_special_tokens
data.dataset_param.preprocess_parameters.use_fast_tokenizer
data.dataset_param.preprocess_parameters.video_fps
data.dataset_param.preprocess_parameters.video_max_pixels
data.dataset_param.preprocess_parameters.video_maxlen
data.dataset_param.preprocess_parameters.video_min_pixels
features.recompute
model.causal_conv1d_implementation
model.gdn_implementation
model.skip_flash_attn_recompute
model.skip_gdn_recompute

```

=====

Successfully imported: mindspeed_mm.fsd.models.qwen3_5

Successfully imported: mindspeed_mm.fsd.data.datasets.huggingface

```

[Rank 0 | Local Rank 0] 2026-09-07 13:34:35,997 INFO [datasets:114] => Disabling Tensorflow because USE_TORCH is set
[Rank 1 | Local Rank 1] 2026-09-07 13:34:36,253 INFO [datasets:114] => Disabling Tensorflow because USE_TORCH is set
[Rank 0 | Local Rank 0] 2026-09-07 13:34:36,331 INFO [mindspeed_mm.fsd.distributed.parallel_state:87] => Parallel state initialized:
[dp] = True | Group size: 2 | device mesh:DeviceMesh('np', [0, 1], mesh_dim_names=('dp',))
[cp] = False | Group size: 1 | device mesh:DeviceMesh('np', [0], mesh_dim_names=('cp',))
[ulysses] = False | Group size: 1 | device mesh:DeviceMesh('np', [0], mesh_dim_names=('ulysses',))
[ring] = False | Group size: 1 | device mesh:DeviceMesh('np', [0], mesh_dim_names=('ring',))
[tp] = False | Group size: 1 | device mesh:DeviceMesh('np', [0], mesh_dim_names=('tp',))
[edp] = False | Group size: 1 | device mesh:DeviceMesh('np', [0], mesh_dim_names=('edp',))
[efsdp] = True | Group size: 2 | device mesh:DeviceMesh('np', [0, 1], mesh_dim_names=('efsdp',))

```

[illegible]

[illegible]

[illegible]

[illegible]

```
[Rank 0 | Local Rank 0] 2026-09-07 13:34:36,898 INFO [mindspeed_mm.fsdp.distributed.fully_shard_parallel:48] => model.visual.blocks.5 set backward p
refetch: ['model.visual.blocks.4']
[Rank 0 | Local Rank 0] 2026-09-07 13:34:36,898 INFO [mindspeed_mm.fsdp.distributed.fully_shard_parallel:48] => model.visual.blocks.4 set backward p
refetch: ['model.visual.blocks.3']
[Rank 0 | Local Rank 0] 2026-09-07 13:34:36,898 INFO [mindspeed_mm.fsdp.distributed.fully_shard_parallel:48] => model.visual.blocks.3 set backward p
refetch: ['model.visual.blocks.2']
[Rank 0 | Local Rank 0] 2026-09-07 13:34:36,898 INFO [mindspeed_mm.fsdp.distributed.fully_shard_parallel:48] => model.visual.blocks.2 set backward p
refetch: ['model.visual.blocks.1']
[Rank 0 | Local Rank 0] 2026-09-07 13:34:36,898 INFO [mindspeed_mm.fsdp.distributed.fully_shard_parallel:48] => model.visual.blocks.1 set backward p
refetch: ['model.visual.blocks.0']
[Rank 0 | Local Rank 0] 2026-09-07 13:34:37,078 INFO [_main__:48] => Prepare data
2026-09-07 13:34:40,493 - template - 756 - INFO - Replace eos token: <|im_end|>
[Rank 0 | Local Rank 0] 2026-09-07 13:34:40,493 INFO [mindspeed_mm.fsdp.data.data_utils.func_utils.template:756] => Replace eos token: <|im_end|>
2026-09-07 13:34:40,690 - template - 756 - INFO - Replace eos token: <|im_end|>
[Rank 1 | Local Rank 1] 2026-09-07 13:34:40,690 INFO [mindspeed_mm.fsdp.data.data_utils.func_utils.template:756] => Replace eos token: <|im_end|>
[Rank 0 | Local Rank 0] 2026-09-07 13:34:43,067 INFO [https:1025] => HTTP Request: HEAD https://s3.amazonaws.com/datasets.huggingface.co/datasets/datasets/json/json.py "HTTP/1.1 200 OK"
```

[illegible]

[illegible]

| consumed samples: 144 | elapsed time per iteration (ms): 510.4 | learning rate: 9.851479E-06 | global batch size: 8 | loss: 1.380249E+00 | grad norm: 11.860 |

[Rank 0 | Local Rank 0] 2026-09-07 13:35:29,810 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:35:29] iteration 19/ 100
| consumed samples: 152 | elapsed time per iteration (ms): 527.5 | learning rate: 9.806308E-06 | global batch size: 8 | loss: 1.460717E+00 | grad norm: 10.009 |

[Rank 0 | Local Rank 0] 2026-09-07 13:35:30,338 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:35:30] iteration 20/ 100
| consumed samples: 160 | elapsed time per iteration (ms): 527.0 | learning rate: 9.755283E-06 | global batch size: 8 | loss: 1.350014E+00 | grad norm: 10.371 |

[Rank 0 | Local Rank 0] 2026-09-07 13:35:30,847 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:35:30] iteration 21/ 100
| consumed samples: 168 | elapsed time per iteration (ms): 508.7 | learning rate: 9.698463E-06 | global batch size: 8 | loss: 1.255464E+00 | grad norm: 10.726 |

[Rank 0 | Local Rank 0] 2026-09-07 13:35:31,374 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:35:31] iteration 22/ 100
| consumed samples: 176 | elapsed time per iteration (ms): 525.6 | learning rate: 9.635919E-06 | global batch size: 8 | loss: 1.432014E+00 | grad norm: 9.800 |

[Rank 0 | Local Rank 0] 2026-09-07 13:35:31,863 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:35:31] iteration 23/ 100
| consumed samples: 184 | elapsed time per iteration (ms): 488.6 | learning rate: 9.567727E-06 | global batch size: 8 | loss: 1.478494E+00 | grad norm: 11.524 |

[Rank 0 | Local Rank 0] 2026-09-07 13:35:32,381 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:35:32] iteration 24/ 100
| consumed samples: 192 | elapsed time per iteration (ms): 516.5 | learning rate: 9.493970E-06 | global batch size: 8 | loss: 1.482918E+00 | grad norm: 9.791 |

[Rank 0 | Local Rank 0] 2026-09-07 13:35:32,886 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:35:32] iteration 25/ 100
| consumed samples: 200 | elapsed time per iteration (ms): 504.7 | learning rate: 9.414738E-06 | global batch size: 8 | loss: 1.421127E+00 | grad norm: 11.052 |

[Rank 0 | Local Rank 0] 2026-09-07 13:35:33,383 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:35:33] iteration 26/ 100
| consumed samples: 208 | elapsed time per iteration (ms): 495.8 | learning rate: 9.330127E-06 | global batch size: 8 | loss: 1.314616E+00 | grad norm: 11.972 |

[Rank 0 | Local Rank 0] 2026-09-07 13:35:33,920 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:35:33] iteration 27/ 100
| consumed samples: 216 | elapsed time per iteration (ms): 536.0 | learning rate: 9.240240E-06 | global batch size: 8 | loss: 1.477171E+00 | grad norm: 10.426 |

[Rank 0 | Local Rank 0] 2026-09-07 13:35:34,438 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:35:34] iteration 28/ 100
| consumed samples: 224 | elapsed time per iteration (ms): 516.8 | learning rate: 9.145188E-06 | global batch size: 8 | loss: 1.352318E+00 | grad norm: 10.358 |

[Rank 0 | Local Rank 0] 2026-09-07 13:35:35,017 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:35:35] iteration 29/ 100
| consumed samples: 232 | elapsed time per iteration (ms): 578.2 | learning rate: 9.045085E-06 | global batch size: 8 | loss: 1.452550E+00 | grad norm: 10.194 |

[Rank 0 | Local Rank 0] 2026-09-07 13:35:35,558 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:35:35] iteration 30/ 100
| consumed samples: 240 | elapsed time per iteration (ms): 538.7 | learning rate: 8.940054E-06 | global batch size: 8 | loss: 1.285747E+00 | grad norm: 9.382 |

[Rank 0 | Local Rank 0] 2026-09-07 13:35:36,135 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:35:36] iteration 31/ 100
| consumed samples: 248 | elapsed time per iteration (ms): 576.0 | learning rate: 8.830222E-06 | global batch size: 8 | loss: 1.320499E+00 | grad norm: 9.487 |

[Rank 0 | Local Rank 0] 2026-09-07 13:35:36,698 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:35:36] iteration 32/ 100
| consumed samples: 256 | elapsed time per iteration (ms): 562.5 | learning rate: 8.715724E-06 | global batch size: 8 | loss: 1.435954E+00 | grad norm: 9.368 |

[Rank 0 | Local Rank 0] 2026-09-07 13:35:37,228 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:35:37] iteration 33/ 100
| consumed samples: 264 | elapsed time per iteration (ms): 528.0 | learning rate: 8.596699E-06 | global batch size: 8 | loss: 1.277608E+00 | grad norm: 9.744 |

[Rank 0 | Local Rank 0] 2026-09-07 13:35:37,760 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:35:37] iteration 34/ 100
| consumed samples: 272 | elapsed time per iteration (ms): 531.7 | learning rate: 8.473292E-06 | global batch size: 8 | loss: 1.294020E+00 | grad norm: 10.339 |

[Rank 0 | Local Rank 0] 2026-09-07 13:35:38,357 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:35:38] iteration 35/ 100
| consumed samples: 280 | elapsed time per iteration (ms): 595.6 | learning rate: 8.345653E-06 | global batch size: 8 | loss: 1.411778E+00 | grad norm: 9.257 |

[Rank 0 | Local Rank 0] 2026-09-07 13:35:39,098 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:35:39] iteration 36/ 100
| consumed samples: 288 | elapsed time per iteration (ms): 740.1 | learning rate: 8.213938E-06 | global batch size: 8 | loss: 1.407220E+00 | grad norm: 9.937 |

[Rank 0 | Local Rank 0] 2026-09-07 13:35:39,667 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:35:39] iteration 37/ 100
| consumed samples: 296 | elapsed time per iteration (ms): 568.1 | learning rate: 8.078307E-06 | global batch size: 8 | loss: 1.294890E+00 | grad norm: 11.838 |

[Rank 0 | Local Rank 0] 2026-09-07 13:35:40,202 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:35:40] iteration 38/ 100
| consumed samples: 304 | elapsed time per iteration (ms): 534.4 | learning rate: 7.938926E-06 | global batch size: 8 | loss: 1.469382E+00 | grad norm: 10.295 |

[Rank 0 | Local Rank 0] 2026-09-07 13:35:40,763 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:35:40] iteration 39/ 100
| consumed samples: 312 | elapsed time per iteration (ms): 559.8 | learning rate: 7.795965E-06 | global batch size: 8 | loss: 1.353907E+00 | grad norm: 9.552 |

[Rank 0 | Local Rank 0] 2026-09-07 13:35:41,309 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:35:41] iteration 40/ 100
| consumed samples: 320 | elapsed time per iteration (ms): 545.8 | learning rate: 7.649596E-06 | global batch size: 8 | loss: 1.514503E+00 | grad norm: 8.942 |

[Rank 0 | Local Rank 0] 2026-09-07 13:35:42,042 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:35:42] iteration 41/ 100
| consumed samples: 328 | elapsed time per iteration (ms): 731.6 | learning rate: 7.500000E-06 | global batch size: 8 | loss: 1.402114E+00 | grad norm: 10.470 |

[Rank 0 | Local Rank 0] 2026-09-07 13:35:42,574 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:35:42] iteration 42/ 100
| consumed samples: 336 | elapsed time per iteration (ms): 530.6 | learning rate: 7.347358E-06 | global batch size: 8 | loss: 1.288841E+00 | grad norm: 9.208 |

[Rank 0 | Local Rank 0] 2026-09-07 13:35:43,057 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:35:43] iteration 43/ 100
| consumed samples: 344 | elapsed time per iteration (ms): 482.5 | learning rate: 7.191856E-06 | global batch size: 8 | loss: 1.322921E+00 | grad norm: 10.316 |

[Rank 0 | Local Rank 0] 2026-09-07 13:35:43,605 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:35:43] iteration 44/ 100
| consumed samples: 352 | elapsed time per iteration (ms): 546.8 | learning rate: 7.033683E-06 | global batch size: 8 | loss: 1.357988E+00 | grad norm: 11.113 |

[Rank 0 | Local Rank 0] 2026-09-07 13:35:44,120 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:35:44] iteration 45/ 100
| consumed samples: 360 | elapsed time per iteration (ms): 513.8 | learning rate: 6.873033E-06 | global batch size: 8 | loss: 1.288817E+00 | grad norm: 9.487 |

[Rank 0 | Local Rank 0] 2026-09-07 13:35:44,619 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:35:44] iteration 46/ 100
| consumed samples: 368 | elapsed time per iteration (ms): 498.7 | learning rate: 6.710101E-06 | global batch size: 8 | loss: 1.328859E+00 | grad norm: 10.447 |

[Rank 0 | Local Rank 0] 2026-09-07 13:35:45,137 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:35:45] iteration 47/ 100
| consumed samples: 376 | elapsed time per iteration (ms): 517.1 | learning rate: 6.545085E-06 | global batch size: 8 | loss: 1.389914E+00 | grad norm: 10.816 |

[Rank 0 | Local Rank 0] 2026-09-07 13:35:45,655 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:35:45] iteration 48/ 100
| consumed samples: 384 | elapsed time per iteration (ms): 516.9 | learning rate: 6.378187E-06 | global batch size: 8 | loss: 1.471767E+00 | grad norm: 9.615 |

[Rank 0 | Local Rank 0] 2026-09-07 13:35:46,164 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:35:46] iteration 49/ 100

| consumed samples: 392 | elapsed time per iteration (ms): 508.1 | learning rate: 6.209609E-06 | global batch size: 8 | loss: 1.271814E+00 | grad norm: 10.056 |

[Rank 0 | Local Rank 0] 2026-09-07 13:35:49,968 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:35:49] iteration 50/ 100
| consumed samples: 400 | elapsed time per iteration (ms): 3803.6 | learning rate: 6.039558E-06 | global batch size: 8 | loss: 1.285854E+00 | grad norm: 9.142 |

[Rank 0 | Local Rank 0] 2026-09-07 13:35:50,492 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:35:50] iteration 51/ 100
| consumed samples: 408 | elapsed time per iteration (ms): 522.9 | learning rate: 5.868241E-06 | global batch size: 8 | loss: 1.360607E+00 | grad norm: 9.242 |

[Rank 0 | Local Rank 0] 2026-09-07 13:35:51,049 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:35:51] iteration 52/ 100
| consumed samples: 416 | elapsed time per iteration (ms): 556.1 | learning rate: 5.695866E-06 | global batch size: 8 | loss: 1.557108E+00 | grad norm: 9.372 |

[Rank 0 | Local Rank 0] 2026-09-07 13:35:51,582 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:35:51] iteration 53/ 100
| consumed samples: 424 | elapsed time per iteration (ms): 531.4 | learning rate: 5.522642E-06 | global batch size: 8 | loss: 1.303630E+00 | grad norm: 10.081 |

[Rank 0 | Local Rank 0] 2026-09-07 13:35:52,100 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:35:52] iteration 54/ 100
| consumed samples: 432 | elapsed time per iteration (ms): 517.5 | learning rate: 5.348782E-06 | global batch size: 8 | loss: 1.280230E+00 | grad norm: 9.592 |

[Rank 0 | Local Rank 0] 2026-09-07 13:35:52,609 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:35:52] iteration 55/ 100
| consumed samples: 440 | elapsed time per iteration (ms): 508.3 | learning rate: 5.174497E-06 | global batch size: 8 | loss: 1.300084E+00 | grad norm: 10.591 |

[Rank 0 | Local Rank 0] 2026-09-07 13:35:53,150 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:35:53] iteration 56/ 100
| consumed samples: 448 | elapsed time per iteration (ms): 539.7 | learning rate: 5.000000E-06 | global batch size: 8 | loss: 1.404183E+00 | grad norm: 9.480 |

[Rank 0 | Local Rank 0] 2026-09-07 13:35:53,687 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:35:53] iteration 57/ 100
| consumed samples: 456 | elapsed time per iteration (ms): 536.3 | learning rate: 4.825503E-06 | global batch size: 8 | loss: 1.391760E+00 | grad norm: 9.649 |

[Rank 0 | Local Rank 0] 2026-09-07 13:35:54,211 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:35:54] iteration 58/ 100
| consumed samples: 464 | elapsed time per iteration (ms): 522.5 | learning rate: 4.651218E-06 | global batch size: 8 | loss: 1.379095E+00 | grad norm: 8.858 |

[Rank 0 | Local Rank 0] 2026-09-07 13:35:54,712 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:35:54] iteration 59/ 100
| consumed samples: 472 | elapsed time per iteration (ms): 500.1 | learning rate: 4.477358E-06 | global batch size: 8 | loss: 1.212878E+00 | grad norm: 10.245 |

[Rank 0 | Local Rank 0] 2026-09-07 13:35:55,205 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:35:55] iteration 60/ 100
| consumed samples: 480 | elapsed time per iteration (ms): 492.7 | learning rate: 4.304134E-06 | global batch size: 8 | loss: 1.270043E+00 | grad norm: 10.545 |

[Rank 0 | Local Rank 0] 2026-09-07 13:35:55,695 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:35:55] iteration 61/ 100
| consumed samples: 488 | elapsed time per iteration (ms): 489.1 | learning rate: 4.131759E-06 | global batch size: 8 | loss: 1.300144E+00 | grad norm: 9.545 |

[Rank 0 | Local Rank 0] 2026-09-07 13:35:56,244 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:35:56] iteration 62/ 100
| consumed samples: 496 | elapsed time per iteration (ms): 548.1 | learning rate: 3.960442E-06 | global batch size: 8 | loss: 1.174947E+00 | grad norm: 11.435 |

[Rank 0 | Local Rank 0] 2026-09-07 13:35:56,784 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:35:56] iteration 63/ 100
| consumed samples: 504 | elapsed time per iteration (ms): 538.8 | learning rate: 3.790391E-06 | global batch size: 8 | loss: 1.404970E+00 | grad norm: 9.713 |

[Rank 0 | Local Rank 0] 2026-09-07 13:35:57,346 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:35:57] iteration 64/ 100
| consumed samples: 512 | elapsed time per iteration (ms): 560.4 | learning rate: 3.621813E-06 | global batch size: 8 | loss: 1.332815E+00 | grad norm: 10.388 |

[Rank 0 | Local Rank 0] 2026-09-07 13:35:57,867 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:35:57] iteration 65/ 100
| consumed samples: 520 | elapsed time per iteration (ms): 520.8 | learning rate: 3.454915E-06 | global batch size: 8 | loss: 1.297799E+00 | grad norm: 10.480 |

[Rank 0 | Local Rank 0] 2026-09-07 13:35:58,747 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:35:58] iteration 66/ 100
| consumed samples: 528 | elapsed time per iteration (ms): 878.4 | learning rate: 3.289899E-06 | global batch size: 8 | loss: 1.415701E+00 | grad norm: 9.583 |

[Rank 0 | Local Rank 0] 2026-09-07 13:35:59,314 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:35:59] iteration 67/ 100
| consumed samples: 536 | elapsed time per iteration (ms): 565.8 | learning rate: 3.126967E-06 | global batch size: 8 | loss: 1.293648E+00 | grad norm: 9.839 |

[Rank 0 | Local Rank 0] 2026-09-07 13:35:59,869 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:35:59] iteration 68/ 100
| consumed samples: 544 | elapsed time per iteration (ms): 553.4 | learning rate: 2.966317E-06 | global batch size: 8 | loss: 1.509343E+00 | grad norm: 10.347 |

[Rank 0 | Local Rank 0] 2026-09-07 13:36:00,434 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:36:00] iteration 69/ 100
| consumed samples: 552 | elapsed time per iteration (ms): 564.2 | learning rate: 2.808144E-06 | global batch size: 8 | loss: 1.425236E+00 | grad norm: 9.373 |

[Rank 0 | Local Rank 0] 2026-09-07 13:36:00,930 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:36:00] iteration 70/ 100
| consumed samples: 560 | elapsed time per iteration (ms): 495.1 | learning rate: 2.652642E-06 | global batch size: 8 | loss: 1.334809E+00 | grad norm: 10.353 |

[Rank 0 | Local Rank 0] 2026-09-07 13:36:01,454 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:36:01] iteration 71/ 100
| consumed samples: 568 | elapsed time per iteration (ms): 522.5 | learning rate: 2.500000E-06 | global batch size: 8 | loss: 1.277023E+00 | grad norm: 10.747 |

[Rank 0 | Local Rank 0] 2026-09-07 13:36:02,486 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:36:02] iteration 72/ 100
| consumed samples: 576 | elapsed time per iteration (ms): 1031.1 | learning rate: 2.350404E-06 | global batch size: 8 | loss: 1.400620E+00 | grad norm: 9.745 |

[Rank 0 | Local Rank 0] 2026-09-07 13:36:03,066 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:36:03] iteration 73/ 100
| consumed samples: 584 | elapsed time per iteration (ms): 580.0 | learning rate: 2.204035E-06 | global batch size: 8 | loss: 1.378598E+00 | grad norm: 9.598 |

[Rank 0 | Local Rank 0] 2026-09-07 13:36:03,627 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:36:03] iteration 74/ 100
| consumed samples: 592 | elapsed time per iteration (ms): 559.1 | learning rate: 2.061074E-06 | global batch size: 8 | loss: 1.327354E+00 | grad norm: 10.270 |

[Rank 0 | Local Rank 0] 2026-09-07 13:36:04,147 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:36:04] iteration 75/ 100
| consumed samples: 600 | elapsed time per iteration (ms): 518.9 | learning rate: 1.921693E-06 | global batch size: 8 | loss: 1.365444E+00 | grad norm: 9.564 |

[Rank 0 | Local Rank 0] 2026-09-07 13:36:04,701 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:36:04] iteration 76/ 100
| consumed samples: 608 | elapsed time per iteration (ms): 553.1 | learning rate: 1.786062E-06 | global batch size: 8 | loss: 1.337648E+00 | grad norm: 9.730 |

[Rank 0 | Local Rank 0] 2026-09-07 13:36:05,237 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:36:05] iteration 77/ 100
| consumed samples: 616 | elapsed time per iteration (ms): 535.0 | learning rate: 1.654347E-06 | global batch size: 8 | loss: 1.403029E+00 | grad norm: 9.563 |

[Rank 0 | Local Rank 0] 2026-09-07 13:36:05,754 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:36:05] iteration 78/ 100
| consumed samples: 624 | elapsed time per iteration (ms): 515.1 | learning rate: 1.526708E-06 | global batch size: 8 | loss: 1.418481E+00 | grad norm: 9.867 |

[Rank 0 | Local Rank 0] 2026-09-07 13:36:06,322 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:36:06] iteration 79/ 100
| consumed samples: 632 | elapsed time per iteration (ms): 567.7 | learning rate: 1.403301E-06 | global batch size: 8 | loss: 1.448472E+00 | grad norm: 9.665 |

[Rank 0 | Local Rank 0] 2026-09-07 13:36:06,878 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 13:36:06] iteration 80/ 100

```
| consumed samples:          640 | elapsed time per iteration (ms): 554.9 | learning rate: 1.284276E-06 | global batch size:      8 | loss: 1.276465E
+00 | grad norm: 8.809 |
[Rank 0 | Local Rank 0] 2026-09-07 13:36:07,388 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 13:36:07] iteration      81/    100
| consumed samples:          648 | elapsed time per iteration (ms): 509.0 | learning rate: 1.169778E-06 | global batch size:      8 | loss: 1.295725E
+00 | grad norm: 10.196 |
[Rank 0 | Local Rank 0] 2026-09-07 13:36:07,968 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 13:36:07] iteration      82/    100
| consumed samples:          656 | elapsed time per iteration (ms): 578.5 | learning rate: 1.059946E-06 | global batch size:      8 | loss: 1.275845E
+00 | grad norm: 9.982 |
[Rank 0 | Local Rank 0] 2026-09-07 13:36:08,491 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 13:36:08] iteration      83/    100
| consumed samples:          664 | elapsed time per iteration (ms): 521.9 | learning rate: 9.549150E-07 | global batch size:      8 | loss: 1.264637E
+00 | grad norm: 9.444 |
[Rank 0 | Local Rank 0] 2026-09-07 13:36:09,016 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 13:36:09] iteration      84/    100
| consumed samples:          672 | elapsed time per iteration (ms): 523.8 | learning rate: 8.548121E-07 | global batch size:      8 | loss: 1.233787E
+00 | grad norm: 9.208 |
[Rank 0 | Local Rank 0] 2026-09-07 13:36:09,537 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 13:36:09] iteration      85/    100
| consumed samples:          680 | elapsed time per iteration (ms): 520.5 | learning rate: 7.597595E-07 | global batch size:      8 | loss: 1.318889E
+00 | grad norm: 11.672 |
[Rank 0 | Local Rank 0] 2026-09-07 13:36:10,022 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 13:36:10] iteration      86/    100
| consumed samples:          688 | elapsed time per iteration (ms): 484.2 | learning rate: 6.698730E-07 | global batch size:      8 | loss: 1.361238E
+00 | grad norm: 10.689 |
[Rank 0 | Local Rank 0] 2026-09-07 13:36:10,560 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 13:36:10] iteration      87/    100
| consumed samples:          696 | elapsed time per iteration (ms): 536.7 | learning rate: 5.852620E-07 | global batch size:      8 | loss: 1.367311E
+00 | grad norm: 8.775 |
[Rank 0 | Local Rank 0] 2026-09-07 13:36:11,091 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 13:36:11] iteration      88/    100
| consumed samples:          704 | elapsed time per iteration (ms): 530.7 | learning rate: 5.060298E-07 | global batch size:      8 | loss: 1.279420E
+00 | grad norm: 10.081 |
[Rank 0 | Local Rank 0] 2026-09-07 13:36:11,596 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 13:36:11] iteration      89/    100
| consumed samples:          712 | elapsed time per iteration (ms): 503.6 | learning rate: 4.322727E-07 | global batch size:      8 | loss: 1.110071E
+00 | grad norm: 9.696 |
[Rank 0 | Local Rank 0] 2026-09-07 13:36:12,117 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 13:36:12] iteration      90/    100
| consumed samples:          720 | elapsed time per iteration (ms): 519.8 | learning rate: 3.640807E-07 | global batch size:      8 | loss: 1.301402E
+00 | grad norm: 9.600 |
[Rank 0 | Local Rank 0] 2026-09-07 13:36:12,666 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 13:36:12] iteration      91/    100
| consumed samples:          728 | elapsed time per iteration (ms): 548.6 | learning rate: 3.015369E-07 | global batch size:      8 | loss: 1.352949E
+00 | grad norm: 8.150 |
[Rank 0 | Local Rank 0] 2026-09-07 13:36:13,138 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 13:36:13] iteration      92/    100
| consumed samples:          736 | elapsed time per iteration (ms): 471.3 | learning rate: 2.447174E-07 | global batch size:      8 | loss: 1.289906E
+00 | grad norm: 10.389 |
[Rank 0 | Local Rank 0] 2026-09-07 13:36:13,645 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 13:36:13] iteration      93/    100
| consumed samples:          744 | elapsed time per iteration (ms): 506.3 | learning rate: 1.936915E-07 | global batch size:      8 | loss: 1.199535E
+00 | grad norm: 10.074 |
[Rank 0 | Local Rank 0] 2026-09-07 13:36:14,188 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 13:36:14] iteration      94/    100
| consumed samples:          752 | elapsed time per iteration (ms): 541.5 | learning rate: 1.485214E-07 | global batch size:      8 | loss: 1.441858E
+00 | grad norm: 9.687 |
[Rank 0 | Local Rank 0] 2026-09-07 13:36:14,693 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 13:36:14] iteration      95/    100
| consumed samples:          760 | elapsed time per iteration (ms): 504.9 | learning rate: 1.092620E-07 | global batch size:      8 | loss: 1.313440E
+00 | grad norm: 9.770 |
[Rank 0 | Local Rank 0] 2026-09-07 13:36:15,502 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 13:36:15] iteration      96/    100
| consumed samples:          768 | elapsed time per iteration (ms): 807.9 | learning rate: 7.596123E-08 | global batch size:      8 | loss: 1.323375E
+00 | grad norm: 8.616 |
[Rank 0 | Local Rank 0] 2026-09-07 13:36:16,011 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 13:36:16] iteration      97/    100
| consumed samples:          776 | elapsed time per iteration (ms): 507.5 | learning rate: 4.865966E-08 | global batch size:      8 | loss: 1.327593E
+00 | grad norm: 9.988 |
[Rank 0 | Local Rank 0] 2026-09-07 13:36:16,474 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 13:36:16] iteration      98/    100
| consumed samples:          784 | elapsed time per iteration (ms): 462.7 | learning rate: 2.739052E-08 | global batch size:      8 | loss: 1.307001E
+00 | grad norm: 11.284 |
[Rank 0 | Local Rank 0] 2026-09-07 13:36:17,016 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 13:36:17] iteration      99/    100
| consumed samples:          792 | elapsed time per iteration (ms): 540.7 | learning rate: 1.217975E-08 | global batch size:      8 | loss: 1.364264E
+00 | grad norm: 10.154 |
[Rank 0 | Local Rank 0] 2026-09-07 13:36:17,546 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 13:36:17] iteration     100/    100
| consumed samples:          800 | elapsed time per iteration (ms): 529.6 | learning rate: 3.045865E-09 | global batch size:      8 | loss: 1.447966E
+00 | grad norm: 10.068 |
/usr/local/lib/python3.10/dist-packages/torch/distributed/checkpoint/filesystem.py:592: UserWarning: Detected an existing checkpoint in save_path/iter_0000100/.metadata, overwriting since self.overwrite=True. Past version 2.5 of PyTorch, `overwrite` will default to False. Set this variable to True to maintain this functionality or False to raise when an existing checkpoint is found.
  warnings.warn(
/usr/local/lib/python3.10/dist-packages/torch/distributed/checkpoint/filesystem.py:592: UserWarning: Detected an existing checkpoint in save_path/iter_0000100/.metadata, overwriting since self.overwrite=True. Past version 2.5 of PyTorch, `overwrite` will default to False. Set this variable to True to maintain this functionality or False to raise when an existing checkpoint is found.
  warnings.warn(
/usr/local/lib/python3.10/dist-packages/torch_npu/distributed/distributed_c10d.py:110: UserWarning: HCCL doesn't support gather at the moment. Implemented with allgather instead.
  warnings.warn("HCCL doesn't support gather at the moment. Implemented with allgather instead.")
[Rank 0 | Local Rank 0] 2026-09-07 13:36:23,450 INFO [mindspeed_mm.fsdp.checkpoint.dcp_checkpointinter:48] => Saved checkpoint to ./save_path/iter_0000100
```

20260907_1257_单die几何m8g1_100步

文件: 20260907_1257_单die几何m8g1_100步.log | 行数: 685 | 大小: 90.4 KB | 单 die 几何 (mbs8/gas1) 100 步
采集说明: 完整日志 (从程序启动开始, 未经裁剪)

```
[INFO] Starting configuration loading...
[INFO] Loaded configuration file: examples/qwen3_5/qwen3_5_0_8B_geo_m8g1_100.yaml
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdp.data.data_utils.func_utils.model_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.model_name_or_path (type: <class 'str'>, value: /root/Qwen3.5-0.8B-hf)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdp.data.data_utils.func_utils.model_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.use_fast_tokenizer (type: <class 'bool'>, value: True)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdp.data.data_utils.func_utils.model_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.split_special_tokens (type: <class 'bool'>, value: False)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdp.data.data_utils.func_utils.model_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.image_max_pixels (type: <class 'int'>, value: 262144)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdp.data.data_utils.func_utils.model_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.image_min_pixels (type: <class 'int'>, value: 1024)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdp.data.data_utils.func_utils.model_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.video_max_pixels (type: <class 'int'>, value: 16384)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdp.data.data_utils.func_utils.model_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.video_min_pixels (type: <class 'int'>, value: 0)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdp.data.data_utils.func_utils.model_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.video_fps (type: <class 'float'>, value: 2.0)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdp.data.data_utils.func_utils.model_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.video_maxlen (type: <class 'int'>, value: 64)
[INFO] Automatically registered YAML field: model.gdn_implementation (type: <class 'str'>, value: triton)
[INFO] Automatically registered YAML field: model.causal_conv1d_implementation (type: <class 'str'>, value: triton)
[INFO] Automatically registered YAML field: model.skip_gdn_recompute (type: <class 'bool'>, value: False)
[INFO] Automatically registered YAML field: model.skip_flash_attn_recompute (type: <class 'bool'>, value: False)
[INFO] Automatically registered YAML field: features.recompute (type: <class 'bool'>, value: True)
=====
BaseArguments Configuration Manager Summary
=====

Configuration class: Arguments
Configuration file: examples/qwen3_5/qwen3_5_0_8B_geo_m8g1_100.yaml
Allow YAML extensions: True
Allow CLI override: True
Strict CLI validation: True
Allow registering YAML fields: True

Defined field count: 208
Dynamic field count: 14
Total field count: 222

===== Configuration Details =====
parallel:
  data_parallel_size: 1
  fully_shard_parallel_size: 1
  fsdp_plan:
    ignored_modules: []
    apply_modules: ['model.visual', 'model.visual.blocks.*', 'model.language_model.embed_tokens', 'model.language_model', 'model.language_model.layers.*', 'lm_head', 'mtp']
    param_dtype: bf16
    reduce_dtype: fp32
    output_dtype: None
    cast_forward_inputs: True
    reshard_after_forward: False
    num_to_forward_prefetch: 1
    num_to_backward_prefetch: 1
    pregather: False
    hook_modules: None
    cpu_offload: False
    tensor_parallel_size: 1
  tp_plan:
    colwise_parallel: []
    rowwise_parallel: []
    sequence_parallel: []
  ring_attention_size: 1
  ulysses_parallel_size: 1
  expert_parallel_size: 1
  expert_fully_shard_parallel_size: 1
  ep_plan:
    apply_modules: []
    use_npu_fused_ops: True
    dispatcher: alltoall
    apply_efsdp_modules: []
  recompute: False
  recompute_plan:
    apply_modules: []
```

```
    use_reentrant: False
    local_rank: 0
    global_rank: 0
    world_size: 1
model:
  model_id: qwen3_5
  model_name_or_path: /root/Qwen3.5-0.8B-hf
  trust_remote_code: True
  attn_implementation: flash_attention_2
  freeze: ['model.visual']
  mtp_num_layers: 0
  mtp_loss_scaling_factor: 0.1
  gdn_implementation: triton
  causal_conv1d_implementation: triton
  skip_gdn_recompute: False
  skip_flash_attn_recompute: False
data:
  dataset_param:
    dataset_type: huggingface
    basic_parameters:
      cache_dir: ./cache_dir/
      template: qwen3_vl_nothink
      chat_template: None
      enable_thinking: False
      dataset_dir: /root/coco_d1/extracted/dataset
      dataset: ['/root/coco_d1/extracted/dataset/annotations_slim.json']
      cutoff_len: 1024
      train_on_prompt: False
      mask_history: False
      streaming: False
      overwrite_cache: False
      preprocessing_batch_size: 1000
      preprocessing_num_workers: 2
      max_samples: None
      tool_format: None
      val_dataset: None
      val_max_samples: None
      val_rate: None
      packing: None
      neat_packing: False
      preprocess_on_fly: True
      async_preprocess: False
      async_preprocess_buffer_size: None
      stage: sft
    preprocess_parameters:
      model_name_or_path: /root/Qwen3.5-0.8B-hf
      use_fast_tokenizer: True
      split_special_tokens: False
      image_resolution: 512
      video_resolution: 128
      image_max_pixels: 262144
      image_min_pixels: 1024
      video_max_pixels: 16384
      video_min_pixels: 0
      video_fps: 2.0
      video_maxlen: 64
      image_do_pan_and_scan: False
      crop_to_patches: False
      use_audio_in_video: False
      audio_sampling_rate: 16000
      trust_remote_code: False
  attr:
    ranking: False
    system: None
    images: images
    videos: None
    audios: None
    prompt: None
    query: None
    response: None
    history: None
    messages: messages
    tools: None
    role_tag: role
    content_tag: content
    user_tag: user
    assistant_tag: assistant
    observation_tag: observation
    function_tag: function_call
    system_tag: system
    chosen: None
    rejected: None
    formatting: sharegpt
  val_dataset_param: None
  dataloader_param:
    dataloader_mode: sampler
    sampler_type: BaseRandomBatchSampler
    shuffle: False
    drop_last: True
    pin_memory: True
    collate_param:
      collator_id: None
      model_name: qwen3vl
```

```

        ignore_pad_token_for_loss: True
        pad_to_multiple_of: 8
        num_workers: 2
        enable_preload: False
training:
  profile:
    enable: False
    start_step: 1
    end_step: 2
    save_path: ./profiling
    record_shapes: True
    with_memory: True
    with_stack: True
    ranks: [0]
    _train_steps: -1
    local_rank: 0
    global_rank: 0
    world_size: 1
    profile_this_rank: False
lora:
  enable: False
  rank: 8
  alpha: 16
  target_modules: ['q_proj', 'k_proj', 'v_proj']
  dropout: 0.0
  init_lora_weights: True
  pretrained_lora_path: None
  lora_target_modules_support: None
  disable_peft_moe_conversion: True
lr: 1e-05
lr_min: 0.0
lr_start: 0.0
weight_decay: 0.0
no_decay_modules: []
no_decay_params: []
optimizer: adamw
mem_fraction_static: 0.8
matched_adamw_rms: 0.2
muon_momentum: 0.95
ns_steps: 5
adam_fused: True
adam_beta1: 0.9
adam_beta2: 0.999
adam_eps: 1e-08
clip_grad: 0.0
clip_grad_foreach: True
micro_batch_size: 8
gradient_accumulation_steps: 1
lr_warmup_ratio: 0.1
lr_decay_style: cosine
lr_decay_ratio: 1.0
enable_full_determinism: False
ckpt_manager: dcp
save_async: False
save_ckpt_dtype: None
load_checkpoint_path: None
save_steps: 0
seed: 42
max_steps: None
init_model_with_meta_device: True
train_iters: 100
train_epochs: None
load: /root/Qwen3.5-0.8B-dcp
load_format: auto
save_format: auto
load_strict: False
load_rank0_and_broadcast: False
no_load_optim: True
no_load_rng: True
save: ./save_path
no_save_optim: True
no_save_rng: True
log_interval: 1
save_interval: 10000
val_interval: 0
val_micro_batch_size: None
use_deter_comp: False
allow_hf32: None
plugin: ['mindspeed_mm/fsdp/models/qwen3_5', 'mindspeed_mm/fsdp/data/datasets/huggingface']
lr_scaling_plan: None
quantization_plan: QuantizeConfig(recipe_name=None, apply_modules=[], ignored_modules=[], quant_converters=[], extra_args={}, enable_fsdp_low_precision_all_gather=True, fsdp_low_precision_all_gather_mode='on-demand')
  _train_steps: -1
  local_rank: 0
  global_rank: 0
  world_size: 1
  global_batch_size: 8
tools:
  profile:
    enable: False
    profile_type: static
    ranks: [0]
    static_param:

```

```

    level: level1
    with_stack: False
    with_memory: False
    record_shapes: False
    with_cpu: True
    save_path: ./profiling
    start_step: 10
    end_step: 11
    data_simplification: False
    aic_metrics_type: PipeUtilization
    analyse_flag: True
memory_profile:
    enable: False
    start_step: 1
    end_step: 2
    save_path: ./memory_snapshot
    dump_ranks: [0]
    stacks: all
    max_entries: None
    mem_info: False
features:
    recompute_plan:
        apply_modules: ['model.visual.blocks.*', 'model.language_model.layers.*']
        use_reentrant: False
    loss_cfg:
        loss_type: default
        router_aux_loss_coef: 0.0
        router_aux_loss_offload: False
    enable_chunk_loss: True
    enable_dynamic_chunk_loss: False
    chunkloss_plan:
        apply_module: lm_head
        chunk_size: 1024
        total_chunk_size: 4096
    enable_activation_offload: True
    activation_offload_plan:
        apply_modules: ['model.visual.blocks.*', 'model.language_model.layers.*']
    enable_chunk_mbs: False
    chunkmbs_plan:
        apply_modules: ['model.language_model.layers.*']
        chunk_mbs: 2
        batch_dim: 0
        chunk_arg_indexes: [0]
        chunk_kwarg_names: ['position_embeddings', 'position_ids', 'rope_deltas', 'attention_mask']
    enable_ep_balance: False
    ep_balance_plan:
        max_dup_experts_num: 2
    skip_moe_pad_tokens: False
    recompute: True
data.dataset_param.preprocess_parameters.model_name_or_path: /root/Qwen3.5-0.8B-hf
data.dataset_param.preprocess_parameters.use_fast_tokenizer: True
data.dataset_param.preprocess_parameters.split_special_tokens: False
data.dataset_param.preprocess_parameters.image_max_pixels: 262144
data.dataset_param.preprocess_parameters.image_min_pixels: 1024
data.dataset_param.preprocess_parameters.video_max_pixels: 16384
data.dataset_param.preprocess_parameters.video_min_pixels: 0
data.dataset_param.preprocess_parameters.video_fps: 2.0
data.dataset_param.preprocess_parameters.video_maxlen: 64

===== Dynamic Fields =====
data.dataset_param.preprocess_parameters.image_max_pixels
data.dataset_param.preprocess_parameters.image_min_pixels
data.dataset_param.preprocess_parameters.model_name_or_path
data.dataset_param.preprocess_parameters.split_special_tokens
data.dataset_param.preprocess_parameters.use_fast_tokenizer
data.dataset_param.preprocess_parameters.video_fps
data.dataset_param.preprocess_parameters.video_max_pixels
data.dataset_param.preprocess_parameters.video_maxlen
data.dataset_param.preprocess_parameters.video_min_pixels
features.recompute
model.causal_conv1d_implementation
model.gdn_implementation
model.skip_flash_attn_recompute
model.skip_gdn_recompute
=====
Successfully imported: mindspeed_mm.fsdpm.models.qwen3_5
Successfully imported: mindspeed_mm.fsdpm.data.datasets.huggingface
[Rank 0 | Local Rank 0] 2026-09-07 12:57:14,528 INFO [datasets:114] => Disabling Tensorflow because USE_TORCH is set
[Rank 0 | Local Rank 0] 2026-09-07 12:57:14,890 INFO [mindspeed_mm.fsdpm.distributed.parallel_state:87] => Parallel state initialized:
[dp] = False | Group size: 1 | device mesh:DeviceMesh('npu', [0], mesh_dim_names=('dp',))
[cp] = False | Group size: 1 | device mesh:DeviceMesh('npu', [0], mesh_dim_names=('cp',))
[ulysses] = False | Group size: 1 | device mesh:DeviceMesh('npu', [0], mesh_dim_names=('ulysses',))
[ring] = False | Group size: 1 | device mesh:DeviceMesh('npu', [0], mesh_dim_names=('ring',))
[tp] = False | Group size: 1 | device mesh:DeviceMesh('npu', [0], mesh_dim_names=('tp',))
[edp] = False | Group size: 1 | device mesh:DeviceMesh('npu', [0], mesh_dim_names=('edp',))
[efsdp] = False | Group size: 1 | device mesh:DeviceMesh('npu', [0], mesh_dim_names=('efsdp',))
[ep] = False | Group size: 1 | device mesh:DeviceMesh('npu', [0], mesh_dim_names=('ep',))
[fsdp] = True | Group size: 1 | device mesh:DeviceMesh('npu', [0], mesh_dim_names=('dp_shard_cp',))

[Rank 0 | Local Rank 0] 2026-09-07 12:57:14,890 INFO [mindspeed_mm.fsdpm.models.modelhub:48] => > Loading AutoConfig from /root/Qwen3.5-0.8B-hf...
[Rank 0 | Local Rank 0] 2026-09-07 12:57:14,893 INFO [mindspeed_mm.fsdpm.models.modelhub:48] => Building transformers model from configuration...
[Rank 0 | Local Rank 0] 2026-09-07 12:57:14,933 INFO [mindspeed_mm.fsdpm.models.qwen3_5.modeling_qwen3_5:48] => Qwen3_5 causal_conv1d use NPU triton

```

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

-dcp/release

[2026-09-07 12:57:27] [WARNING] [29097] profiler.py: Invalid parameter export_type: None, reset it to text.

[2026-09-07 12:57:27] [WARNING] [29097] profiler.py: Profiler won't be using warmup, this can skew profiler results

/root/MindSpeed-MM/mindspeed_mm/fsdp/features/memory/async_offload.py:25: UserWarning: TypedStorage is deprecated. It will be removed in the future and UntypedStorage will be the only storage class. This should only matter to you if you are using storages directly. To access UntypedStorage directly, use tensor.untyped_storage() instead of tensor.storage()

if tensor.storage().size() <= 0:

[rank0]:[W907 12:57:28.323114500 compiler_depend.ts:335] Warning: Cannot create tensor with interal format while allow_internel_format=False, tensor will be created with base format. (function operator())

[Rank 0 | Local Rank 0] 2026-09-07 12:57:33,996 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:57:33] iteration 1/ 100
| consumed samples: 8 | elapsed time per iteration (ms): 6727.3 | learning rate: 0.000000E+00 | global batch size: 8 | loss: 1.930328
+00 | grad norm: 133.630 |

[Rank 0 | Local Rank 0] 2026-09-07 12:57:34,991 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:57:34] iteration 2/ 100
| consumed samples: 16 | elapsed time per iteration (ms): 993.6 | learning rate: 1.000000E-06 | global batch size: 8 | loss: 1.788462E
+00 | grad norm: 112.819 |

[Rank 0] (after 2 iterations) memory (MB) | allocated: 9009.94775390625 | max allocated: 29596.4375 | reserved: 34204.0 | max reserved: 34204.0

[Rank 0 | Local Rank 0] 2026-09-07 12:57:35,744 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:57:35] iteration 3/ 100
| consumed samples: 24 | elapsed time per iteration (ms): 751.2 | learning rate: 2.000000E-06 | global batch size: 8 | loss: 1.779180E
+00 | grad norm: 88.238 |

[Rank 0 | Local Rank 0] 2026-09-07 12:57:36,450 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:57:36] iteration 4/ 100
| consumed samples: 32 | elapsed time per iteration (ms): 705.1 | learning rate: 3.000000E-06 | global batch size: 8 | loss: 1.505999E
+00 | grad norm: 57.634 |

[Rank 0 | Local Rank 0] 2026-09-07 12:57:37,152 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:57:37] iteration 5/ 100
| consumed samples: 40 | elapsed time per iteration (ms): 701.3 | learning rate: 4.000000E-06 | global batch size: 8 | loss: 1.409337E
+00 | grad norm: 23.024 |

[Rank 0 | Local Rank 0] 2026-09-07 12:57:37,871 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:57:37] iteration 6/ 100
| consumed samples: 48 | elapsed time per iteration (ms): 718.5 | learning rate: 5.000000E-06 | global batch size: 8 | loss: 1.499240E
+00 | grad norm: 15.791 |

[Rank 0 | Local Rank 0] 2026-09-07 12:57:38,923 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:57:38] iteration 7/ 100
| consumed samples: 56 | elapsed time per iteration (ms): 1051.6 | learning rate: 6.000000E-06 | global batch size: 8 | loss: 1.409245
+00 | grad norm: 15.761 |

[Rank 0 | Local Rank 0] 2026-09-07 12:57:39,696 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:57:39] iteration 8/ 100
| consumed samples: 64 | elapsed time per iteration (ms): 772.3 | learning rate: 7.000000E-06 | global batch size: 8 | loss: 1.460790E
+00 | grad norm: 12.575 |

[Rank 0 | Local Rank 0] 2026-09-07 12:57:40,500 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:57:40] iteration 9/ 100
| consumed samples: 72 | elapsed time per iteration (ms): 802.6 | learning rate: 8.000000E-06 | global batch size: 8 | loss: 1.277361E
+00 | grad norm: 12.856 |

[Rank 0 | Local Rank 0] 2026-09-07 12:57:41,306 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:57:41] iteration 10/ 100
| consumed samples: 80 | elapsed time per iteration (ms): 805.5 | learning rate: 9.000000E-06 | global batch size: 8 | loss: 1.410658E
+00 | grad norm: 10.894 |

[Rank 0 | Local Rank 0] 2026-09-07 12:57:42,063 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:57:42] iteration 11/ 100
| consumed samples: 88 | elapsed time per iteration (ms): 755.8 | learning rate: 1.000000E-05 | global batch size: 8 | loss: 1.433983E
+00 | grad norm: 12.057 |

[Rank 0 | Local Rank 0] 2026-09-07 12:57:43,013 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:57:43] iteration 12/ 100
| consumed samples: 96 | elapsed time per iteration (ms): 949.5 | learning rate: 9.996954E-06 | global batch size: 8 | loss: 1.431983E
+00 | grad norm: 11.358 |

[Rank 0 | Local Rank 0] 2026-09-07 12:57:43,763 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:57:43] iteration 13/ 100
| consumed samples: 104 | elapsed time per iteration (ms): 749.1 | learning rate: 9.987820E-06 | global batch size: 8 | loss: 1.465686E
+00 | grad norm: 9.937 |

[Rank 0 | Local Rank 0] 2026-09-07 12:57:44,535 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:57:44] iteration 14/ 100
| consumed samples: 112 | elapsed time per iteration (ms): 770.6 | learning rate: 9.972609E-06 | global batch size: 8 | loss: 1.448743E
+00 | grad norm: 10.955 |

[Rank 0 | Local Rank 0] 2026-09-07 12:57:45,457 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:57:45] iteration 15/ 100
| consumed samples: 120 | elapsed time per iteration (ms): 920.9 | learning rate: 9.951340E-06 | global batch size: 8 | loss: 1.426714E
+00 | grad norm: 9.847 |

[Rank 0 | Local Rank 0] 2026-09-07 12:57:46,151 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:57:46] iteration 16/ 100
| consumed samples: 128 | elapsed time per iteration (ms): 693.4 | learning rate: 9.924039E-06 | global batch size: 8 | loss: 1.433776E
+00 | grad norm: 11.050 |

[Rank 0 | Local Rank 0] 2026-09-07 12:57:46,877 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:57:46] iteration 17/ 100
| consumed samples: 136 | elapsed time per iteration (ms): 725.1 | learning rate: 9.890738E-06 | global batch size: 8 | loss: 1.411894E
+00 | grad norm: 10.341 |

[Rank 0 | Local Rank 0] 2026-09-07 12:57:47,604 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:57:47] iteration 18/ 100
| consumed samples: 144 | elapsed time per iteration (ms): 726.1 | learning rate: 9.851479E-06 | global batch size: 8 | loss: 1.365979E
+00 | grad norm: 11.710 |

[Rank 0 | Local Rank 0] 2026-09-07 12:57:48,329 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:57:48] iteration 19/ 100
| consumed samples: 152 | elapsed time per iteration (ms): 724.6 | learning rate: 9.806308E-06 | global batch size: 8 | loss: 1.451362E
+00 | grad norm: 9.986 |

[Rank 0 | Local Rank 0] 2026-09-07 12:57:49,066 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:57:49] iteration 20/ 100
| consumed samples: 160 | elapsed time per iteration (ms): 736.2 | learning rate: 9.755283E-06 | global batch size: 8 | loss: 1.339798E
+00 | grad norm: 10.246 |

[Rank 0 | Local Rank 0] 2026-09-07 12:57:49,743 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:57:49] iteration 21/ 100
| consumed samples: 168 | elapsed time per iteration (ms): 676.2 | learning rate: 9.698463E-06 | global batch size: 8 | loss: 1.254943E
+00 | grad norm: 10.706 |

[Rank 0 | Local Rank 0] 2026-09-07 12:57:50,465 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:57:50] iteration 22/ 100
| consumed samples: 176 | elapsed time per iteration (ms): 721.4 | learning rate: 9.635919E-06 | global batch size: 8 | loss: 1.433029E
+00 | grad norm: 9.829 |

[Rank 0 | Local Rank 0] 2026-09-07 12:57:51,152 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:57:51] iteration 23/ 100
| consumed samples: 184 | elapsed time per iteration (ms): 685.7 | learning rate: 9.567727E-06 | global batch size: 8 | loss: 1.480780E
+00 | grad norm: 11.492 |

[Rank 0 | Local Rank 0] 2026-09-07 12:57:51,859 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:57:51] iteration 24/ 100
| consumed samples: 192 | elapsed time per iteration (ms): 706.6 | learning rate: 9.493970E-06 | global batch size: 8 | loss: 1.487493E
+00 | grad norm: 9.735 |

[Rank 0 | Local Rank 0] 2026-09-07 12:57:52,544 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:57:52] iteration 25/ 100
| consumed samples: 200 | elapsed time per iteration (ms): 684.3 | learning rate: 9.414738E-06 | global batch size: 8 | loss: 1.434240E
+00 | grad norm: 10.932 |

[Rank 0 | Local Rank 0] 2026-09-07 12:57:53,174 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:57:53] iteration 26/ 100
| consumed samples: 208 | elapsed time per iteration (ms): 629.4 | learning rate: 9.330127E-06 | global batch size: 8 | loss: 1.309418E
+00 | grad norm: 11.898 |

[Rank 0 | Local Rank 0] 2026-09-07 12:57:53,905 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:57:53] iteration 27/ 100
| consumed samples: 216 | elapsed time per iteration (ms): 729.9 | learning rate: 9.240240E-06 | global batch size: 8 | loss: 1.468136E
+00 | grad norm: 10.299 |

[Rank 0 | Local Rank 0] 2026-09-07 12:57:54,631 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:57:54] iteration 28/ 100
| consumed samples: 224 | elapsed time per iteration (ms): 725.4 | learning rate: 9.145188E-06 | global batch size: 8 | loss: 1.372576E

```
+00 | grad norm: 10.380 |
[Rank 0 | Local Rank 0] 2026-09-07 12:57:55,389 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:57:55] iteration      29/    100
| consumed samples:      232 | elapsed time per iteration (ms): 757.0 | learning rate: 9.045085E-06 | global batch size:    8 | loss: 1.455504E
+00 | grad norm: 10.246 |
[Rank 0 | Local Rank 0] 2026-09-07 12:57:56,114 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:57:56] iteration      30/    100
| consumed samples:      240 | elapsed time per iteration (ms): 724.3 | learning rate: 8.940054E-06 | global batch size:    8 | loss: 1.294052E
+00 | grad norm: 9.437 |
[Rank 0 | Local Rank 0] 2026-09-07 12:57:56,881 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:57:56] iteration      31/    100
| consumed samples:      248 | elapsed time per iteration (ms): 765.8 | learning rate: 8.830222E-06 | global batch size:    8 | loss: 1.321291E
+00 | grad norm: 9.522 |
[Rank 0 | Local Rank 0] 2026-09-07 12:57:57,624 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:57:57] iteration      32/    100
| consumed samples:      256 | elapsed time per iteration (ms): 742.9 | learning rate: 8.715724E-06 | global batch size:    8 | loss: 1.435563E
+00 | grad norm: 9.331 |
[Rank 0 | Local Rank 0] 2026-09-07 12:57:58,386 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:57:58] iteration      33/    100
| consumed samples:      264 | elapsed time per iteration (ms): 761.4 | learning rate: 8.596699E-06 | global batch size:    8 | loss: 1.277295E
+00 | grad norm: 9.736 |
[Rank 0 | Local Rank 0] 2026-09-07 12:57:59,122 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:57:59] iteration      34/    100
| consumed samples:      272 | elapsed time per iteration (ms): 735.1 | learning rate: 8.473292E-06 | global batch size:    8 | loss: 1.292857E
+00 | grad norm: 10.319 |
[Rank 0 | Local Rank 0] 2026-09-07 12:57:59,921 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:57:59] iteration      35/    100
| consumed samples:      280 | elapsed time per iteration (ms): 797.9 | learning rate: 8.345653E-06 | global batch size:    8 | loss: 1.421362E
+00 | grad norm: 9.337 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:00,692 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:58:00] iteration      36/    100
| consumed samples:      288 | elapsed time per iteration (ms): 769.8 | learning rate: 8.213938E-06 | global batch size:    8 | loss: 1.397483E
+00 | grad norm: 9.974 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:01,462 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:58:01] iteration      37/    100
| consumed samples:      296 | elapsed time per iteration (ms): 769.7 | learning rate: 8.078307E-06 | global batch size:    8 | loss: 1.383438E
+00 | grad norm: 11.104 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:02,241 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:58:02] iteration      38/    100
| consumed samples:      304 | elapsed time per iteration (ms): 777.2 | learning rate: 7.938926E-06 | global batch size:    8 | loss: 1.474795E
+00 | grad norm: 10.309 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:03,010 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:58:03] iteration      39/    100
| consumed samples:      312 | elapsed time per iteration (ms): 768.2 | learning rate: 7.795965E-06 | global batch size:    8 | loss: 1.353377E
+00 | grad norm: 9.529 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:03,981 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:58:03] iteration      40/    100
| consumed samples:      320 | elapsed time per iteration (ms): 970.5 | learning rate: 7.649596E-06 | global batch size:    8 | loss: 1.512490E
+00 | grad norm: 8.928 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:04,724 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:58:04] iteration      41/    100
| consumed samples:      328 | elapsed time per iteration (ms): 742.4 | learning rate: 7.500000E-06 | global batch size:    8 | loss: 1.404383E
+00 | grad norm: 10.423 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:05,485 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:58:05] iteration      42/    100
| consumed samples:      336 | elapsed time per iteration (ms): 759.9 | learning rate: 7.347358E-06 | global batch size:    8 | loss: 1.305056E
+00 | grad norm: 9.322 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:06,161 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:58:06] iteration      43/    100
| consumed samples:      344 | elapsed time per iteration (ms): 675.0 | learning rate: 7.191856E-06 | global batch size:    8 | loss: 1.327726E
+00 | grad norm: 10.303 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:06,926 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:58:06] iteration      44/    100
| consumed samples:      352 | elapsed time per iteration (ms): 764.8 | learning rate: 7.033683E-06 | global batch size:    8 | loss: 1.354272E
+00 | grad norm: 10.308 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:07,664 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:58:07] iteration      45/    100
| consumed samples:      360 | elapsed time per iteration (ms): 736.6 | learning rate: 6.873033E-06 | global batch size:    8 | loss: 1.288647E
+00 | grad norm: 9.499 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:08,368 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:58:08] iteration      46/    100
| consumed samples:      368 | elapsed time per iteration (ms): 703.5 | learning rate: 6.710101E-06 | global batch size:    8 | loss: 1.351867E
+00 | grad norm: 10.619 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:09,117 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:58:09] iteration      47/    100
| consumed samples:      376 | elapsed time per iteration (ms): 748.2 | learning rate: 6.545085E-06 | global batch size:    8 | loss: 1.448191E
+00 | grad norm: 10.640 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:09,848 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:58:09] iteration      48/    100
| consumed samples:      384 | elapsed time per iteration (ms): 730.3 | learning rate: 6.378187E-06 | global batch size:    8 | loss: 1.486479E
+00 | grad norm: 9.621 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:10,577 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:58:10] iteration      49/    100
| consumed samples:      392 | elapsed time per iteration (ms): 728.1 | learning rate: 6.209609E-06 | global batch size:    8 | loss: 1.271842E
+00 | grad norm: 10.032 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:11,520 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:58:11] iteration      50/    100
| consumed samples:      400 | elapsed time per iteration (ms): 942.1 | learning rate: 6.039558E-06 | global batch size:    8 | loss: 1.276974E
+00 | grad norm: 9.070 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:12,268 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:58:12] iteration      51/    100
| consumed samples:      408 | elapsed time per iteration (ms): 747.2 | learning rate: 5.868241E-06 | global batch size:    8 | loss: 1.373032E
+00 | grad norm: 9.273 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:13,036 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:58:13] iteration      52/    100
| consumed samples:      416 | elapsed time per iteration (ms): 767.0 | learning rate: 5.695866E-06 | global batch size:    8 | loss: 1.558424E
+00 | grad norm: 9.389 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:13,767 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:58:13] iteration      53/    100
| consumed samples:      424 | elapsed time per iteration (ms): 730.1 | learning rate: 5.522642E-06 | global batch size:    8 | loss: 1.303517E
+00 | grad norm: 10.037 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:14,504 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:58:14] iteration      54/    100
| consumed samples:      432 | elapsed time per iteration (ms): 736.1 | learning rate: 5.348782E-06 | global batch size:    8 | loss: 1.282160E
+00 | grad norm: 9.527 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:15,226 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:58:15] iteration      55/    100
| consumed samples:      440 | elapsed time per iteration (ms): 721.1 | learning rate: 5.174497E-06 | global batch size:    8 | loss: 1.298149E
+00 | grad norm: 10.541 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:15,995 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:58:15] iteration      56/    100
| consumed samples:      448 | elapsed time per iteration (ms): 768.1 | learning rate: 5.000000E-06 | global batch size:    8 | loss: 1.428514E
+00 | grad norm: 9.234 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:16,798 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:58:16] iteration      57/    100
| consumed samples:      456 | elapsed time per iteration (ms): 802.3 | learning rate: 4.825503E-06 | global batch size:    8 | loss: 1.386111E
+00 | grad norm: 9.614 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:17,522 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:58:17] iteration      58/    100
| consumed samples:      464 | elapsed time per iteration (ms): 723.0 | learning rate: 4.651218E-06 | global batch size:    8 | loss: 1.364428E
+00 | grad norm: 8.742 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:18,252 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:58:18] iteration      59/    100
| consumed samples:      472 | elapsed time per iteration (ms): 729.3 | learning rate: 4.477358E-06 | global batch size:    8 | loss: 1.208515E
```

```
+00 | grad norm: 10.223 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:18,976 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:58:18] iteration      60/    100
| consumed samples:      480 | elapsed time per iteration (ms): 722.8 | learning rate: 4.304134E-06 | global batch size:      8 | loss: 1.268237E
+00 | grad norm: 10.472 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:19,706 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:58:19] iteration      61/    100
| consumed samples:      488 | elapsed time per iteration (ms): 729.4 | learning rate: 4.131759E-06 | global batch size:      8 | loss: 1.309227E
+00 | grad norm: 9.585 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:20,467 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:58:20] iteration      62/    100
| consumed samples:      496 | elapsed time per iteration (ms): 760.4 | learning rate: 3.960442E-06 | global batch size:      8 | loss: 1.196219E
+00 | grad norm: 10.985 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:21,218 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:58:21] iteration      63/    100
| consumed samples:      504 | elapsed time per iteration (ms): 749.3 | learning rate: 3.790391E-06 | global batch size:      8 | loss: 1.417329E
+00 | grad norm: 9.743 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:21,984 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:58:21] iteration      64/    100
| consumed samples:      512 | elapsed time per iteration (ms): 764.7 | learning rate: 3.621813E-06 | global batch size:      8 | loss: 1.352126E
+00 | grad norm: 10.170 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:22,700 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:58:22] iteration      65/    100
| consumed samples:      520 | elapsed time per iteration (ms): 715.6 | learning rate: 3.454915E-06 | global batch size:      8 | loss: 1.308774E
+00 | grad norm: 10.250 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:23,463 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:58:23] iteration      66/    100
| consumed samples:      528 | elapsed time per iteration (ms): 761.6 | learning rate: 3.289899E-06 | global batch size:      8 | loss: 1.410771E
+00 | grad norm: 9.517 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:24,235 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:58:24] iteration      67/    100
| consumed samples:      536 | elapsed time per iteration (ms): 771.2 | learning rate: 3.126967E-06 | global batch size:      8 | loss: 1.291760E
+00 | grad norm: 9.705 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:25,018 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:58:25] iteration      68/    100
| consumed samples:      544 | elapsed time per iteration (ms): 782.6 | learning rate: 2.966317E-06 | global batch size:      8 | loss: 1.534115E
+00 | grad norm: 10.443 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:25,838 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:58:25] iteration      69/    100
| consumed samples:      552 | elapsed time per iteration (ms): 818.7 | learning rate: 2.808144E-06 | global batch size:      8 | loss: 1.443359E
+00 | grad norm: 9.389 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:26,550 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:58:26] iteration      70/    100
| consumed samples:      560 | elapsed time per iteration (ms): 710.4 | learning rate: 2.652642E-06 | global batch size:      8 | loss: 1.331196E
+00 | grad norm: 10.317 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:27,686 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:58:27] iteration      71/    100
| consumed samples:      568 | elapsed time per iteration (ms): 1135.9 | learning rate: 2.500000E-06 | global batch size:      8 | loss: 1.286824
E+00 | grad norm: 10.824 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:28,425 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:58:28] iteration      72/    100
| consumed samples:      576 | elapsed time per iteration (ms): 737.1 | learning rate: 2.359404E-06 | global batch size:      8 | loss: 1.402632E
+00 | grad norm: 9.761 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:29,170 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:58:29] iteration      73/    100
| consumed samples:      584 | elapsed time per iteration (ms): 744.6 | learning rate: 2.204035E-06 | global batch size:      8 | loss: 1.378191E
+00 | grad norm: 9.595 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:29,936 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:58:29] iteration      74/    100
| consumed samples:      592 | elapsed time per iteration (ms): 764.7 | learning rate: 2.061074E-06 | global batch size:      8 | loss: 1.324130E
+00 | grad norm: 10.044 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:30,629 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:58:30] iteration      75/    100
| consumed samples:      600 | elapsed time per iteration (ms): 692.4 | learning rate: 1.921693E-06 | global batch size:      8 | loss: 1.362146E
+00 | grad norm: 9.509 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:31,408 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:58:31] iteration      76/    100
| consumed samples:      608 | elapsed time per iteration (ms): 778.0 | learning rate: 1.786062E-06 | global batch size:      8 | loss: 1.338994E
+00 | grad norm: 9.396 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:32,142 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:58:32] iteration      77/    100
| consumed samples:      616 | elapsed time per iteration (ms): 733.1 | learning rate: 1.654347E-06 | global batch size:      8 | loss: 1.415432E
+00 | grad norm: 9.746 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:32,845 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:58:32] iteration      78/    100
| consumed samples:      624 | elapsed time per iteration (ms): 702.2 | learning rate: 1.526708E-06 | global batch size:      8 | loss: 1.423982E
+00 | grad norm: 9.896 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:33,651 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:58:33] iteration      79/    100
| consumed samples:      632 | elapsed time per iteration (ms): 804.9 | learning rate: 1.403301E-06 | global batch size:      8 | loss: 1.468774E
+00 | grad norm: 9.475 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:34,417 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:58:34] iteration      80/    100
| consumed samples:      640 | elapsed time per iteration (ms): 764.8 | learning rate: 1.284276E-06 | global batch size:      8 | loss: 1.277077E
+00 | grad norm: 8.813 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:35,116 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:58:35] iteration      81/    100
| consumed samples:      648 | elapsed time per iteration (ms): 697.8 | learning rate: 1.169778E-06 | global batch size:      8 | loss: 1.309318E
+00 | grad norm: 10.262 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:35,907 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:58:35] iteration      82/    100
| consumed samples:      656 | elapsed time per iteration (ms): 790.3 | learning rate: 1.059946E-06 | global batch size:      8 | loss: 1.319990E
+00 | grad norm: 9.836 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:36,641 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:58:36] iteration      83/    100
| consumed samples:      664 | elapsed time per iteration (ms): 732.9 | learning rate: 9.549150E-07 | global batch size:      8 | loss: 1.266413E
+00 | grad norm: 9.414 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:37,404 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:58:37] iteration      84/    100
| consumed samples:      672 | elapsed time per iteration (ms): 762.2 | learning rate: 8.548121E-07 | global batch size:      8 | loss: 1.235728E
+00 | grad norm: 9.165 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:38,179 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:58:38] iteration      85/    100
| consumed samples:      680 | elapsed time per iteration (ms): 773.6 | learning rate: 7.597595E-07 | global batch size:      8 | loss: 1.374440E
+00 | grad norm: 11.422 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:38,889 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:58:38] iteration      86/    100
| consumed samples:      688 | elapsed time per iteration (ms): 709.1 | learning rate: 6.698730E-07 | global batch size:      8 | loss: 1.372502E
+00 | grad norm: 10.687 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:39,661 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:58:39] iteration      87/    100
| consumed samples:      696 | elapsed time per iteration (ms): 770.5 | learning rate: 5.852620E-07 | global batch size:      8 | loss: 1.368567E
+00 | grad norm: 8.825 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:40,425 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:58:40] iteration      88/    100
| consumed samples:      704 | elapsed time per iteration (ms): 762.7 | learning rate: 5.060298E-07 | global batch size:      8 | loss: 1.274261E
+00 | grad norm: 10.057 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:41,180 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:58:41] iteration      89/    100
| consumed samples:      712 | elapsed time per iteration (ms): 753.9 | learning rate: 4.322727E-07 | global batch size:      8 | loss: 1.133601E
+00 | grad norm: 9.861 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:41,937 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:58:41] iteration      90/    100
| consumed samples:      720 | elapsed time per iteration (ms): 756.9 | learning rate: 3.640807E-07 | global batch size:      8 | loss: 1.295144E
```

```
+00 | grad norm: 9.588 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:42,793 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:58:42] iteration      91/    100
| consumed samples:      728 | elapsed time per iteration (ms): 854.8 | learning rate: 3.015369E-07 | global batch size:    8 | loss: 1.352168E
+00 | grad norm: 8.143 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:43,447 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:58:43] iteration      92/    100
| consumed samples:      736 | elapsed time per iteration (ms): 653.6 | learning rate: 2.447174E-07 | global batch size:    8 | loss: 1.286163E
+00 | grad norm: 10.368 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:44,179 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:58:44] iteration      93/    100
| consumed samples:      744 | elapsed time per iteration (ms): 730.2 | learning rate: 1.936915E-07 | global batch size:    8 | loss: 1.205287E
+00 | grad norm: 10.175 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:45,009 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:58:45] iteration      94/    100
| consumed samples:      752 | elapsed time per iteration (ms): 829.6 | learning rate: 1.485214E-07 | global batch size:    8 | loss: 1.439860E
+00 | grad norm: 9.666 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:45,767 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:58:45] iteration      95/    100
| consumed samples:      760 | elapsed time per iteration (ms): 756.5 | learning rate: 1.092620E-07 | global batch size:    8 | loss: 1.322350E
+00 | grad norm: 9.837 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:46,626 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:58:46] iteration      96/    100
| consumed samples:      768 | elapsed time per iteration (ms): 858.2 | learning rate: 7.596123E-08 | global batch size:    8 | loss: 1.324017E
+00 | grad norm: 8.599 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:47,380 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:58:47] iteration      97/    100
| consumed samples:      776 | elapsed time per iteration (ms): 753.1 | learning rate: 4.865966E-08 | global batch size:    8 | loss: 1.343088E
+00 | grad norm: 10.059 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:48,044 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:58:48] iteration      98/    100
| consumed samples:      784 | elapsed time per iteration (ms): 663.1 | learning rate: 2.739052E-08 | global batch size:    8 | loss: 1.368604E
+00 | grad norm: 11.544 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:48,826 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:58:48] iteration      99/    100
| consumed samples:      792 | elapsed time per iteration (ms): 781.2 | learning rate: 1.217975E-08 | global batch size:    8 | loss: 1.357631E
+00 | grad norm: 9.909 |
[Rank 0 | Local Rank 0] 2026-09-07 12:58:49,849 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:58:49] iteration     100/    100
| consumed samples:      800 | elapsed time per iteration (ms): 1021.5 | learning rate: 3.045865E-09 | global batch size:    8 | loss: 1.440428E
+00 | grad norm: 9.364 |
/usr/local/lib/python3.10/dist-packages/torch/distributed/checkpoint/filesystem.py:592: UserWarning: Detected an existing checkpoint in save_path/iter_0000100/.metadata, overwriting since self.overwrite=True. Past version 2.5 of PyTorch, `overwrite` will default to False. Set this variable to True to maintain this functionality or False to raise when an existing checkpoint is found.
  warnings.warn(
[Rank 0 | Local Rank 0] 2026-09-07 12:58:58,881 INFO [mindspeed_mm.fsdp.checkpoint.dcp_checkpoint:48] => Saved checkpoint to ./save_path/iter_0000100
```

20260907_跨卡数值等价_coco100

文件: 20260907_跨卡数值等价_coco100.log | 行数: 685 | 大小: 90.4 KB | 跨卡数值等价性验证 (910B4 vs 910C)
采集说明: 完整日志 (从程序启动开始, 未经裁剪)

```
[INFO] Starting configuration loading...
[INFO] Loaded configuration file: examples/qwen3_5/qwen3_5_0_8B_final_config.yaml
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdps.data.data_utils.func_utils.model_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.model_name_or_path (type: <class 'str'>, value: /root/Qwen3.5-0.8B-hf)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdps.data.data_utils.func_utils.model_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.use_fast_tokenizer (type: <class 'bool'>, value: True)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdps.data.data_utils.func_utils.model_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.split_special_tokens (type: <class 'bool'>, value: False)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdps.data.data_utils.func_utils.model_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.image_max_pixels (type: <class 'int'>, value: 262144)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdps.data.data_utils.func_utils.model_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.image_min_pixels (type: <class 'int'>, value: 1024)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdps.data.data_utils.func_utils.model_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.video_max_pixels (type: <class 'int'>, value: 16384)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdps.data.data_utils.func_utils.model_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.video_min_pixels (type: <class 'int'>, value: 0)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdps.data.data_utils.func_utils.model_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.video_fps (type: <class 'float'>, value: 2.0)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdps.data.data_utils.func_utils.model_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.video_maxlen (type: <class 'int'>, value: 64)
[INFO] Automatically registered YAML field: model.gdn_implementation (type: <class 'str'>, value: triton)
[INFO] Automatically registered YAML field: model.causal_conv1d_implementation (type: <class 'str'>, value: triton)
[INFO] Automatically registered YAML field: model.skip_gdn_recompute (type: <class 'bool'>, value: False)
[INFO] Automatically registered YAML field: model.skip_flash_attn_recompute (type: <class 'bool'>, value: False)
[INFO] Automatically registered YAML field: features.recompute (type: <class 'bool'>, value: True)
=====
BaseArguments Configuration Manager Summary
=====

Configuration class: Arguments
Configuration file: examples/qwen3_5/qwen3_5_0_8B_final_config.yaml
Allow YAML extensions: True
Allow CLI override: True
Strict CLI validation: True
Allow registering YAML fields: True

Defined field count: 208
Dynamic field count: 14
Total field count: 222

===== Configuration Details =====
parallel:
  data_parallel_size: 1
  fully_shard_parallel_size: 1
  fsdp_plan:
    ignored_modules: []
    apply_modules: ['model.visual', 'model.visual.blocks.*', 'model.language_model.embed_tokens', 'model.language_model', 'model.language_model.layers.*', 'lm_head', 'mtp']
    param_dtype: bf16
    reduce_dtype: fp32
    output_dtype: None
    cast_forward_inputs: True
    reshard_after_forward: False
    num_to_forward_prefetch: 1
    num_to_backward_prefetch: 1
    pregather: False
    hook_modules: None
    cpu_offload: False
    tensor_parallel_size: 1
  tp_plan:
    colwise_parallel: []
    rowwise_parallel: []
    sequence_parallel: []
  ring_attention_size: 1
  ulysses_parallel_size: 1
  expert_parallel_size: 1
  expert_fully_shard_parallel_size: 1
  ep_plan:
    apply_modules: []
    use_npu_fused_ops: True
    dispatcher: alltoall
    apply_efsdps_modules: []
  recompute: False
  recompute_plan:
    apply_modules: []
```

```
    use_reentrant: False
    local_rank: 0
    global_rank: 0
    world_size: 1
model:
  model_id: qwen3_5
  model_name_or_path: /root/Qwen3.5-0.8B-hf
  trust_remote_code: True
  attn_implementation: flash_attention_2
  freeze: ['model.visual']
  mtp_num_layers: 0
  mtp_loss_scaling_factor: 0.1
  gdn_implementation: triton
  causal_conv1d_implementation: triton
  skip_gdn_recompute: False
  skip_flash_attn_recompute: False
data:
  dataset_param:
    dataset_type: huggingface
    basic_parameters:
      cache_dir: ./cache_dir/
      template: qwen3_vl_nothink
      chat_template: None
      enable_thinking: False
      dataset_dir: /root/coco_d1/extracted/dataset
      dataset: ['/root/coco_d1/extracted/dataset/annotations_slim.json']
      cutoff_len: 1024
      train_on_prompt: False
      mask_history: False
      streaming: False
      overwrite_cache: False
      preprocessing_batch_size: 1000
      preprocessing_num_workers: 2
      max_samples: None
      tool_format: None
      val_dataset: None
      val_max_samples: None
      val_rate: None
      packing: None
      neat_packing: False
      preprocess_on_fly: True
      async_preprocess: False
      async_preprocess_buffer_size: None
      stage: sft
    preprocess_parameters:
      model_name_or_path: /root/Qwen3.5-0.8B-hf
      use_fast_tokenizer: True
      split_special_tokens: False
      image_resolution: 512
      video_resolution: 128
      image_max_pixels: 262144
      image_min_pixels: 1024
      video_max_pixels: 16384
      video_min_pixels: 0
      video_fps: 2.0
      video_maxlen: 64
      image_do_pan_and_scan: False
      crop_to_patches: False
      use_audio_in_video: False
      audio_sampling_rate: 16000
      trust_remote_code: False
  attr:
    ranking: False
    system: None
    images: images
    videos: None
    audios: None
    prompt: None
    query: None
    response: None
    history: None
    messages: messages
    tools: None
    role_tag: role
    content_tag: content
    user_tag: user
    assistant_tag: assistant
    observation_tag: observation
    function_tag: function_call
    system_tag: system
    chosen: None
    rejected: None
    formatting: sharegpt
  val_dataset_param: None
  dataloader_param:
    dataloader_mode: sampler
    sampler_type: BaseRandomBatchSampler
    shuffle: False
    drop_last: True
    pin_memory: True
    collate_param:
      collator_id: None
      model_name: qwen3vl
```

```

        ignore_pad_token_for_loss: True
        pad_to_multiple_of: 8
        num_workers: 2
        enable_preload: False
training:
  profile:
    enable: False
    start_step: 1
    end_step: 2
    save_path: ./profiling
    record_shapes: True
    with_memory: True
    with_stack: True
    ranks: [0]
    _train_steps: -1
    local_rank: 0
    global_rank: 0
    world_size: 1
    profile_this_rank: False
lora:
  enable: False
  rank: 8
  alpha: 16
  target_modules: ['q_proj', 'k_proj', 'v_proj']
  dropout: 0.0
  init_lora_weights: True
  pretrained_lora_path: None
  lora_target_modules_support: None
  disable_peft_moe_conversion: True
lr: 1e-05
lr_min: 0.0
lr_start: 0.0
weight_decay: 0.0
no_decay_modules: []
no_decay_params: []
optimizer: adamw
mem_fraction_static: 0.8
matched_adamw_rms: 0.2
muon_momentum: 0.95
ns_steps: 5
adam_fused: True
adam_beta1: 0.9
adam_beta2: 0.999
adam_eps: 1e-08
clip_grad: 0.0
clip_grad_foreach: True
micro_batch_size: 2
gradient_accumulation_steps: 4
lr_warmup_ratio: 0.1
lr_decay_style: cosine
lr_decay_ratio: 1.0
enable_full_determinism: False
ckpt_manager: dcp
save_async: False
save_ckpt_dtype: None
load_checkpoint_path: None
save_steps: 0
seed: 42
max_steps: None
init_model_with_meta_device: True
train_iters: 100
train_epochs: None
load: /root/Qwen3.5-0.8B-dcp
load_format: auto
save_format: auto
load_strict: False
load_rank0_and_broadcast: False
no_load_optim: True
no_load_rng: True
save: ./save_path
no_save_optim: True
no_save_rng: True
log_interval: 1
save_interval: 10000
val_interval: 0
val_micro_batch_size: None
use_deter_comp: False
allow_hf32: None
plugin: ['mindspeed_mm/fsdp/models/qwen3_5', 'mindspeed_mm/fsdp/data/datasets/huggingface']
lr_scaling_plan: None
quantization_plan: QuantizeConfig(recipe_name=None, apply_modules=[], ignored_modules=[], quant_converters=[], extra_args={}, enable_fsdp_low_precision_all_gather=True, fsdp_low_precision_all_gather_mode='on-demand')
_train_steps: -1
local_rank: 0
global_rank: 0
world_size: 1
global_batch_size: 8
tools:
  profile:
    enable: False
    profile_type: static
    ranks: [0]
    static_param:

```

```

    level: level1
    with_stack: False
    with_memory: False
    record_shapes: False
    with_cpu: True
    save_path: ./profiling
    start_step: 10
    end_step: 11
    data_simplification: False
    aic_metrics_type: PipeUtilization
    analyse_flag: True
memory_profile:
    enable: False
    start_step: 1
    end_step: 2
    save_path: ./memory_snapshot
    dump_ranks: [0]
    stacks: all
    max_entries: None
    mem_info: False
features:
    recompute_plan:
        apply_modules: ['model.visual.blocks.*', 'model.language_model.layers.*']
        use_reentrant: False
    loss_cfg:
        loss_type: default
        router_aux_loss_coef: 0.0
        router_aux_loss_offload: False
    enable_chunk_loss: True
    enable_dynamic_chunk_loss: False
    chunkloss_plan:
        apply_module: lm_head
        chunk_size: 1024
        total_chunk_size: 4096
    enable_activation_offload: True
    activation_offload_plan:
        apply_modules: ['model.visual.blocks.*', 'model.language_model.layers.*']
    enable_chunk_mbs: False
    chunkmbs_plan:
        apply_modules: ['model.language_model.layers.*']
        chunk_mbs: 2
        batch_dim: 0
        chunk_arg_indexes: [0]
        chunk_kwarg_names: ['position_embeddings', 'position_ids', 'rope_deltas', 'attention_mask']
    enable_ep_balance: False
    ep_balance_plan:
        max_dup_experts_num: 2
    skip_moe_pad_tokens: False
    recompute: True
data.dataset_param.preprocess_parameters.model_name_or_path: /root/Qwen3.5-0.8B-hf
data.dataset_param.preprocess_parameters.use_fast_tokenizer: True
data.dataset_param.preprocess_parameters.split_special_tokens: False
data.dataset_param.preprocess_parameters.image_max_pixels: 262144
data.dataset_param.preprocess_parameters.image_min_pixels: 1024
data.dataset_param.preprocess_parameters.video_max_pixels: 16384
data.dataset_param.preprocess_parameters.video_min_pixels: 0
data.dataset_param.preprocess_parameters.video_fps: 2.0
data.dataset_param.preprocess_parameters.video_maxlen: 64

===== Dynamic Fields =====
data.dataset_param.preprocess_parameters.image_max_pixels
data.dataset_param.preprocess_parameters.image_min_pixels
data.dataset_param.preprocess_parameters.model_name_or_path
data.dataset_param.preprocess_parameters.split_special_tokens
data.dataset_param.preprocess_parameters.use_fast_tokenizer
data.dataset_param.preprocess_parameters.video_fps
data.dataset_param.preprocess_parameters.video_max_pixels
data.dataset_param.preprocess_parameters.video_maxlen
data.dataset_param.preprocess_parameters.video_min_pixels
features.recompute
model.causal_conv1d_implementation
model.gdn_implementation
model.skip_flash_attn_recompute
model.skip_gdn_recompute
=====
Successfully imported: mindspeed_mm.fsd.models.qwen3_5
Successfully imported: mindspeed_mm.fsd.data.datasets.huggingface
[Rank 0 | Local Rank 0] 2026-09-07 12:12:42,213 INFO [datasets:114] => Disabling Tensorflow because USE_TORCH is set
[Rank 0 | Local Rank 0] 2026-09-07 12:12:42,564 INFO [mindspeed_mm.fsd.distributed.parallel_state:87] => Parallel state initialized:
[dp] = False | Group size: 1 | device mesh:DeviceMesh('npu', [0], mesh_dim_names=('dp',))
[cp] = False | Group size: 1 | device mesh:DeviceMesh('npu', [0], mesh_dim_names=('cp',))
[ulysses] = False | Group size: 1 | device mesh:DeviceMesh('npu', [0], mesh_dim_names=('ulysses',))
[ring] = False | Group size: 1 | device mesh:DeviceMesh('npu', [0], mesh_dim_names=('ring',))
[tp] = False | Group size: 1 | device mesh:DeviceMesh('npu', [0], mesh_dim_names=('tp',))
[edp] = False | Group size: 1 | device mesh:DeviceMesh('npu', [0], mesh_dim_names=('edp',))
[efsdp] = False | Group size: 1 | device mesh:DeviceMesh('npu', [0], mesh_dim_names=('efsdp',))
[ep] = False | Group size: 1 | device mesh:DeviceMesh('npu', [0], mesh_dim_names=('ep',))
[fsdp] = True | Group size: 1 | device mesh:DeviceMesh('npu', [0], mesh_dim_names=('dp_shard_cp',))

[Rank 0 | Local Rank 0] 2026-09-07 12:12:42,564 INFO [mindspeed_mm.fsd.models.modelhub:48] => > Loading AutoConfig from /root/Qwen3.5-0.8B-hf...
[Rank 0 | Local Rank 0] 2026-09-07 12:12:42,567 INFO [mindspeed_mm.fsd.models.modelhub:48] => Building transformers model from configuration...
[Rank 0 | Local Rank 0] 2026-09-07 12:12:42,606 INFO [mindspeed_mm.fsd.models.qwen3_5.modeling_qwen3_5:48] => Qwen3_5 causal_conv1d use NPU triton

```

[illegible]

[illegible]

[illegible]

[illegible]

-dcp/release

[2026-09-07 12:12:53] [WARNING] [15876] profiler.py: Invalid parameter export_type: None, reset it to text.

[2026-09-07 12:12:53] [WARNING] [15876] profiler.py: Profiler won't be using warmup, this can skew profiler results

/root/MindSpeed-MM/mindspeed_mm/fsdp/features/memory/async_offload.py:25: UserWarning: TypedStorage is deprecated. It will be removed in the future and UntypedStorage will be the only storage class. This should only matter to you if you are using storages directly. To access UntypedStorage directly, use tensor.untyped_storage() instead of tensor.storage()

if tensor.storage().size() <= 0:

[rank0]:[W907 12:12:54.254017740 compiler_depend.ts:335] Warning: Cannot create tensor with interal format while allow_internel_format=False, tensor will be created with base format. (function operator())

[Rank 0 | Local Rank 0] 2026-09-07 12:13:00,237 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:13:00] iteration 1/ 100
| consumed samples: 8 | elapsed time per iteration (ms): 6717.6 | learning rate: 0.000000E+00 | global batch size: 8 | loss: 1.927377
E+00 | grad norm: 137.957 |

[Rank 0 | Local Rank 0] 2026-09-07 12:13:02,419 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:13:02] iteration 2/ 100
| consumed samples: 16 | elapsed time per iteration (ms): 2180.9 | learning rate: 1.000000E-06 | global batch size: 8 | loss: 1.786724
E+00 | grad norm: 114.677 |

[Rank 0] (after 2 iterations) memory (MB) | allocated: 9001.8911328125 | max allocated: 20300.5595703125 | reserved: 22028.0 | max reserved: 22028.0

[Rank 0 | Local Rank 0] 2026-09-07 12:13:04,052 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:13:04] iteration 3/ 100
| consumed samples: 24 | elapsed time per iteration (ms): 1631.3 | learning rate: 2.000000E-06 | global batch size: 8 | loss: 1.788325
E+00 | grad norm: 91.525 |

[Rank 0 | Local Rank 0] 2026-09-07 12:13:05,656 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:13:05] iteration 4/ 100
| consumed samples: 32 | elapsed time per iteration (ms): 1603.5 | learning rate: 3.000000E-06 | global batch size: 8 | loss: 1.476821
E+00 | grad norm: 61.991 |

[Rank 0 | Local Rank 0] 2026-09-07 12:13:07,173 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:13:07] iteration 5/ 100
| consumed samples: 40 | elapsed time per iteration (ms): 1516.1 | learning rate: 4.000000E-06 | global batch size: 8 | loss: 1.461596
E+00 | grad norm: 26.809 |

[Rank 0 | Local Rank 0] 2026-09-07 12:13:08,735 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:13:08] iteration 6/ 100
| consumed samples: 48 | elapsed time per iteration (ms): 1561.5 | learning rate: 5.000000E-06 | global batch size: 8 | loss: 1.404129
E+00 | grad norm: 17.341 |

[Rank 0 | Local Rank 0] 2026-09-07 12:13:10,279 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:13:10] iteration 7/ 100
| consumed samples: 56 | elapsed time per iteration (ms): 1543.1 | learning rate: 6.000000E-06 | global batch size: 8 | loss: 1.330969
E+00 | grad norm: 15.968 |

[Rank 0 | Local Rank 0] 2026-09-07 12:13:11,941 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:13:11] iteration 8/ 100
| consumed samples: 64 | elapsed time per iteration (ms): 1661.3 | learning rate: 7.000000E-06 | global batch size: 8 | loss: 1.455358
E+00 | grad norm: 11.901 |

[Rank 0 | Local Rank 0] 2026-09-07 12:13:13,649 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:13:13] iteration 9/ 100
| consumed samples: 72 | elapsed time per iteration (ms): 1706.4 | learning rate: 8.000000E-06 | global batch size: 8 | loss: 1.286244
E+00 | grad norm: 12.865 |

[Rank 0 | Local Rank 0] 2026-09-07 12:13:15,528 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:13:15] iteration 10/ 100
| consumed samples: 80 | elapsed time per iteration (ms): 1878.7 | learning rate: 9.000000E-06 | global batch size: 8 | loss: 1.419681
E+00 | grad norm: 11.338 |

[Rank 0 | Local Rank 0] 2026-09-07 12:13:19,623 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:13:19] iteration 11/ 100
| consumed samples: 88 | elapsed time per iteration (ms): 4094.4 | learning rate: 1.000000E-05 | global batch size: 8 | loss: 1.349043
E+00 | grad norm: 13.225 |

[Rank 0 | Local Rank 0] 2026-09-07 12:13:39,893 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:13:39] iteration 12/ 100
| consumed samples: 96 | elapsed time per iteration (ms): 20268.4 | learning rate: 9.996954E-06 | global batch size: 8 | loss: 1.43592
4E+00 | grad norm: 11.133 |

[Rank 0 | Local Rank 0] 2026-09-07 12:13:41,471 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:13:41] iteration 13/ 100
| consumed samples: 104 | elapsed time per iteration (ms): 1578.0 | learning rate: 9.987820E-06 | global batch size: 8 | loss: 1.441790
E+00 | grad norm: 10.067 |

[Rank 0 | Local Rank 0] 2026-09-07 12:13:43,178 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:13:43] iteration 14/ 100
| consumed samples: 112 | elapsed time per iteration (ms): 1705.2 | learning rate: 9.972609E-06 | global batch size: 8 | loss: 1.418822
E+00 | grad norm: 11.087 |

[Rank 0 | Local Rank 0] 2026-09-07 12:14:03,836 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:14:03] iteration 15/ 100
| consumed samples: 120 | elapsed time per iteration (ms): 20655.3 | learning rate: 9.951340E-06 | global batch size: 8 | loss: 1.41059
1E+00 | grad norm: 9.928 |

[Rank 0 | Local Rank 0] 2026-09-07 12:14:05,537 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:14:05] iteration 16/ 100
| consumed samples: 128 | elapsed time per iteration (ms): 1700.3 | learning rate: 9.924039E-06 | global batch size: 8 | loss: 1.459432
E+00 | grad norm: 11.813 |

[Rank 0 | Local Rank 0] 2026-09-07 12:14:07,214 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:14:07] iteration 17/ 100
| consumed samples: 136 | elapsed time per iteration (ms): 1675.8 | learning rate: 9.890738E-06 | global batch size: 8 | loss: 1.361190
E+00 | grad norm: 11.016 |

[Rank 0 | Local Rank 0] 2026-09-07 12:14:09,271 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:14:09] iteration 18/ 100
| consumed samples: 144 | elapsed time per iteration (ms): 2055.7 | learning rate: 9.851479E-06 | global batch size: 8 | loss: 1.280451
E+00 | grad norm: 11.839 |

[Rank 0 | Local Rank 0] 2026-09-07 12:14:10,986 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:14:10] iteration 19/ 100
| consumed samples: 152 | elapsed time per iteration (ms): 1713.4 | learning rate: 9.806308E-06 | global batch size: 8 | loss: 1.453921
E+00 | grad norm: 10.118 |

[Rank 0 | Local Rank 0] 2026-09-07 12:14:12,656 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:14:12] iteration 20/ 100
| consumed samples: 160 | elapsed time per iteration (ms): 1668.9 | learning rate: 9.755283E-06 | global batch size: 8 | loss: 1.254092
E+00 | grad norm: 11.733 |

[Rank 0 | Local Rank 0] 2026-09-07 12:14:14,327 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:14:14] iteration 21/ 100
| consumed samples: 168 | elapsed time per iteration (ms): 1670.6 | learning rate: 9.698463E-06 | global batch size: 8 | loss: 1.202713
E+00 | grad norm: 10.784 |

[Rank 0 | Local Rank 0] 2026-09-07 12:14:16,006 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:14:16] iteration 22/ 100
| consumed samples: 176 | elapsed time per iteration (ms): 1676.9 | learning rate: 9.635919E-06 | global batch size: 8 | loss: 1.435384
E+00 | grad norm: 10.204 |

[Rank 0 | Local Rank 0] 2026-09-07 12:14:17,698 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:14:17] iteration 23/ 100
| consumed samples: 184 | elapsed time per iteration (ms): 1691.3 | learning rate: 9.567727E-06 | global batch size: 8 | loss: 1.433935
E+00 | grad norm: 12.778 |

[Rank 0 | Local Rank 0] 2026-09-07 12:14:19,412 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:14:19] iteration 24/ 100
| consumed samples: 192 | elapsed time per iteration (ms): 1712.9 | learning rate: 9.493970E-06 | global batch size: 8 | loss: 1.491424
E+00 | grad norm: 9.886 |

[Rank 0 | Local Rank 0] 2026-09-07 12:14:21,095 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:14:21] iteration 25/ 100
| consumed samples: 200 | elapsed time per iteration (ms): 1681.5 | learning rate: 9.414738E-06 | global batch size: 8 | loss: 1.421190
E+00 | grad norm: 10.971 |

[Rank 0 | Local Rank 0] 2026-09-07 12:14:23,148 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:14:23] iteration 26/ 100
| consumed samples: 208 | elapsed time per iteration (ms): 2052.0 | learning rate: 9.330127E-06 | global batch size: 8 | loss: 1.297737
E+00 | grad norm: 12.047 |

[Rank 0 | Local Rank 0] 2026-09-07 12:14:24,813 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:14:24] iteration 27/ 100
| consumed samples: 216 | elapsed time per iteration (ms): 1664.8 | learning rate: 9.240240E-06 | global batch size: 8 | loss: 1.455574
E+00 | grad norm: 11.042 |

[Rank 0 | Local Rank 0] 2026-09-07 12:14:26,473 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:14:26] iteration 28/ 100

| consumed samples: 224 | elapsed time per iteration (ms): 1658.5 | learning rate: 9.145188E-06 | global batch size: 8 | loss: 1.317277
E+00 | grad norm: 10.620 |
[Rank 0 | Local Rank 0] 2026-09-07 12:14:28,152 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:14:28] iteration 29/ 100
| consumed samples: 232 | elapsed time per iteration (ms): 1678.4 | learning rate: 9.045085E-06 | global batch size: 8 | loss: 1.393960
E+00 | grad norm: 10.124 |
[Rank 0 | Local Rank 0] 2026-09-07 12:14:29,833 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:14:29] iteration 30/ 100
| consumed samples: 240 | elapsed time per iteration (ms): 1679.6 | learning rate: 8.940054E-06 | global batch size: 8 | loss: 1.329519
E+00 | grad norm: 10.060 |
[Rank 0 | Local Rank 0] 2026-09-07 12:14:34,056 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:14:34] iteration 31/ 100
| consumed samples: 248 | elapsed time per iteration (ms): 4221.4 | learning rate: 8.830222E-06 | global batch size: 8 | loss: 1.329337
E+00 | grad norm: 9.477 |
[Rank 0 | Local Rank 0] 2026-09-07 12:14:35,739 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:14:35] iteration 32/ 100
| consumed samples: 256 | elapsed time per iteration (ms): 1682.9 | learning rate: 8.715724E-06 | global batch size: 8 | loss: 1.456220
E+00 | grad norm: 9.604 |
[Rank 0 | Local Rank 0] 2026-09-07 12:15:04,402 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:15:04] iteration 33/ 100
| consumed samples: 264 | elapsed time per iteration (ms): 28661.4 | learning rate: 8.596699E-06 | global batch size: 8 | loss: 1.22639
5E+00 | grad norm: 9.960 |
[Rank 0 | Local Rank 0] 2026-09-07 12:15:06,443 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:15:06] iteration 34/ 100
| consumed samples: 272 | elapsed time per iteration (ms): 2040.1 | learning rate: 8.473292E-06 | global batch size: 8 | loss: 1.290874
E+00 | grad norm: 10.651 |
[Rank 0 | Local Rank 0] 2026-09-07 12:15:10,819 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:15:10] iteration 35/ 100
| consumed samples: 280 | elapsed time per iteration (ms): 4374.8 | learning rate: 8.345653E-06 | global batch size: 8 | loss: 1.417334
E+00 | grad norm: 9.156 |
[Rank 0 | Local Rank 0] 2026-09-07 12:15:12,542 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:15:12] iteration 36/ 100
| consumed samples: 288 | elapsed time per iteration (ms): 1722.7 | learning rate: 8.213938E-06 | global batch size: 8 | loss: 1.382835
E+00 | grad norm: 9.872 |
[Rank 0 | Local Rank 0] 2026-09-07 12:15:14,241 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:15:14] iteration 37/ 100
| consumed samples: 296 | elapsed time per iteration (ms): 1697.5 | learning rate: 8.078307E-06 | global batch size: 8 | loss: 1.338235
E+00 | grad norm: 11.182 |
[Rank 0 | Local Rank 0] 2026-09-07 12:15:18,429 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:15:18] iteration 38/ 100
| consumed samples: 304 | elapsed time per iteration (ms): 4187.0 | learning rate: 7.938926E-06 | global batch size: 8 | loss: 1.436761
E+00 | grad norm: 11.801 |
[Rank 0 | Local Rank 0] 2026-09-07 12:15:20,158 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:15:20] iteration 39/ 100
| consumed samples: 312 | elapsed time per iteration (ms): 1728.4 | learning rate: 7.795965E-06 | global batch size: 8 | loss: 1.341184
E+00 | grad norm: 9.663 |
[Rank 0 | Local Rank 0] 2026-09-07 12:15:21,904 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:15:21] iteration 40/ 100
| consumed samples: 320 | elapsed time per iteration (ms): 1744.9 | learning rate: 7.649596E-06 | global batch size: 8 | loss: 1.515426
E+00 | grad norm: 9.012 |
[Rank 0 | Local Rank 0] 2026-09-07 12:15:23,658 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:15:23] iteration 41/ 100
| consumed samples: 328 | elapsed time per iteration (ms): 1752.5 | learning rate: 7.500000E-06 | global batch size: 8 | loss: 1.391082
E+00 | grad norm: 10.875 |
[Rank 0 | Local Rank 0] 2026-09-07 12:15:25,617 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:15:25] iteration 42/ 100
| consumed samples: 336 | elapsed time per iteration (ms): 1958.5 | learning rate: 7.347358E-06 | global batch size: 8 | loss: 1.242836
E+00 | grad norm: 9.676 |
[Rank 0 | Local Rank 0] 2026-09-07 12:15:27,323 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:15:27] iteration 43/ 100
| consumed samples: 344 | elapsed time per iteration (ms): 1705.1 | learning rate: 7.191856E-06 | global batch size: 8 | loss: 1.328679
E+00 | grad norm: 10.423 |
[Rank 0 | Local Rank 0] 2026-09-07 12:15:29,019 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:15:29] iteration 44/ 100
| consumed samples: 352 | elapsed time per iteration (ms): 1694.7 | learning rate: 7.033683E-06 | global batch size: 8 | loss: 1.332474
E+00 | grad norm: 10.610 |
[Rank 0 | Local Rank 0] 2026-09-07 12:15:30,731 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:15:30] iteration 45/ 100
| consumed samples: 360 | elapsed time per iteration (ms): 1711.1 | learning rate: 6.873033E-06 | global batch size: 8 | loss: 1.251700
E+00 | grad norm: 10.244 |
[Rank 0 | Local Rank 0] 2026-09-07 12:15:32,433 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:15:32] iteration 46/ 100
| consumed samples: 368 | elapsed time per iteration (ms): 1701.2 | learning rate: 6.710101E-06 | global batch size: 8 | loss: 1.353570
E+00 | grad norm: 10.704 |
[Rank 0 | Local Rank 0] 2026-09-07 12:15:34,188 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:15:34] iteration 47/ 100
| consumed samples: 376 | elapsed time per iteration (ms): 1754.2 | learning rate: 6.545085E-06 | global batch size: 8 | loss: 1.421278
E+00 | grad norm: 11.102 |
[Rank 0 | Local Rank 0] 2026-09-07 12:15:35,963 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:15:35] iteration 48/ 100
| consumed samples: 384 | elapsed time per iteration (ms): 1773.5 | learning rate: 6.378187E-06 | global batch size: 8 | loss: 1.491345
E+00 | grad norm: 9.716 |
[Rank 0 | Local Rank 0] 2026-09-07 12:15:37,635 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:15:37] iteration 49/ 100
| consumed samples: 392 | elapsed time per iteration (ms): 1671.6 | learning rate: 6.209609E-06 | global batch size: 8 | loss: 1.220450
E+00 | grad norm: 10.492 |
[Rank 0 | Local Rank 0] 2026-09-07 12:15:58,219 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:15:58] iteration 50/ 100
| consumed samples: 400 | elapsed time per iteration (ms): 20583.4 | learning rate: 6.039558E-06 | global batch size: 8 | loss: 1.20825
5E+00 | grad norm: 9.280 |
[Rank 0 | Local Rank 0] 2026-09-07 12:15:59,951 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:15:59] iteration 51/ 100
| consumed samples: 408 | elapsed time per iteration (ms): 1730.7 | learning rate: 5.868241E-06 | global batch size: 8 | loss: 1.372874
E+00 | grad norm: 9.552 |
[Rank 0 | Local Rank 0] 2026-09-07 12:16:01,671 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:16:01] iteration 52/ 100
| consumed samples: 416 | elapsed time per iteration (ms): 1718.9 | learning rate: 5.695866E-06 | global batch size: 8 | loss: 1.537614
E+00 | grad norm: 9.645 |
[Rank 0 | Local Rank 0] 2026-09-07 12:16:03,344 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:16:03] iteration 53/ 100
| consumed samples: 424 | elapsed time per iteration (ms): 1671.3 | learning rate: 5.522642E-06 | global batch size: 8 | loss: 1.302565
E+00 | grad norm: 10.097 |
[Rank 0 | Local Rank 0] 2026-09-07 12:16:05,022 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:16:05] iteration 54/ 100
| consumed samples: 432 | elapsed time per iteration (ms): 1677.1 | learning rate: 5.348782E-06 | global batch size: 8 | loss: 1.303162
E+00 | grad norm: 10.237 |
[Rank 0 | Local Rank 0] 2026-09-07 12:16:06,716 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:16:06] iteration 55/ 100
| consumed samples: 440 | elapsed time per iteration (ms): 1692.8 | learning rate: 5.174497E-06 | global batch size: 8 | loss: 1.233353
E+00 | grad norm: 11.005 |
[Rank 0 | Local Rank 0] 2026-09-07 12:16:08,449 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:16:08] iteration 56/ 100
| consumed samples: 448 | elapsed time per iteration (ms): 1731.9 | learning rate: 5.000000E-06 | global batch size: 8 | loss: 1.431562
E+00 | grad norm: 9.322 |
[Rank 0 | Local Rank 0] 2026-09-07 12:16:10,155 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:16:10] iteration 57/ 100
| consumed samples: 456 | elapsed time per iteration (ms): 1704.5 | learning rate: 4.825503E-06 | global batch size: 8 | loss: 1.388738
E+00 | grad norm: 9.530 |
[Rank 0 | Local Rank 0] 2026-09-07 12:16:11,990 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:16:11] iteration 58/ 100
| consumed samples: 464 | elapsed time per iteration (ms): 1834.1 | learning rate: 4.651218E-06 | global batch size: 8 | loss: 1.369282
E+00 | grad norm: 8.703 |
[Rank 0 | Local Rank 0] 2026-09-07 12:16:13,605 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:16:13] iteration 59/ 100

| consumed samples: 472 | elapsed time per iteration (ms): 1614.0 | learning rate: 4.477358E-06 | global batch size: 8 | loss: 1.178877
E+00 | grad norm: 9.918 |
[Rank 0 | Local Rank 0] 2026-09-07 12:16:15,397 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:16:15] iteration 60/ 100
| consumed samples: 480 | elapsed time per iteration (ms): 1790.9 | learning rate: 4.304134E-06 | global batch size: 8 | loss: 1.260648
E+00 | grad norm: 10.947 |
[Rank 0 | Local Rank 0] 2026-09-07 12:16:17,038 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:16:17] iteration 61/ 100
| consumed samples: 488 | elapsed time per iteration (ms): 1639.6 | learning rate: 4.131759E-06 | global batch size: 8 | loss: 1.262179
E+00 | grad norm: 9.495 |
[Rank 0 | Local Rank 0] 2026-09-07 12:16:18,733 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:16:18] iteration 62/ 100
| consumed samples: 496 | elapsed time per iteration (ms): 1694.6 | learning rate: 3.960442E-06 | global batch size: 8 | loss: 1.183804
E+00 | grad norm: 11.220 |
[Rank 0 | Local Rank 0] 2026-09-07 12:16:20,279 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:16:20] iteration 63/ 100
| consumed samples: 504 | elapsed time per iteration (ms): 1545.4 | learning rate: 3.790391E-06 | global batch size: 8 | loss: 1.438879
E+00 | grad norm: 10.484 |
[Rank 0 | Local Rank 0] 2026-09-07 12:16:22,027 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:16:22] iteration 64/ 100
| consumed samples: 512 | elapsed time per iteration (ms): 1747.0 | learning rate: 3.621813E-06 | global batch size: 8 | loss: 1.324483
E+00 | grad norm: 10.251 |
[Rank 0 | Local Rank 0] 2026-09-07 12:16:23,781 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:16:23] iteration 65/ 100
| consumed samples: 520 | elapsed time per iteration (ms): 1752.7 | learning rate: 3.454915E-06 | global batch size: 8 | loss: 1.293153
E+00 | grad norm: 10.395 |
[Rank 0 | Local Rank 0] 2026-09-07 12:16:25,635 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:16:25] iteration 66/ 100
| consumed samples: 528 | elapsed time per iteration (ms): 1853.1 | learning rate: 3.289899E-06 | global batch size: 8 | loss: 1.348541
E+00 | grad norm: 9.897 |
[Rank 0 | Local Rank 0] 2026-09-07 12:16:27,369 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:16:27] iteration 67/ 100
| consumed samples: 536 | elapsed time per iteration (ms): 1732.8 | learning rate: 3.126967E-06 | global batch size: 8 | loss: 1.227077
E+00 | grad norm: 9.963 |
[Rank 0 | Local Rank 0] 2026-09-07 12:16:29,140 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:16:29] iteration 68/ 100
| consumed samples: 544 | elapsed time per iteration (ms): 1770.1 | learning rate: 2.966317E-06 | global batch size: 8 | loss: 1.504360
E+00 | grad norm: 11.506 |
[Rank 0 | Local Rank 0] 2026-09-07 12:16:30,904 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:16:30] iteration 69/ 100
| consumed samples: 552 | elapsed time per iteration (ms): 1763.2 | learning rate: 2.808144E-06 | global batch size: 8 | loss: 1.446333
E+00 | grad norm: 10.177 |
[Rank 0 | Local Rank 0] 2026-09-07 12:16:32,563 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:16:32] iteration 70/ 100
| consumed samples: 560 | elapsed time per iteration (ms): 1658.3 | learning rate: 2.652642E-06 | global batch size: 8 | loss: 1.277611
E+00 | grad norm: 10.548 |
[Rank 0 | Local Rank 0] 2026-09-07 12:16:34,210 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:16:34] iteration 71/ 100
| consumed samples: 568 | elapsed time per iteration (ms): 1646.1 | learning rate: 2.500000E-06 | global batch size: 8 | loss: 1.294862
E+00 | grad norm: 10.787 |
[Rank 0 | Local Rank 0] 2026-09-07 12:16:35,761 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:16:35] iteration 72/ 100
| consumed samples: 576 | elapsed time per iteration (ms): 1550.2 | learning rate: 2.350404E-06 | global batch size: 8 | loss: 1.402811
E+00 | grad norm: 10.187 |
[Rank 0 | Local Rank 0] 2026-09-07 12:16:37,466 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:16:37] iteration 73/ 100
| consumed samples: 584 | elapsed time per iteration (ms): 1704.0 | learning rate: 2.204035E-06 | global batch size: 8 | loss: 1.361734
E+00 | grad norm: 9.674 |
[Rank 0 | Local Rank 0] 2026-09-07 12:16:39,218 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:16:39] iteration 74/ 100
| consumed samples: 592 | elapsed time per iteration (ms): 1750.3 | learning rate: 2.061074E-06 | global batch size: 8 | loss: 1.287258
E+00 | grad norm: 10.670 |
[Rank 0 | Local Rank 0] 2026-09-07 12:16:41,137 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:16:41] iteration 75/ 100
| consumed samples: 600 | elapsed time per iteration (ms): 1918.3 | learning rate: 1.921693E-06 | global batch size: 8 | loss: 1.366915
E+00 | grad norm: 10.931 |
[Rank 0 | Local Rank 0] 2026-09-07 12:16:42,758 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:16:42] iteration 76/ 100
| consumed samples: 608 | elapsed time per iteration (ms): 1619.9 | learning rate: 1.786062E-06 | global batch size: 8 | loss: 1.346393
E+00 | grad norm: 9.323 |
[Rank 0 | Local Rank 0] 2026-09-07 12:16:44,366 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:16:44] iteration 77/ 100
| consumed samples: 616 | elapsed time per iteration (ms): 1606.5 | learning rate: 1.654347E-06 | global batch size: 8 | loss: 1.403700
E+00 | grad norm: 9.937 |
[Rank 0 | Local Rank 0] 2026-09-07 12:16:46,005 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:16:46] iteration 78/ 100
| consumed samples: 624 | elapsed time per iteration (ms): 1638.7 | learning rate: 1.526708E-06 | global batch size: 8 | loss: 1.396697
E+00 | grad norm: 9.960 |
[Rank 0 | Local Rank 0] 2026-09-07 12:16:47,730 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:16:47] iteration 79/ 100
| consumed samples: 632 | elapsed time per iteration (ms): 1722.9 | learning rate: 1.403301E-06 | global batch size: 8 | loss: 1.428493
E+00 | grad norm: 9.751 |
[Rank 0 | Local Rank 0] 2026-09-07 12:16:49,381 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:16:49] iteration 80/ 100
| consumed samples: 640 | elapsed time per iteration (ms): 1649.8 | learning rate: 1.284276E-06 | global batch size: 8 | loss: 1.248427
E+00 | grad norm: 8.869 |
[Rank 0 | Local Rank 0] 2026-09-07 12:16:51,030 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:16:51] iteration 81/ 100
| consumed samples: 648 | elapsed time per iteration (ms): 1648.2 | learning rate: 1.169778E-06 | global batch size: 8 | loss: 1.286088
E+00 | grad norm: 10.884 |
[Rank 0 | Local Rank 0] 2026-09-07 12:16:52,684 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:16:52] iteration 82/ 100
| consumed samples: 656 | elapsed time per iteration (ms): 1652.4 | learning rate: 1.059946E-06 | global batch size: 8 | loss: 1.222421
E+00 | grad norm: 10.358 |
[Rank 0 | Local Rank 0] 2026-09-07 12:16:54,324 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:16:54] iteration 83/ 100
| consumed samples: 664 | elapsed time per iteration (ms): 1639.4 | learning rate: 9.549150E-07 | global batch size: 8 | loss: 1.258622
E+00 | grad norm: 9.470 |
[Rank 0 | Local Rank 0] 2026-09-07 12:16:56,233 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:16:56] iteration 84/ 100
| consumed samples: 672 | elapsed time per iteration (ms): 1907.9 | learning rate: 8.548121E-07 | global batch size: 8 | loss: 1.183182
E+00 | grad norm: 9.400 |
[Rank 0 | Local Rank 0] 2026-09-07 12:16:57,860 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:16:57] iteration 85/ 100
| consumed samples: 680 | elapsed time per iteration (ms): 1625.4 | learning rate: 7.597595E-07 | global batch size: 8 | loss: 1.212902
E+00 | grad norm: 11.578 |
[Rank 0 | Local Rank 0] 2026-09-07 12:16:59,509 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:16:59] iteration 86/ 100
| consumed samples: 688 | elapsed time per iteration (ms): 1648.6 | learning rate: 6.698730E-07 | global batch size: 8 | loss: 1.353531
E+00 | grad norm: 11.083 |
[Rank 0 | Local Rank 0] 2026-09-07 12:17:01,164 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:17:01] iteration 87/ 100
| consumed samples: 696 | elapsed time per iteration (ms): 1653.2 | learning rate: 5.852620E-07 | global batch size: 8 | loss: 1.317314
E+00 | grad norm: 9.117 |
[Rank 0 | Local Rank 0] 2026-09-07 12:17:02,809 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:17:02] iteration 88/ 100
| consumed samples: 704 | elapsed time per iteration (ms): 1644.5 | learning rate: 5.060298E-07 | global batch size: 8 | loss: 1.177502
E+00 | grad norm: 10.028 |
[Rank 0 | Local Rank 0] 2026-09-07 12:17:04,468 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:17:04] iteration 89/ 100
| consumed samples: 712 | elapsed time per iteration (ms): 1657.1 | learning rate: 4.322727E-07 | global batch size: 8 | loss: 1.067771
E+00 | grad norm: 10.123 |
[Rank 0 | Local Rank 0] 2026-09-07 12:17:06,143 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-07 12:17:06] iteration 90/ 100

```
| consumed samples:          720 | elapsed time per iteration (ms): 1673.9 | learning rate: 3.640807E-07 | global batch size:      8 | loss: 1.241448
E+00 | grad norm: 9.870 |
[Rank 0 | Local Rank 0] 2026-09-07 12:17:07,827 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:17:07] iteration      91/    100
| consumed samples:          728 | elapsed time per iteration (ms): 1682.8 | learning rate: 3.015369E-07 | global batch size:      8 | loss: 1.354908
E+00 | grad norm: 8.372 |
[Rank 0 | Local Rank 0] 2026-09-07 12:17:09,746 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:17:09] iteration      92/    100
| consumed samples:          736 | elapsed time per iteration (ms): 1918.0 | learning rate: 2.447174E-07 | global batch size:      8 | loss: 1.286086
E+00 | grad norm: 10.419 |
[Rank 0 | Local Rank 0] 2026-09-07 12:17:11,381 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:17:11] iteration      93/    100
| consumed samples:          744 | elapsed time per iteration (ms): 1633.1 | learning rate: 1.936915E-07 | global batch size:      8 | loss: 1.163156
E+00 | grad norm: 10.211 |
[Rank 0 | Local Rank 0] 2026-09-07 12:17:12,995 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:17:12] iteration      94/    100
| consumed samples:          752 | elapsed time per iteration (ms): 1613.0 | learning rate: 1.485214E-07 | global batch size:      8 | loss: 1.373765
E+00 | grad norm: 9.903 |
[Rank 0 | Local Rank 0] 2026-09-07 12:17:14,588 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:17:14] iteration      95/    100
| consumed samples:          760 | elapsed time per iteration (ms): 1592.6 | learning rate: 1.092620E-07 | global batch size:      8 | loss: 1.310215
E+00 | grad norm: 10.078 |
[Rank 0 | Local Rank 0] 2026-09-07 12:17:18,818 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:17:18] iteration      96/    100
| consumed samples:          768 | elapsed time per iteration (ms): 4228.7 | learning rate: 7.596123E-08 | global batch size:      8 | loss: 1.285187
E+00 | grad norm: 9.445 |
[Rank 0 | Local Rank 0] 2026-09-07 12:17:20,386 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:17:20] iteration      97/    100
| consumed samples:          776 | elapsed time per iteration (ms): 1567.1 | learning rate: 4.865966E-08 | global batch size:      8 | loss: 1.362708
E+00 | grad norm: 10.481 |
[Rank 0 | Local Rank 0] 2026-09-07 12:17:22,000 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:17:21] iteration      98/    100
| consumed samples:          784 | elapsed time per iteration (ms): 1612.4 | learning rate: 2.739052E-08 | global batch size:      8 | loss: 1.360065
E+00 | grad norm: 11.857 |
[Rank 0 | Local Rank 0] 2026-09-07 12:17:23,640 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:17:23] iteration      99/    100
| consumed samples:          792 | elapsed time per iteration (ms): 1639.0 | learning rate: 1.217975E-08 | global batch size:      8 | loss: 1.279637
E+00 | grad norm: 9.897 |
[Rank 0 | Local Rank 0] 2026-09-07 12:17:25,511 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-07 12:17:25] iteration     100/    100
| consumed samples:          800 | elapsed time per iteration (ms): 1870.4 | learning rate: 3.045865E-09 | global batch size:      8 | loss: 1.423572
E+00 | grad norm: 9.439 |
/usr/local/lib/python3.10/dist-packages/torch_npu/distributed/distributed_c10d.py:110: UserWarning: HCCL doesn't support gather at the moment. Implemented with allgather instead.
  warnings.warn("HCCL doesn't support gather at the moment. Implemented with allgather instead.")
[Rank 0 | Local Rank 0] 2026-09-07 12:17:30,208 INFO [mindspeed_mm.fsdp.checkpoint.dcp_checkpointner:48] => Saved checkpoint to ./save_path/iter_0000
100
```

20260910_0949_优化轮E1_关闭显存节省开关

文件: 20260910_0949_优化轮E1_关闭显存节省开关.log | 行数: 632 | 大小: 80.1 KB | 性能优化轮 E1 (关闭显存节省开关)
采集说明: 完整日志 (从程序启动开始, 未经裁剪)

```
W0910 09:49:12.584000 4184 torch/distributed/run.py:766] *****
W0910 09:49:12.584000 4184 torch/distributed/run.py:766] *****
W0910 09:49:12.584000 4184 torch/distributed/run.py:766] Setting OMP_NUM_THREADS environment variable for each process to be 1 in default, to avoid
your system being overloaded, please further tune the variable for optimal performance in your application as needed.
W0910 09:49:12.584000 4184 torch/distributed/run.py:766] *****
[INFO] Starting configuration loading...
[INFO] Loaded configuration file: examples/qwen3_5/e1_norecomp100.yaml
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdp.data.data_utils.func_utils.mod
el_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.model_name_or_path (type: <class 'str'>, value: /root/Qwen3.5-
0.8B-hf)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdp.data.data_utils.func_utils.mod
el_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.use_fast_tokenizer (type: <class 'bool'>, value: True)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdp.data.data_utils.func_utils.mod
el_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.split_special_tokens (type: <class 'bool'>, value: False)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdp.data.data_utils.func_utils.mod
el_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.image_max_pixels (type: <class 'int'>, value: 262144)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdp.data.data_utils.func_utils.mod
el_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.image_min_pixels (type: <class 'int'>, value: 1024)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdp.data.data_utils.func_utils.mod
el_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.video_max_pixels (type: <class 'int'>, value: 16384)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdp.data.data_utils.func_utils.mod
el_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.video_min_pixels (type: <class 'int'>, value: 0)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdp.data.data_utils.func_utils.mod
el_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.video_fps (type: <class 'float'>, value: 2.0)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdp.data.data_utils.func_utils.mod
el_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.video_maxlen (type: <class 'int'>, value: 64)
[INFO] Automatically registered YAML field: model.gdn_implementation (type: <class 'str'>, value: triton)
[INFO] Automatically registered YAML field: model.causal_conv1d_implementation (type: <class 'str'>, value: triton)
[INFO] Automatically registered YAML field: model.skip_gdn_recompute (type: <class 'bool'>, value: False)
[INFO] Automatically registered YAML field: model.skip_flash_attn_recompute (type: <class 'bool'>, value: False)
[INFO] Automatically registered YAML field: features.recompute (type: <class 'bool'>, value: False)
=====
BaseArguments Configuration Manager Summary
=====

Configuration class: Arguments
Configuration file: examples/qwen3_5/e1_norecomp100.yaml
Allow YAML extensions: True
Allow CLI override: True
Strict CLI validation: True
Allow registering YAML fields: True

Defined field count: 208
Dynamic field count: 14
Total field count: 222

===== Configuration Details =====
parallel:
  data_parallel_size: 2
  fully_shard_parallel_size: 2
  fsdp_plan:
    ignored_modules: []
    apply_modules: ['model.visual', 'model.visual.blocks.*', 'model.language_model.embed_tokens', 'model.language_model', 'model.language_model.la
yers.*'], 'lm_head', 'mtp']
    param_dtype: bf16
    reduce_dtype: fp32
    output_dtype: None
    cast_forward_inputs: True
    reshard_after_forward: False
    num_to_forward_prefetch: 1
    num_to_backward_prefetch: 1
    pregather: False
    hook_modules: None
    cpu_offload: False
  tensor_parallel_size: 1
  tp_plan:
    colwise_parallel: []
    rowwise_parallel: []
    sequence_parallel: []
  ring_attention_size: 1
  ulysses_parallel_size: 1
  expert_parallel_size: 1
  expert_fully_shard_parallel_size: 2
  ep_plan:
    apply_modules: []
    use_npu_fused_ops: True
```

```

    dispatcher: alltoall
    apply_efsd_modules: []
recompute: False
recompute_plan:
    apply_modules: []
    use_reentrant: False
local_rank: 0
global_rank: 0
world_size: 2
model:
    model_id: qwen3_5
    model_name_or_path: /root/Qwen3.5-0.8B-hf
    trust_remote_code: True
    attn_implementation: flash_attention_2
    freeze: ['model.visual']
    mtp_num_layers: 0
    mtp_loss_scaling_factor: 0.1
    gdn_implementation: triton
    causal_conv1d_implementation: triton
    skip_gdn_recompute: False
    skip_flash_attn_recompute: False
data:
    dataset_param:
        dataset_type: huggingface
        basic_parameters:
            cache_dir: ./cache_dir/
            template: qwen3_vl_nothink
            chat_template: None
            enable_thinking: False
            dataset_dir: /root/data/coco
            dataset: ['/root/data/output_llava_coco_data.json']
            cutoff_len: 1024
            train_on_prompt: False
            mask_history: False
            streaming: False
            overwrite_cache: False
            preprocessing_batch_size: 1000
            preprocessing_num_workers: 16
            max_samples: None
            tool_format: None
            val_dataset: None
            val_max_samples: None
            val_rate: None
            packing: None
            neat_packing: False
            preprocess_on_fly: True
            async_preprocess: False
            async_preprocess_buffer_size: None
            stage: sft
        preprocess_parameters:
            model_name_or_path: /root/Qwen3.5-0.8B-hf
            use_fast_tokenizer: True
            split_special_tokens: False
            image_resolution: 512
            video_resolution: 128
            image_max_pixels: 262144
            image_min_pixels: 1024
            video_max_pixels: 16384
            video_min_pixels: 0
            video_fps: 2.0
            video_maxlen: 64
            image_do_pan_and_scan: False
            crop_to_patches: False
            use_audio_in_video: False
            audio_sampling_rate: 16000
            trust_remote_code: False
    attr:
        ranking: False
        system: None
        images: images
        videos: None
        audios: None
        prompt: None
        query: None
        response: None
        history: None
        messages: messages
        tools: None
        role_tag: role
        content_tag: content
        user_tag: user
        assistant_tag: assistant
        observation_tag: observation
        function_tag: function_call
        system_tag: system
        chosen: None
        rejected: None
        formatting: sharegpt
val_dataset_param: None
dataloading_param:
    dataloader_mode: sampler
    sampler_type: BaseRandomBatchSampler
    shuffle: False

```

```

drop_last: True
pin_memory: True
collate_param:
  collator_id: None
  model_name: qwen3v1
  ignore_pad_token_for_loss: True
  pad_to_multiple_of: 8
num_workers: 8
enable_preload: False
training:
  profile:
    enable: False
    start_step: 1
    end_step: 2
    save_path: ./profiling
    record_shapes: True
    with_memory: True
    with_stack: True
    ranks: [0]
    _train_steps: -1
    local_rank: 0
    global_rank: 0
    world_size: 2
    profile_this_rank: False
lora:
  enable: False
  rank: 8
  alpha: 16
  target_modules: ['q_proj', 'k_proj', 'v_proj']
  dropout: 0.0
  init_lora_weights: True
  pretrained_lora_path: None
  lora_target_modules_support: None
  disable_peft_moe_conversion: True
lr: 1e-05
lr_min: 0.0
lr_start: 0.0
weight_decay: 0.0
no_decay_modules: []
no_decay_params: []
optimizer: adamw
mem_fraction_static: 0.8
matched_adamw_rms: 0.2
muon_momentum: 0.95
ns_steps: 5
adam_fused: True
adam_beta1: 0.9
adam_beta2: 0.999
adam_eps: 1e-08
clip_grad: 0.0
clip_grad_foreach: True
micro_batch_size: 4
gradient_accumulation_steps: 1
lr_warmup_ratio: 0.1
lr_decay_style: cosine
lr_decay_ratio: 1.0
enable_full_determinism: False
ckpt_manager: dcp
save_async: False
save_ckpt_dtype: None
load_checkpoint_path: None
save_steps: 0
seed: 42
max_steps: None
init_model_with_meta_device: True
train_iters: 100
train_epochs: None
load: /root/Qwen3.5-0.8B-dcp
load_format: auto
save_format: auto
load_strict: False
load_rank0_and_broadcast: False
no_load_optim: True
no_load_rng: True
save: ./save_path
no_save_optim: True
no_save_rng: True
log_interval: 1
save_interval: 10000
val_interval: 0
val_micro_batch_size: None
use_deter_comp: False
allow_hf32: None
plugin: ['mindspeed_mm/fsdp/models/qwen3_5', 'mindspeed_mm/fsdp/data/datasets/huggingface']
lr_scaling_plan: None
quantization_plan: QuantizeConfig(recipe_name=None, apply_modules=[], ignored_modules=[], quant_converters=[], extra_args={}, enable_fsdp_low_precision_all_gather=True, fsdp_low_precision_all_gather_mode='on-demand')
_train_steps: -1
local_rank: 0
global_rank: 0
world_size: 2
global_batch_size: 8
tools:

```

```

profile:
  enable: False
  profile_type: static
  ranks: [0]
  static_param:
    level: level1
    with_stack: False
    with_memory: False
    record_shapes: False
    with_cpu: True
    save_path: ./profiling
    start_step: 10
    end_step: 11
    data_simplification: False
    aic_metrics_type: PipeUtilization
    analyse_flag: True
memory_profile:
  enable: False
  start_step: 1
  end_step: 2
  save_path: ./memory_snapshot
  dump_ranks: [0]
  stacks: all
  max_entries: None
  mem_info: False
features:
  recompute_plan:
    apply_modules: ['model.visual.blocks.*', 'model.language_model.layers.*']
    use_reentrant: False
  loss_cfg:
    loss_type: default
    router_aux_loss_coef: 0.0
    router_aux_loss_offload: False
  enable_chunk_loss: False
  enable_dynamic_chunk_loss: False
  chunkloss_plan:
    apply_module: lm_head
    chunk_size: 1024
    total_chunk_size: 4096
  enable_activation_offload: False
  activation_offload_plan:
    apply_modules: ['model.visual.blocks.*', 'model.language_model.layers.*']
  enable_chunk_mbs: False
  chunkmbs_plan:
    apply_modules: ['model.language_model.layers.*']
    chunk_mbs: 2
    batch_dim: 0
    chunk_arg_indexes: [0]
    chunk_kwang_names: ['position_embeddings', 'position_ids', 'rope_deltas', 'attention_mask']
  enable_ep_balance: False
  ep_balance_plan:
    max_dup_experts_num: 2
  skip_moe_pad_tokens: False
  recompute: False
data.dataset_param.preprocess_parameters.model_name_or_path: /root/Qwen3.5-0.8B-hf
data.dataset_param.preprocess_parameters.use_fast_tokenizer: True
data.dataset_param.preprocess_parameters.split_special_tokens: False
data.dataset_param.preprocess_parameters.image_max_pixels: 262144
data.dataset_param.preprocess_parameters.image_min_pixels: 1024
data.dataset_param.preprocess_parameters.video_max_pixels: 16384
data.dataset_param.preprocess_parameters.video_min_pixels: 0
data.dataset_param.preprocess_parameters.video_fps: 2.0
data.dataset_param.preprocess_parameters.video_maxlen: 64

```

===== Dynamic Fields =====

```

data.dataset_param.preprocess_parameters.image_max_pixels
data.dataset_param.preprocess_parameters.image_min_pixels
data.dataset_param.preprocess_parameters.model_name_or_path
data.dataset_param.preprocess_parameters.split_special_tokens
data.dataset_param.preprocess_parameters.use_fast_tokenizer
data.dataset_param.preprocess_parameters.video_fps
data.dataset_param.preprocess_parameters.video_max_pixels
data.dataset_param.preprocess_parameters.video_maxlen
data.dataset_param.preprocess_parameters.video_min_pixels
features.recompute
model.causal_conv1d_implementation
model.gdn_implementation
model.skip_flash_attn_recompute
model.skip_gdn_recompute

```

=====

Successfully imported: mindspeed_mm.fsd.models.qwen3_5

Successfully imported: mindspeed_mm.fsd.data.datasets.huggingface

[Rank 1 | Local Rank 1] 2026-09-10 09:49:20,191 INFO [datasets:114] => Disabling Tensorflow because USE_TORCH is set

[Rank 0 | Local Rank 0] 2026-09-10 09:49:20,338 INFO [datasets:114] => Disabling Tensorflow because USE_TORCH is set

[Rank 0 | Local Rank 0] 2026-09-10 09:49:20,684 INFO [mindspeed_mm.fsd.distributed.parallel_state:87] => Parallel state initialized:

[dp] = True | Group size: 2 | device mesh:DeviceMesh('np', [0, 1], mesh_dim_names=('dp',))

[cp] = False | Group size: 1 | device mesh:DeviceMesh('np', [0], mesh_dim_names=('cp',))

[ulysses] = False | Group size: 1 | device mesh:DeviceMesh('np', [0], mesh_dim_names=('ulysses',))

[ring] = False | Group size: 1 | device mesh:DeviceMesh('np', [0], mesh_dim_names=('ring',))

[tp] = False | Group size: 1 | device mesh:DeviceMesh('np', [0], mesh_dim_names=('tp',))

[edp] = False | Group size: 1 | device mesh:DeviceMesh('np', [0], mesh_dim_names=('edp',))

[efsdp] = True | Group size: 2 | device mesh:DeviceMesh('np', [0, 1], mesh_dim_names=('efsdp',))

[illegible]

[illegible]

[illegible]


```
+00 | grad norm: 12.834 |
[Rank 0 | Local Rank 0] 2026-09-10 09:49:45,625 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-10 09:49:45] iteration      10/    100
| consumed samples:      80 | elapsed time per iteration (ms): 634.3 | learning rate: 9.000000E-06 | global batch size: 8 | loss: 1.406257E
+00 | grad norm: 10.902 |
[Rank 0 | Local Rank 0] 2026-09-10 09:49:46,096 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-10 09:49:46] iteration      11/    100
| consumed samples:      88 | elapsed time per iteration (ms): 470.4 | learning rate: 1.000000E-05 | global batch size: 8 | loss: 1.433005E
+00 | grad norm: 12.064 |
[Rank 0 | Local Rank 0] 2026-09-10 09:49:46,705 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-10 09:49:46] iteration      12/    100
| consumed samples:      96 | elapsed time per iteration (ms): 608.1 | learning rate: 9.996954E-06 | global batch size: 8 | loss: 1.401638E
+00 | grad norm: 11.561 |
[Rank 0 | Local Rank 0] 2026-09-10 09:49:47,169 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-10 09:49:47] iteration      13/    100
| consumed samples:     104 | elapsed time per iteration (ms): 461.9 | learning rate: 9.987820E-06 | global batch size: 8 | loss: 1.467012E
+00 | grad norm: 10.043 |
[Rank 0 | Local Rank 0] 2026-09-10 09:49:47,665 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-10 09:49:47] iteration      14/    100
| consumed samples:     112 | elapsed time per iteration (ms): 495.2 | learning rate: 9.972609E-06 | global batch size: 8 | loss: 1.471380E
+00 | grad norm: 11.257 |
[Rank 0 | Local Rank 0] 2026-09-10 09:49:48,249 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-10 09:49:48] iteration      15/    100
| consumed samples:     120 | elapsed time per iteration (ms): 582.9 | learning rate: 9.951340E-06 | global batch size: 8 | loss: 1.418947E
+00 | grad norm: 9.819 |
[Rank 0 | Local Rank 0] 2026-09-10 09:49:48,695 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-10 09:49:48] iteration      16/    100
| consumed samples:     128 | elapsed time per iteration (ms): 444.7 | learning rate: 9.924039E-06 | global batch size: 8 | loss: 1.429941E
+00 | grad norm: 11.276 |
[Rank 0 | Local Rank 0] 2026-09-10 09:49:49,166 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-10 09:49:49] iteration      17/    100
| consumed samples:     136 | elapsed time per iteration (ms): 470.1 | learning rate: 9.890738E-06 | global batch size: 8 | loss: 1.410848E
+00 | grad norm: 10.318 |
[Rank 0 | Local Rank 0] 2026-09-10 09:49:49,597 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-10 09:49:49] iteration      18/    100
| consumed samples:     144 | elapsed time per iteration (ms): 430.0 | learning rate: 9.851479E-06 | global batch size: 8 | loss: 1.379771E
+00 | grad norm: 11.859 |
[Rank 0 | Local Rank 0] 2026-09-10 09:49:50,054 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-10 09:49:50] iteration      19/    100
| consumed samples:     152 | elapsed time per iteration (ms): 455.1 | learning rate: 9.806308E-06 | global batch size: 8 | loss: 1.459756E
+00 | grad norm: 10.009 |
[Rank 0 | Local Rank 0] 2026-09-10 09:49:50,520 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-10 09:49:50] iteration      20/    100
| consumed samples:     160 | elapsed time per iteration (ms): 464.6 | learning rate: 9.755283E-06 | global batch size: 8 | loss: 1.350282E
+00 | grad norm: 10.378 |
[Rank 0 | Local Rank 0] 2026-09-10 09:49:50,976 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-10 09:49:50] iteration      21/    100
| consumed samples:     168 | elapsed time per iteration (ms): 454.9 | learning rate: 9.698463E-06 | global batch size: 8 | loss: 1.255694E
+00 | grad norm: 10.727 |
[Rank 0 | Local Rank 0] 2026-09-10 09:49:51,442 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-10 09:49:51] iteration      22/    100
| consumed samples:     176 | elapsed time per iteration (ms): 465.2 | learning rate: 9.635919E-06 | global batch size: 8 | loss: 1.431852E
+00 | grad norm: 9.803 |
[Rank 0 | Local Rank 0] 2026-09-10 09:49:51,885 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-10 09:49:51] iteration      23/    100
| consumed samples:     184 | elapsed time per iteration (ms): 442.1 | learning rate: 9.567727E-06 | global batch size: 8 | loss: 1.479843E
+00 | grad norm: 11.539 |
[Rank 0 | Local Rank 0] 2026-09-10 09:49:52,343 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-10 09:49:52] iteration      24/    100
| consumed samples:     192 | elapsed time per iteration (ms): 456.6 | learning rate: 9.493970E-06 | global batch size: 8 | loss: 1.482672E
+00 | grad norm: 9.782 |
[Rank 0 | Local Rank 0] 2026-09-10 09:49:52,773 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-10 09:49:52] iteration      25/    100
| consumed samples:     200 | elapsed time per iteration (ms): 428.7 | learning rate: 9.414738E-06 | global batch size: 8 | loss: 1.420358E
+00 | grad norm: 11.043 |
[Rank 0 | Local Rank 0] 2026-09-10 09:49:53,184 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-10 09:49:53] iteration      26/    100
| consumed samples:     208 | elapsed time per iteration (ms): 409.6 | learning rate: 9.330127E-06 | global batch size: 8 | loss: 1.313674E
+00 | grad norm: 11.962 |
[Rank 0 | Local Rank 0] 2026-09-10 09:49:53,638 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-10 09:49:53] iteration      27/    100
| consumed samples:     216 | elapsed time per iteration (ms): 453.9 | learning rate: 9.240240E-06 | global batch size: 8 | loss: 1.477658E
+00 | grad norm: 10.432 |
[Rank 0 | Local Rank 0] 2026-09-10 09:49:54,084 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-10 09:49:54] iteration      28/    100
| consumed samples:     224 | elapsed time per iteration (ms): 445.0 | learning rate: 9.145188E-06 | global batch size: 8 | loss: 1.352777E
+00 | grad norm: 10.369 |
[Rank 0 | Local Rank 0] 2026-09-10 09:49:54,528 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-10 09:49:54] iteration      29/    100
| consumed samples:     232 | elapsed time per iteration (ms): 442.8 | learning rate: 9.045085E-06 | global batch size: 8 | loss: 1.451961E
+00 | grad norm: 10.201 |
[Rank 0 | Local Rank 0] 2026-09-10 09:49:54,954 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-10 09:49:54] iteration      30/    100
| consumed samples:     240 | elapsed time per iteration (ms): 425.4 | learning rate: 8.940054E-06 | global batch size: 8 | loss: 1.285751E
+00 | grad norm: 9.373 |
[Rank 0 | Local Rank 0] 2026-09-10 09:49:55,475 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-10 09:49:55] iteration      31/    100
| consumed samples:     248 | elapsed time per iteration (ms): 519.6 | learning rate: 8.830222E-06 | global batch size: 8 | loss: 1.319847E
+00 | grad norm: 9.474 |
[Rank 0 | Local Rank 0] 2026-09-10 09:49:55,906 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-10 09:49:55] iteration      32/    100
| consumed samples:     256 | elapsed time per iteration (ms): 430.2 | learning rate: 8.715724E-06 | global batch size: 8 | loss: 1.436016E
+00 | grad norm: 9.372 |
[Rank 0 | Local Rank 0] 2026-09-10 09:49:56,361 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-10 09:49:56] iteration      33/    100
| consumed samples:     264 | elapsed time per iteration (ms): 453.7 | learning rate: 8.596699E-06 | global batch size: 8 | loss: 1.278239E
+00 | grad norm: 9.758 |
[Rank 0 | Local Rank 0] 2026-09-10 09:49:56,798 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-10 09:49:56] iteration      34/    100
| consumed samples:     272 | elapsed time per iteration (ms): 436.5 | learning rate: 8.473292E-06 | global batch size: 8 | loss: 1.292705E
+00 | grad norm: 10.339 |
[Rank 0 | Local Rank 0] 2026-09-10 09:49:57,315 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-10 09:49:57] iteration      35/    100
| consumed samples:     280 | elapsed time per iteration (ms): 515.7 | learning rate: 8.345653E-06 | global batch size: 8 | loss: 1.411542E
+00 | grad norm: 9.255 |
[Rank 0 | Local Rank 0] 2026-09-10 09:49:57,778 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-10 09:49:57] iteration      36/    100
| consumed samples:     288 | elapsed time per iteration (ms): 461.7 | learning rate: 8.213938E-06 | global batch size: 8 | loss: 1.407445E
+00 | grad norm: 9.932 |
[Rank 0 | Local Rank 0] 2026-09-10 09:49:58,246 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-10 09:49:58] iteration      37/    100
| consumed samples:     296 | elapsed time per iteration (ms): 467.7 | learning rate: 8.078307E-06 | global batch size: 8 | loss: 1.293699E
+00 | grad norm: 11.826 |
[Rank 0 | Local Rank 0] 2026-09-10 09:49:58,714 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-10 09:49:58] iteration      38/    100
| consumed samples:     304 | elapsed time per iteration (ms): 467.0 | learning rate: 7.938926E-06 | global batch size: 8 | loss: 1.469754E
+00 | grad norm: 10.296 |
[Rank 0 | Local Rank 0] 2026-09-10 09:49:59,183 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-10 09:49:59] iteration      39/    100
| consumed samples:     312 | elapsed time per iteration (ms): 467.2 | learning rate: 7.795965E-06 | global batch size: 8 | loss: 1.354045E
+00 | grad norm: 9.548 |
[Rank 0 | Local Rank 0] 2026-09-10 09:49:59,637 INFO [mindspeed_mm.fsdsp.train.train_engine:48] => [2026-09-10 09:49:59] iteration      40/    100
| consumed samples:     320 | elapsed time per iteration (ms): 453.6 | learning rate: 7.649596E-06 | global batch size: 8 | loss: 1.514582E
```

```
+00 | grad norm: 8.940 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:00,096 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 09:50:00] iteration 41/ 100
| consumed samples: 328 | elapsed time per iteration (ms): 457.9 | learning rate: 7.500000E-06 | global batch size: 8 | loss: 1.402503E
+00 | grad norm: 10.470 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:00,561 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 09:50:00] iteration 42/ 100
| consumed samples: 336 | elapsed time per iteration (ms): 464.3 | learning rate: 7.347358E-06 | global batch size: 8 | loss: 1.288520E
+00 | grad norm: 9.212 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:00,987 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 09:50:00] iteration 43/ 100
| consumed samples: 344 | elapsed time per iteration (ms): 424.1 | learning rate: 7.191856E-06 | global batch size: 8 | loss: 1.323021E
+00 | grad norm: 10.306 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:01,448 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 09:50:01] iteration 44/ 100
| consumed samples: 352 | elapsed time per iteration (ms): 460.1 | learning rate: 7.033683E-06 | global batch size: 8 | loss: 1.359287E
+00 | grad norm: 11.091 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:01,873 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 09:50:01] iteration 45/ 100
| consumed samples: 360 | elapsed time per iteration (ms): 424.6 | learning rate: 6.873033E-06 | global batch size: 8 | loss: 1.289140E
+00 | grad norm: 9.465 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:02,284 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 09:50:02] iteration 46/ 100
| consumed samples: 368 | elapsed time per iteration (ms): 409.9 | learning rate: 6.710101E-06 | global batch size: 8 | loss: 1.328894E
+00 | grad norm: 10.447 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:02,722 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 09:50:02] iteration 47/ 100
| consumed samples: 376 | elapsed time per iteration (ms): 436.9 | learning rate: 6.545085E-06 | global batch size: 8 | loss: 1.390787E
+00 | grad norm: 10.827 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:03,679 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 09:50:03] iteration 48/ 100
| consumed samples: 384 | elapsed time per iteration (ms): 956.2 | learning rate: 6.378187E-06 | global batch size: 8 | loss: 1.471148E
+00 | grad norm: 9.606 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:04,108 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 09:50:04] iteration 49/ 100
| consumed samples: 392 | elapsed time per iteration (ms): 427.0 | learning rate: 6.209609E-06 | global batch size: 8 | loss: 1.271755E
+00 | grad norm: 10.066 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:07,968 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 09:50:07] iteration 50/ 100
| consumed samples: 400 | elapsed time per iteration (ms): 3858.2 | learning rate: 6.039558E-06 | global batch size: 8 | loss: 1.284707
E+00 | grad norm: 9.127 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:08,406 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 09:50:08] iteration 51/ 100
| consumed samples: 408 | elapsed time per iteration (ms): 437.4 | learning rate: 5.868241E-06 | global batch size: 8 | loss: 1.360355E
+00 | grad norm: 9.243 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:08,889 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 09:50:08] iteration 52/ 100
| consumed samples: 416 | elapsed time per iteration (ms): 482.2 | learning rate: 5.695866E-06 | global batch size: 8 | loss: 1.556308E
+00 | grad norm: 9.373 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:09,342 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 09:50:09] iteration 53/ 100
| consumed samples: 424 | elapsed time per iteration (ms): 451.4 | learning rate: 5.522642E-06 | global batch size: 8 | loss: 1.302439E
+00 | grad norm: 10.075 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:09,792 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 09:50:09] iteration 54/ 100
| consumed samples: 432 | elapsed time per iteration (ms): 449.2 | learning rate: 5.348782E-06 | global batch size: 8 | loss: 1.279586E
+00 | grad norm: 9.589 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:10,239 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 09:50:10] iteration 55/ 100
| consumed samples: 440 | elapsed time per iteration (ms): 445.6 | learning rate: 5.174497E-06 | global batch size: 8 | loss: 1.300742E
+00 | grad norm: 10.616 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:10,702 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 09:50:10] iteration 56/ 100
| consumed samples: 448 | elapsed time per iteration (ms): 462.1 | learning rate: 5.000000E-06 | global batch size: 8 | loss: 1.402216E
+00 | grad norm: 9.485 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:11,170 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 09:50:11] iteration 57/ 100
| consumed samples: 456 | elapsed time per iteration (ms): 467.8 | learning rate: 4.825503E-06 | global batch size: 8 | loss: 1.392442E
+00 | grad norm: 9.644 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:11,607 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 09:50:11] iteration 58/ 100
| consumed samples: 464 | elapsed time per iteration (ms): 435.1 | learning rate: 4.651218E-06 | global batch size: 8 | loss: 1.378122E
+00 | grad norm: 8.848 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:12,031 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 09:50:12] iteration 59/ 100
| consumed samples: 472 | elapsed time per iteration (ms): 422.8 | learning rate: 4.477358E-06 | global batch size: 8 | loss: 1.212556E
+00 | grad norm: 10.249 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:12,446 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 09:50:12] iteration 60/ 100
| consumed samples: 480 | elapsed time per iteration (ms): 414.8 | learning rate: 4.304134E-06 | global batch size: 8 | loss: 1.269083E
+00 | grad norm: 10.539 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:12,869 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 09:50:12] iteration 61/ 100
| consumed samples: 488 | elapsed time per iteration (ms): 421.2 | learning rate: 4.131759E-06 | global batch size: 8 | loss: 1.300264E
+00 | grad norm: 9.550 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:13,311 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 09:50:13] iteration 62/ 100
| consumed samples: 496 | elapsed time per iteration (ms): 441.8 | learning rate: 3.960442E-06 | global batch size: 8 | loss: 1.174930E
+00 | grad norm: 11.484 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:13,747 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 09:50:13] iteration 63/ 100
| consumed samples: 504 | elapsed time per iteration (ms): 433.9 | learning rate: 3.790391E-06 | global batch size: 8 | loss: 1.406163E
+00 | grad norm: 9.716 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:14,194 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 09:50:14] iteration 64/ 100
| consumed samples: 512 | elapsed time per iteration (ms): 446.7 | learning rate: 3.621813E-06 | global batch size: 8 | loss: 1.334237E
+00 | grad norm: 10.402 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:14,614 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 09:50:14] iteration 65/ 100
| consumed samples: 520 | elapsed time per iteration (ms): 418.9 | learning rate: 3.454915E-06 | global batch size: 8 | loss: 1.296929E
+00 | grad norm: 10.465 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:15,063 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 09:50:15] iteration 66/ 100
| consumed samples: 528 | elapsed time per iteration (ms): 448.0 | learning rate: 3.289899E-06 | global batch size: 8 | loss: 1.415588E
+00 | grad norm: 9.577 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:15,515 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 09:50:15] iteration 67/ 100
| consumed samples: 536 | elapsed time per iteration (ms): 451.0 | learning rate: 3.126967E-06 | global batch size: 8 | loss: 1.293933E
+00 | grad norm: 9.836 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:15,964 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 09:50:15] iteration 68/ 100
| consumed samples: 544 | elapsed time per iteration (ms): 447.8 | learning rate: 2.966317E-06 | global batch size: 8 | loss: 1.508645E
+00 | grad norm: 10.325 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:16,424 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 09:50:16] iteration 69/ 100
| consumed samples: 552 | elapsed time per iteration (ms): 459.4 | learning rate: 2.808144E-06 | global batch size: 8 | loss: 1.426050E
+00 | grad norm: 9.362 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:16,839 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 09:50:16] iteration 70/ 100
| consumed samples: 560 | elapsed time per iteration (ms): 413.6 | learning rate: 2.652642E-06 | global batch size: 8 | loss: 1.333285E
+00 | grad norm: 10.353 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:17,296 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-10 09:50:17] iteration 71/ 100
| consumed samples: 568 | elapsed time per iteration (ms): 456.0 | learning rate: 2.500000E-06 | global batch size: 8 | loss: 1.278005E
```

```
+00 | grad norm: 10.738 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:17,762 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-10 09:50:17] iteration      72/    100
| consumed samples:      576 | elapsed time per iteration (ms): 464.8 | learning rate: 2.350404E-06 | global batch size:    8 | loss: 1.400092E
+00 | grad norm: 9.749 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:18,219 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-10 09:50:18] iteration      73/    100
| consumed samples:      584 | elapsed time per iteration (ms): 456.2 | learning rate: 2.204035E-06 | global batch size:    8 | loss: 1.378221E
+00 | grad norm: 9.594 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:18,688 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-10 09:50:18] iteration      74/    100
| consumed samples:      592 | elapsed time per iteration (ms): 468.0 | learning rate: 2.061074E-06 | global batch size:    8 | loss: 1.327019E
+00 | grad norm: 10.213 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:19,100 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-10 09:50:19] iteration      75/    100
| consumed samples:      600 | elapsed time per iteration (ms): 411.5 | learning rate: 1.921693E-06 | global batch size:    8 | loss: 1.366259E
+00 | grad norm: 9.568 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:19,564 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-10 09:50:19] iteration      76/    100
| consumed samples:      608 | elapsed time per iteration (ms): 462.8 | learning rate: 1.786062E-06 | global batch size:    8 | loss: 1.336825E
+00 | grad norm: 9.708 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:20,024 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-10 09:50:20] iteration      77/    100
| consumed samples:      616 | elapsed time per iteration (ms): 458.5 | learning rate: 1.654347E-06 | global batch size:    8 | loss: 1.400733E
+00 | grad norm: 9.561 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:20,463 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-10 09:50:20] iteration      78/    100
| consumed samples:      624 | elapsed time per iteration (ms): 437.4 | learning rate: 1.526708E-06 | global batch size:    8 | loss: 1.419183E
+00 | grad norm: 9.869 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:20,933 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-10 09:50:20] iteration      79/    100
| consumed samples:      632 | elapsed time per iteration (ms): 469.8 | learning rate: 1.403301E-06 | global batch size:    8 | loss: 1.447845E
+00 | grad norm: 9.662 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:21,387 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-10 09:50:21] iteration      80/    100
| consumed samples:      640 | elapsed time per iteration (ms): 452.3 | learning rate: 1.284276E-06 | global batch size:    8 | loss: 1.276349E
+00 | grad norm: 8.811 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:21,815 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-10 09:50:21] iteration      81/    100
| consumed samples:      648 | elapsed time per iteration (ms): 427.4 | learning rate: 1.169778E-06 | global batch size:    8 | loss: 1.295292E
+00 | grad norm: 10.192 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:22,274 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-10 09:50:22] iteration      82/    100
| consumed samples:      656 | elapsed time per iteration (ms): 458.2 | learning rate: 1.059946E-06 | global batch size:    8 | loss: 1.274825E
+00 | grad norm: 9.983 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:22,707 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-10 09:50:22] iteration      83/    100
| consumed samples:      664 | elapsed time per iteration (ms): 431.5 | learning rate: 9.549150E-07 | global batch size:    8 | loss: 1.264797E
+00 | grad norm: 9.445 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:23,151 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-10 09:50:23] iteration      84/    100
| consumed samples:      672 | elapsed time per iteration (ms): 443.0 | learning rate: 8.548121E-07 | global batch size:    8 | loss: 1.233971E
+00 | grad norm: 9.197 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:24,000 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-10 09:50:23] iteration      85/    100
| consumed samples:      680 | elapsed time per iteration (ms): 846.9 | learning rate: 7.597595E-07 | global batch size:    8 | loss: 1.319490E
+00 | grad norm: 11.693 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:24,426 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-10 09:50:24] iteration      86/    100
| consumed samples:      688 | elapsed time per iteration (ms): 424.8 | learning rate: 6.698730E-07 | global batch size:    8 | loss: 1.360535E
+00 | grad norm: 10.703 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:24,893 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-10 09:50:24] iteration      87/    100
| consumed samples:      696 | elapsed time per iteration (ms): 465.8 | learning rate: 5.852620E-07 | global batch size:    8 | loss: 1.367505E
+00 | grad norm: 8.778 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:25,353 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-10 09:50:25] iteration      88/    100
| consumed samples:      704 | elapsed time per iteration (ms): 459.4 | learning rate: 5.060298E-07 | global batch size:    8 | loss: 1.278803E
+00 | grad norm: 10.091 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:25,820 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-10 09:50:25] iteration      89/    100
| consumed samples:      712 | elapsed time per iteration (ms): 466.0 | learning rate: 4.322727E-07 | global batch size:    8 | loss: 1.110864E
+00 | grad norm: 9.701 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:26,283 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-10 09:50:26] iteration      90/    100
| consumed samples:      720 | elapsed time per iteration (ms): 461.5 | learning rate: 3.640807E-07 | global batch size:    8 | loss: 1.301911E
+00 | grad norm: 9.602 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:26,779 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-10 09:50:26] iteration      91/    100
| consumed samples:      728 | elapsed time per iteration (ms): 495.1 | learning rate: 3.015369E-07 | global batch size:    8 | loss: 1.352304E
+00 | grad norm: 8.144 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:27,216 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-10 09:50:27] iteration      92/    100
| consumed samples:      736 | elapsed time per iteration (ms): 435.6 | learning rate: 2.447174E-07 | global batch size:    8 | loss: 1.289772E
+00 | grad norm: 10.380 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:27,679 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-10 09:50:27] iteration      93/    100
| consumed samples:      744 | elapsed time per iteration (ms): 462.1 | learning rate: 1.936915E-07 | global batch size:    8 | loss: 1.200276E
+00 | grad norm: 10.073 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:28,162 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-10 09:50:28] iteration      94/    100
| consumed samples:      752 | elapsed time per iteration (ms): 482.0 | learning rate: 1.485214E-07 | global batch size:    8 | loss: 1.441172E
+00 | grad norm: 9.691 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:28,610 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-10 09:50:28] iteration      95/    100
| consumed samples:      760 | elapsed time per iteration (ms): 446.7 | learning rate: 1.092620E-07 | global batch size:    8 | loss: 1.313918E
+00 | grad norm: 9.774 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:29,156 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-10 09:50:29] iteration      96/    100
| consumed samples:      768 | elapsed time per iteration (ms): 544.6 | learning rate: 7.596123E-08 | global batch size:    8 | loss: 1.322249E
+00 | grad norm: 8.624 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:29,597 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-10 09:50:29] iteration      97/    100
| consumed samples:      776 | elapsed time per iteration (ms): 440.6 | learning rate: 4.865966E-08 | global batch size:    8 | loss: 1.329803E
+00 | grad norm: 10.000 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:30,001 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-10 09:50:30] iteration      98/    100
| consumed samples:      784 | elapsed time per iteration (ms): 402.7 | learning rate: 2.739052E-08 | global batch size:    8 | loss: 1.307319E
+00 | grad norm: 11.291 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:30,458 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-10 09:50:30] iteration      99/    100
| consumed samples:      792 | elapsed time per iteration (ms): 455.5 | learning rate: 1.217975E-08 | global batch size:    8 | loss: 1.363874E
+00 | grad norm: 10.150 |
[Rank 0 | Local Rank 0] 2026-09-10 09:50:30,906 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-10 09:50:30] iteration     100/    100
| consumed samples:      800 | elapsed time per iteration (ms): 447.0 | learning rate: 3.045865E-09 | global batch size:    8 | loss: 1.448290E
+00 | grad norm: 10.069 |
/usr/local/lib/python3.10/dist-packages/torch/distributed/checkpoint/filesystem.py:592: UserWarning: Detected an existing checkpoint in save_path/it
er_0000100/.metadata, overwriting since self.overwrite=True. Past version 2.5 of PyTorch, 'overwrite' will default to False. Set this variable to Tr
ue to maintain this functionality or False to raise when an existing checkpoint is found.
  warnings.warn(
/usr/local/lib/python3.10/dist-packages/torch/distributed/checkpoint/filesystem.py:592: UserWarning: Detected an existing checkpoint in save_path/it
```

```
er_0000100/.metadata, overwriting since self.overwrite=True. Past version 2.5 of PyTorch, `overwrite` will default to False. Set this variable to True to maintain this functionality or False to raise when an existing checkpoint is found.
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/torch_npu/distributed/distributed_c10d.py:110: UserWarning: HCCL doesn't support gather at the moment. Implemented with allgather instead.
    warnings.warn("HCCL doesn't support gather at the moment. Implemented with allgather instead.")
/usr/local/lib/python3.10/dist-packages/torch/distributed/checkpoint/filesystem.py:111: UserWarning: TypedStorage is deprecated. It will be removed in the future and UntypedStorage will be the only storage class. This should only matter to you if you are using storages directly. To access UntypedStorage directly, use tensor.untyped_storage() instead of tensor.storage()
    if tensor.storage().size() != tensor.numel():
/usr/local/lib/python3.10/dist-packages/torch/distributed/checkpoint/filesystem.py:111: UserWarning: TypedStorage is deprecated. It will be removed in the future and UntypedStorage will be the only storage class. This should only matter to you if you are using storages directly. To access UntypedStorage directly, use tensor.untyped_storage() instead of tensor.storage()
    if tensor.storage().size() != tensor.numel():
[Rank 0 | Local Rank 0] 2026-09-10 09:50:36,894 INFO [mindspeed_mm.fsdcp.checkpoint.dcp_checkpointinter:48] => Saved checkpoint to ./save_path/iter_0000100
```

20260904_A2_910B4_最终配置100步

文件: 20260904_A2_910B4_最终配置100步.log | 行数: 682 | 大小: 90.7 KB | A2 (910B4 单卡) 最终配置 100 步 —— 跨卡对照基线
采集说明: 完整日志 (从程序启动开始, 未经裁剪)

```
[INFO] Starting configuration loading...
[INFO] Loaded configuration file: examples/qwen3_5/qwen3_5_0_8B_ze2.yaml
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdp.data.data_utils.func_utils.model_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.model_name_or_path (type: <class 'str'>, value: /root/Qwen3.5-0.8B-hf)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdp.data.data_utils.func_utils.model_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.use_fast_tokenizer (type: <class 'bool'>, value: True)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdp.data.data_utils.func_utils.model_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.split_special_tokens (type: <class 'bool'>, value: False)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdp.data.data_utils.func_utils.model_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.image_max_pixels (type: <class 'int'>, value: 262144)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdp.data.data_utils.func_utils.model_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.image_min_pixels (type: <class 'int'>, value: 1024)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdp.data.data_utils.func_utils.model_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.video_max_pixels (type: <class 'int'>, value: 16384)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdp.data.data_utils.func_utils.model_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.video_min_pixels (type: <class 'int'>, value: 0)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdp.data.data_utils.func_utils.model_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.video_fps (type: <class 'float'>, value: 2.0)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdp.data.data_utils.func_utils.model_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.video_maxlen (type: <class 'int'>, value: 64)
[INFO] Automatically registered YAML field: model.gdn_implementation (type: <class 'str'>, value: triton)
[INFO] Automatically registered YAML field: model.causal_conv1d_implementation (type: <class 'str'>, value: triton)
[INFO] Automatically registered YAML field: model.skip_gdn_recompute (type: <class 'bool'>, value: False)
[INFO] Automatically registered YAML field: model.skip_flash_attn_recompute (type: <class 'bool'>, value: False)
[INFO] Automatically registered YAML field: features.recompute (type: <class 'bool'>, value: True)
=====
BaseArguments Configuration Manager Summary
=====

Configuration class: Arguments
Configuration file: examples/qwen3_5/qwen3_5_0_8B_ze2.yaml
Allow YAML extensions: True
Allow CLI override: True
Strict CLI validation: True
Allow registering YAML fields: True

Defined field count: 203
Dynamic field count: 14
Total field count: 217

===== Configuration Details =====
parallel:
  data_parallel_size: 1
  fully_shard_parallel_size: 1
  fsdp_plan:
    ignored_modules: []
    apply_modules: ['model.visual', 'model.visual.blocks.*', 'model.language_model.embed_tokens', 'model.language_model', 'model.language_model.layers.*', 'lm_head', 'mtp']
    param_dtype: bf16
    reduce_dtype: fp32
    output_dtype: None
    cast_forward_inputs: True
    reshard_after_forward: False
    num_to_forward_prefetch: 1
    num_to_backward_prefetch: 1
    pregather: False
    hook_modules: None
    cpu_offload: False
  tensor_parallel_size: 1
  tp_plan:
    colwise_parallel: []
    rowwise_parallel: []
    sequence_parallel: []
  ring_attention_size: 1
  ulysses_parallel_size: 1
  expert_parallel_size: 1
  expert_fully_shard_parallel_size: 1
  ep_plan:
    apply_modules: []
    use_npu_fused_ops: True
    dispatcher: alltoall
    apply_efsdp_modules: []
  recompute: False
  recompute_plan:
    apply_modules: []
```

```
    use_reentrant: False
    local_rank: 0
    global_rank: 0
    world_size: 1
model:
  model_id: qwen3_5
  model_name_or_path: /root/Qwen3.5-0.8B-hf
  trust_remote_code: True
  attn_implementation: flash_attention_2
  freeze: ['model.visual']
  mtp_num_layers: 0
  mtp_loss_scaling_factor: 0.1
  gdn_implementation: triton
  causal_conv1d_implementation: triton
  skip_gdn_recompute: False
  skip_flash_attn_recompute: False
data:
  dataset_param:
    dataset_type: huggingface
    basic_parameters:
      cache_dir: ./cache_dir/
      template: qwen3_vl_nothink
      chat_template: None
      enable_thinking: False
      dataset_dir: /root/coco_d1/extracted/dataset
      dataset: ['/root/coco_d1/extracted/dataset/annotations_slim.json']
      cutoff_len: 1024
      train_on_prompt: False
      mask_history: False
      streaming: False
      overwrite_cache: False
      preprocessing_batch_size: 1000
      preprocessing_num_workers: 2
      max_samples: None
      tool_format: None
      val_dataset: None
      val_max_samples: None
      val_rate: None
      packing: None
      neat_packing: False
      preprocess_on_fly: True
      async_preprocess: False
      async_preprocess_buffer_size: None
      stage: sft
    preprocess_parameters:
      model_name_or_path: /root/Qwen3.5-0.8B-hf
      use_fast_tokenizer: True
      split_special_tokens: False
      image_resolution: 512
      video_resolution: 128
      image_max_pixels: 262144
      image_min_pixels: 1024
      video_max_pixels: 16384
      video_min_pixels: 0
      video_fps: 2.0
      video_maxlen: 64
      image_do_pan_and_scan: False
      crop_to_patches: False
      use_audio_in_video: False
      audio_sampling_rate: 16000
      trust_remote_code: False
  attr:
    ranking: False
    system: None
    images: images
    videos: None
    audios: None
    prompt: None
    query: None
    response: None
    history: None
    messages: messages
    tools: None
    role_tag: role
    content_tag: content
    user_tag: user
    assistant_tag: assistant
    observation_tag: observation
    function_tag: function_call
    system_tag: system
    chosen: None
    rejected: None
    formatting: sharegpt
  val_dataset_param: None
  dataloader_param:
    dataloader_mode: sampler
    sampler_type: BaseRandomBatchSampler
    shuffle: False
    drop_last: True
    pin_memory: True
    collate_param:
      collator_id: None
      model_name: qwen3vl
```

```

        ignore_pad_token_for_loss: True
        pad_to_multiple_of: 8
        num_workers: 2
        enable_preload: False
training:
  profile:
    enable: False
    start_step: 1
    end_step: 2
    save_path: ./profiling
    record_shapes: True
    with_memory: True
    with_stack: True
    ranks: [0]
    _train_steps: -1
    local_rank: 0
    global_rank: 0
    world_size: 1
    profile_this_rank: False
lora:
  enable: False
  rank: 8
  alpha: 16
  target_modules: ['q_proj', 'k_proj', 'v_proj']
  dropout: 0.0
  init_lora_weights: True
  pretrained_lora_path: None
  lora_target_modules_support: None
lr: 1e-05
lr_min: 0.0
lr_start: 0.0
weight_decay: 0.0
no_decay_modules: []
no_decay_params: []
optimizer: adamw
matched_adamw_rms: 0.2
muon_momentum: 0.95
ns_steps: 5
adam_fused: True
adam_beta1: 0.9
adam_beta2: 0.999
adam_eps: 1e-08
clip_grad: 0.0
clip_grad_foreach: True
micro_batch_size: 2
gradient_accumulation_steps: 4
lr_warmup_ratio: 0.1
lr_decay_style: cosine
lr_decay_ratio: 1.0
enable_full_determinism: False
ckpt_manager: dcp
save_async: False
load_checkpoint_path: None
save_steps: 0
seed: 42
max_steps: None
init_model_with_meta_device: True
train_iters: 100
train_epochs: None
load: /root/Qwen3.5-0.8B-dcp
load_format: auto
load_strict: False
load_rank0_and_broadcast: False
no_load_optim: True
no_load_rng: True
save: ./save_path
no_save_optim: True
no_save_rng: True
log_interval: 1
save_interval: 10000
val_interval: 0
val_micro_batch_size: None
use_deter_comp: False
allow_hf32: None
plugin: ['mindspeed_mm/fsdp/models/qwen3_5', 'mindspeed_mm/fsdp/data/datasets/huggingface']
lr_scaling_plan: None
quantization_plan: QuantizeConfig(recipe_name=None, apply_modules=[], ignored_modules=[], quant_converters=[], extra_args={}, enable_fsdp_low_precision_all_gather=True, fsdp_low_precision_all_gather_mode='on-demand')
  _train_steps: -1
  local_rank: 0
  global_rank: 0
  world_size: 1
  global_batch_size: 8
tools:
  profile:
    enable: False
    profile_type: static
    ranks: [0]
    static_param:
      level: level1
      with_stack: False
      with_memory: False
      record_shapes: False

```

```

with_cpu: True
save_path: ./profiling
start_step: 10
end_step: 11
data_simplification: False
aic_metrics_type: PipeUtilization
analyse_flag: True
memory_profile:
  enable: False
  start_step: 1
  end_step: 2
  save_path: ./memory_snapshot
  dump_ranks: [0]
  stacks: all
  max_entries: None
  mem_info: False
features:
  recompute_plan:
    apply_modules: ['model.visual.blocks.*', 'model.language_model.layers.*']
    use_reentrant: False
  loss_cfg:
    loss_type: default
    router_aux_loss_coef: 0.0
    router_aux_loss_offload: False
  enable_chunk_loss: True
  enable_dynamic_chunk_loss: False
  chunkloss_plan:
    apply_module: lm_head
    chunk_size: 1024
    total_chunk_size: 4096
  enable_activation_offload: True
  activation_offload_plan:
    apply_modules: ['model.visual.blocks.*', 'model.language_model.layers.*']
  enable_chunk_mbs: False
  chunkmbs_plan:
    apply_modules: ['model.language_model.layers.*']
    chunk_mbs: 2
    batch_dim: 0
    chunk_arg_indexes: [0]
    chunk_kwargs_names: ['position_embeddings', 'position_ids', 'rope_deltas', 'attention_mask']
  enable_ep_balance: False
  ep_balance_plan:
    max_dup_experts_num: 2
  recompute: True
data.dataset_param.preprocess_parameters.model_name_or_path: /root/Qwen3.5-0.8B-hf
data.dataset_param.preprocess_parameters.use_fast_tokenizer: True
data.dataset_param.preprocess_parameters.split_special_tokens: False
data.dataset_param.preprocess_parameters.image_max_pixels: 262144
data.dataset_param.preprocess_parameters.image_min_pixels: 1024
data.dataset_param.preprocess_parameters.video_max_pixels: 16384
data.dataset_param.preprocess_parameters.video_min_pixels: 0
data.dataset_param.preprocess_parameters.video_fps: 2.0
data.dataset_param.preprocess_parameters.video_maxlen: 64

```

===== Dynamic Fields =====

```

data.dataset_param.preprocess_parameters.image_max_pixels
data.dataset_param.preprocess_parameters.image_min_pixels
data.dataset_param.preprocess_parameters.model_name_or_path
data.dataset_param.preprocess_parameters.split_special_tokens
data.dataset_param.preprocess_parameters.use_fast_tokenizer
data.dataset_param.preprocess_parameters.video_fps
data.dataset_param.preprocess_parameters.video_max_pixels
data.dataset_param.preprocess_parameters.video_maxlen
data.dataset_param.preprocess_parameters.video_min_pixels
features.recompute
model.causal_conv1d_implementation
model.gdn_implementation
model.skip_flash_attn_recompute
model.skip_gdn_recompute

```

=====

Successfully imported: mindspeed_mm.fsdps.models.qwen3_5

Successfully imported: mindspeed_mm.fsdps.data.datasets.huggingface

[Rank 0 | Local Rank 0] 2026-09-03 07:57:30,781 INFO [datasets:114] => Disabling Tensorflow because USE_TORCH is set

[Rank 0 | Local Rank 0] 2026-09-03 07:57:31,085 INFO [mindspeed_mm.fsdps.distributed.parallel_state:87] => Parallel state initialized:

```

[dp] = False | Group size: 1 | device mesh:DeviceMesh('npu', [0], mesh_dim_names=('dp',))
[cp] = False | Group size: 1 | device mesh:DeviceMesh('npu', [0], mesh_dim_names=('cp',))
[ulysses] = False | Group size: 1 | device mesh:DeviceMesh('npu', [0], mesh_dim_names=('ulysses',))
[ring] = False | Group size: 1 | device mesh:DeviceMesh('npu', [0], mesh_dim_names=('ring',))
[tp] = False | Group size: 1 | device mesh:DeviceMesh('npu', [0], mesh_dim_names=('tp',))
[edp] = False | Group size: 1 | device mesh:DeviceMesh('npu', [0], mesh_dim_names=('edp',))
[efsdp] = False | Group size: 1 | device mesh:DeviceMesh('npu', [0], mesh_dim_names=('efsdp',))
[ep] = False | Group size: 1 | device mesh:DeviceMesh('npu', [0], mesh_dim_names=('ep',))
[fsdp] = True | Group size: 1 | device mesh:DeviceMesh('npu', [0], mesh_dim_names=('dp_shard_cp',))

```

[Rank 0 | Local Rank 0] 2026-09-03 07:57:31,086 INFO [mindspeed_mm.fsdps.models.modelhub:48] => > Loading AutoConfig from /root/Qwen3.5-0.8B-hf...

[Rank 0 | Local Rank 0] 2026-09-03 07:57:31,090 INFO [mindspeed_mm.fsdps.models.modelhub:48] => Building transformers model from configuration...

[Rank 0 | Local Rank 0] 2026-09-03 07:57:31,180 INFO [mindspeed_mm.fsdps.models.qwen3_5.modeling_qwen3_5:48] => Qwen3_5 causal_conv1d use NPU triton ops

[Rank 0 | Local Rank 0] 2026-09-03 07:57:31,357 INFO [mindspeed_mm.fsdps.models.qwen3_5.modeling_qwen3_5:48] => Qwen3_5GatedDeltaNet use NPU triton f used ops

[Rank 0 | Local Rank 0] 2026-09-03 07:57:31,357 WARNING [mindspeed_mm.fsdps.models.qwen3_5.modeling_qwen3_5:327] => The fast path is not available because one of the required library is not installed. Falling back to torch implementation.

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

```
ctly, use tensor.untyped_storage() instead of tensor.storage()
  if tensor.storage().size() <= 0:
[rank0]:[W903 07:57:50.154915630 compiler_depend.ts:335] Warning: Cannot create tensor with interal format while allow_internel_format=False, tensor
will be created with base format. (function operator())
[Rank 0 | Local Rank 0] 2026-09-03 07:57:58,693 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 07:57:58] iteration      1/   100
| consumed samples:      8 | elapsed time per iteration (ms): 9918.9 | learning rate: 0.000000E+00 | global batch size:   8 | loss: 1.928093
E+00 | grad norm: 137.134 |
[Rank 0 | Local Rank 0] 2026-09-03 07:58:01,806 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 07:58:01] iteration      2/   100
| consumed samples:     16 | elapsed time per iteration (ms): 3110.8 | learning rate: 1.000000E-06 | global batch size:   8 | loss: 1.787479
E+00 | grad norm: 114.441 |
[Rank 0] (after 2 iterations) memory (MB) | allocated: 9001.89111328125 | max allocated: 20300.5595703125 | reserved: 21968.0 | max reserved: 21968.
0
[Rank 0 | Local Rank 0] 2026-09-03 07:58:04,793 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 07:58:04] iteration      3/   100
| consumed samples:     24 | elapsed time per iteration (ms): 2985.2 | learning rate: 2.000000E-06 | global batch size:   8 | loss: 1.788652
E+00 | grad norm: 91.294 |
[Rank 0 | Local Rank 0] 2026-09-03 07:58:07,806 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 07:58:07] iteration      4/   100
| consumed samples:     32 | elapsed time per iteration (ms): 3011.1 | learning rate: 3.000000E-06 | global batch size:   8 | loss: 1.478453
E+00 | grad norm: 62.176 |
[Rank 0 | Local Rank 0] 2026-09-03 07:58:10,819 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 07:58:10] iteration      5/   100
| consumed samples:     40 | elapsed time per iteration (ms): 3011.5 | learning rate: 4.000000E-06 | global batch size:   8 | loss: 1.460526
E+00 | grad norm: 26.904 |
[Rank 0 | Local Rank 0] 2026-09-03 07:58:13,823 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 07:58:13] iteration      6/   100
| consumed samples:     48 | elapsed time per iteration (ms): 3002.4 | learning rate: 5.000000E-06 | global batch size:   8 | loss: 1.403486
E+00 | grad norm: 17.358 |
[Rank 0 | Local Rank 0] 2026-09-03 07:58:16,712 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 07:58:16] iteration      7/   100
| consumed samples:     56 | elapsed time per iteration (ms): 2887.4 | learning rate: 6.000000E-06 | global batch size:   8 | loss: 1.332761
E+00 | grad norm: 15.949 |
[Rank 0 | Local Rank 0] 2026-09-03 07:58:19,658 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 07:58:19] iteration      8/   100
| consumed samples:     64 | elapsed time per iteration (ms): 2944.2 | learning rate: 7.000000E-06 | global batch size:   8 | loss: 1.453734
E+00 | grad norm: 11.894 |
[Rank 0 | Local Rank 0] 2026-09-03 07:58:23,236 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 07:58:23] iteration      9/   100
| consumed samples:     72 | elapsed time per iteration (ms): 3576.2 | learning rate: 8.000000E-06 | global batch size:   8 | loss: 1.285888
E+00 | grad norm: 12.836 |
[Rank 0 | Local Rank 0] 2026-09-03 07:58:26,278 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 07:58:26] iteration     10/   100
| consumed samples:     80 | elapsed time per iteration (ms): 3040.0 | learning rate: 9.000000E-06 | global batch size:   8 | loss: 1.421257
E+00 | grad norm: 11.377 |
[Rank 0 | Local Rank 0] 2026-09-03 07:58:29,334 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 07:58:29] iteration     11/   100
| consumed samples:     88 | elapsed time per iteration (ms): 3054.9 | learning rate: 1.000000E-05 | global batch size:   8 | loss: 1.348977
E+00 | grad norm: 13.195 |
[Rank 0 | Local Rank 0] 2026-09-03 07:58:32,419 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 07:58:32] iteration     12/   100
| consumed samples:     96 | elapsed time per iteration (ms): 3082.9 | learning rate: 9.996954E-06 | global batch size:   8 | loss: 1.435174
E+00 | grad norm: 11.115 |
[Rank 0 | Local Rank 0] 2026-09-03 07:58:35,398 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 07:58:35] iteration     13/   100
| consumed samples:    104 | elapsed time per iteration (ms): 2977.7 | learning rate: 9.987820E-06 | global batch size:   8 | loss: 1.441972
E+00 | grad norm: 10.084 |
[Rank 0 | Local Rank 0] 2026-09-03 07:58:38,343 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 07:58:38] iteration     14/   100
| consumed samples:    112 | elapsed time per iteration (ms): 2942.7 | learning rate: 9.972609E-06 | global batch size:   8 | loss: 1.417805
E+00 | grad norm: 11.076 |
[Rank 0 | Local Rank 0] 2026-09-03 07:58:41,284 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 07:58:41] iteration     15/   100
| consumed samples:    120 | elapsed time per iteration (ms): 2939.6 | learning rate: 9.951340E-06 | global batch size:   8 | loss: 1.410591
E+00 | grad norm: 9.920 |
[Rank 0 | Local Rank 0] 2026-09-03 07:58:44,154 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 07:58:44] iteration     16/   100
| consumed samples:    128 | elapsed time per iteration (ms): 2868.7 | learning rate: 9.924039E-06 | global batch size:   8 | loss: 1.460164
E+00 | grad norm: 11.802 |
[Rank 0 | Local Rank 0] 2026-09-03 07:58:47,438 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 07:58:47] iteration     17/   100
| consumed samples:    136 | elapsed time per iteration (ms): 3282.6 | learning rate: 9.890738E-06 | global batch size:   8 | loss: 1.361181
E+00 | grad norm: 11.015 |
[Rank 0 | Local Rank 0] 2026-09-03 07:58:50,371 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 07:58:50] iteration     18/   100
| consumed samples:    144 | elapsed time per iteration (ms): 2931.7 | learning rate: 9.851479E-06 | global batch size:   8 | loss: 1.280402
E+00 | grad norm: 11.909 |
[Rank 0 | Local Rank 0] 2026-09-03 07:58:53,330 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 07:58:53] iteration     19/   100
| consumed samples:    152 | elapsed time per iteration (ms): 2957.2 | learning rate: 9.806308E-06 | global batch size:   8 | loss: 1.453319
E+00 | grad norm: 10.115 |
[Rank 0 | Local Rank 0] 2026-09-03 07:58:56,223 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 07:58:56] iteration     20/   100
| consumed samples:    160 | elapsed time per iteration (ms): 2891.0 | learning rate: 9.755283E-06 | global batch size:   8 | loss: 1.253502
E+00 | grad norm: 11.712 |
[Rank 0 | Local Rank 0] 2026-09-03 07:58:59,158 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 07:58:59] iteration     21/   100
| consumed samples:    168 | elapsed time per iteration (ms): 2933.2 | learning rate: 9.698463E-06 | global batch size:   8 | loss: 1.202060
E+00 | grad norm: 10.782 |
[Rank 0 | Local Rank 0] 2026-09-03 07:59:02,093 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 07:59:02] iteration     22/   100
| consumed samples:    176 | elapsed time per iteration (ms): 2933.1 | learning rate: 9.635919E-06 | global batch size:   8 | loss: 1.435723
E+00 | grad norm: 10.193 |
[Rank 0 | Local Rank 0] 2026-09-03 07:59:04,978 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 07:59:04] iteration     23/   100
| consumed samples:    184 | elapsed time per iteration (ms): 2883.4 | learning rate: 9.567727E-06 | global batch size:   8 | loss: 1.435873
E+00 | grad norm: 12.801 |
[Rank 0 | Local Rank 0] 2026-09-03 07:59:07,919 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 07:59:07] iteration     24/   100
| consumed samples:    192 | elapsed time per iteration (ms): 2939.7 | learning rate: 9.493970E-06 | global batch size:   8 | loss: 1.490324
E+00 | grad norm: 9.869 |
[Rank 0 | Local Rank 0] 2026-09-03 07:59:11,573 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 07:59:11] iteration     25/   100
| consumed samples:    200 | elapsed time per iteration (ms): 3652.6 | learning rate: 9.414738E-06 | global batch size:   8 | loss: 1.419205
E+00 | grad norm: 10.951 |
[Rank 0 | Local Rank 0] 2026-09-03 07:59:14,623 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 07:59:14] iteration     26/   100
| consumed samples:    208 | elapsed time per iteration (ms): 3048.7 | learning rate: 9.330127E-06 | global batch size:   8 | loss: 1.296585
E+00 | grad norm: 12.014 |
[Rank 0 | Local Rank 0] 2026-09-03 07:59:17,652 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 07:59:17] iteration     27/   100
| consumed samples:    216 | elapsed time per iteration (ms): 3027.5 | learning rate: 9.240240E-06 | global batch size:   8 | loss: 1.454204
E+00 | grad norm: 11.011 |
[Rank 0 | Local Rank 0] 2026-09-03 07:59:20,769 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 07:59:20] iteration     28/   100
| consumed samples:    224 | elapsed time per iteration (ms): 3115.3 | learning rate: 9.145188E-06 | global batch size:   8 | loss: 1.317015
E+00 | grad norm: 10.639 |
[Rank 0 | Local Rank 0] 2026-09-03 07:59:23,815 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 07:59:23] iteration     29/   100
| consumed samples:    232 | elapsed time per iteration (ms): 3043.8 | learning rate: 9.045085E-06 | global batch size:   8 | loss: 1.395229
E+00 | grad norm: 10.124 |
```

[Rank 0 | Local Rank 0] 2026-09-03 07:59:26,862 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-03 07:59:26] iteration 30/ 100
| consumed samples: 240 | elapsed time per iteration (ms): 3045.3 | learning rate: 8.940054E-06 | global batch size: 8 | loss: 1.327958
E+00 | grad norm: 10.062 |
[Rank 0 | Local Rank 0] 2026-09-03 07:59:29,882 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-03 07:59:29] iteration 31/ 100
| consumed samples: 248 | elapsed time per iteration (ms): 3018.5 | learning rate: 8.830222E-06 | global batch size: 8 | loss: 1.329273
E+00 | grad norm: 9.496 |
[Rank 0 | Local Rank 0] 2026-09-03 07:59:33,017 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-03 07:59:33] iteration 32/ 100
| consumed samples: 256 | elapsed time per iteration (ms): 3133.7 | learning rate: 8.715724E-06 | global batch size: 8 | loss: 1.456537
E+00 | grad norm: 9.625 |
[Rank 0 | Local Rank 0] 2026-09-03 07:59:35,960 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-03 07:59:35] iteration 33/ 100
| consumed samples: 264 | elapsed time per iteration (ms): 2940.5 | learning rate: 8.596699E-06 | global batch size: 8 | loss: 1.227463
E+00 | grad norm: 9.974 |
[Rank 0 | Local Rank 0] 2026-09-03 07:59:39,515 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-03 07:59:39] iteration 34/ 100
| consumed samples: 272 | elapsed time per iteration (ms): 3553.8 | learning rate: 8.473292E-06 | global batch size: 8 | loss: 1.290263
E+00 | grad norm: 10.650 |
[Rank 0 | Local Rank 0] 2026-09-03 07:59:42,749 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-03 07:59:42] iteration 35/ 100
| consumed samples: 280 | elapsed time per iteration (ms): 3231.5 | learning rate: 8.345653E-06 | global batch size: 8 | loss: 1.417686
E+00 | grad norm: 9.162 |
[Rank 0 | Local Rank 0] 2026-09-03 07:59:45,826 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-03 07:59:45] iteration 36/ 100
| consumed samples: 288 | elapsed time per iteration (ms): 3076.0 | learning rate: 8.213938E-06 | global batch size: 8 | loss: 1.381018
E+00 | grad norm: 9.860 |
[Rank 0 | Local Rank 0] 2026-09-03 07:59:48,827 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-03 07:59:48] iteration 37/ 100
| consumed samples: 296 | elapsed time per iteration (ms): 2999.9 | learning rate: 8.078307E-06 | global batch size: 8 | loss: 1.338350
E+00 | grad norm: 11.196 |
[Rank 0 | Local Rank 0] 2026-09-03 07:59:51,880 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-03 07:59:51] iteration 38/ 100
| consumed samples: 304 | elapsed time per iteration (ms): 3050.4 | learning rate: 7.938926E-06 | global batch size: 8 | loss: 1.435714
E+00 | grad norm: 11.801 |
[Rank 0 | Local Rank 0] 2026-09-03 07:59:55,002 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-03 07:59:55] iteration 39/ 100
| consumed samples: 312 | elapsed time per iteration (ms): 3120.7 | learning rate: 7.795965E-06 | global batch size: 8 | loss: 1.341281
E+00 | grad norm: 9.661 |
[Rank 0 | Local Rank 0] 2026-09-03 07:59:58,096 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-03 07:59:58] iteration 40/ 100
| consumed samples: 320 | elapsed time per iteration (ms): 3091.6 | learning rate: 7.649596E-06 | global batch size: 8 | loss: 1.514152
E+00 | grad norm: 9.013 |
[Rank 0 | Local Rank 0] 2026-09-03 08:00:01,161 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-03 08:00:01] iteration 41/ 100
| consumed samples: 328 | elapsed time per iteration (ms): 3063.7 | learning rate: 7.500000E-06 | global batch size: 8 | loss: 1.391773
E+00 | grad norm: 10.879 |
[Rank 0 | Local Rank 0] 2026-09-03 08:00:04,234 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-03 08:00:04] iteration 42/ 100
| consumed samples: 336 | elapsed time per iteration (ms): 3071.5 | learning rate: 7.347358E-06 | global batch size: 8 | loss: 1.241714
E+00 | grad norm: 9.663 |
[Rank 0 | Local Rank 0] 2026-09-03 08:00:07,995 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-03 08:00:07] iteration 43/ 100
| consumed samples: 344 | elapsed time per iteration (ms): 3758.9 | learning rate: 7.191856E-06 | global batch size: 8 | loss: 1.328652
E+00 | grad norm: 10.425 |
[Rank 0 | Local Rank 0] 2026-09-03 08:00:11,149 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-03 08:00:11] iteration 44/ 100
| consumed samples: 352 | elapsed time per iteration (ms): 3152.2 | learning rate: 7.033683E-06 | global batch size: 8 | loss: 1.330024
E+00 | grad norm: 10.591 |
[Rank 0 | Local Rank 0] 2026-09-03 08:00:14,266 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-03 08:00:14] iteration 45/ 100
| consumed samples: 360 | elapsed time per iteration (ms): 3115.4 | learning rate: 6.873033E-06 | global batch size: 8 | loss: 1.251678
E+00 | grad norm: 10.240 |
[Rank 0 | Local Rank 0] 2026-09-03 08:00:17,377 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-03 08:00:17] iteration 46/ 100
| consumed samples: 368 | elapsed time per iteration (ms): 3108.9 | learning rate: 6.710101E-06 | global batch size: 8 | loss: 1.354432
E+00 | grad norm: 10.738 |
[Rank 0 | Local Rank 0] 2026-09-03 08:00:20,407 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-03 08:00:20] iteration 47/ 100
| consumed samples: 376 | elapsed time per iteration (ms): 3028.6 | learning rate: 6.545085E-06 | global batch size: 8 | loss: 1.419736
E+00 | grad norm: 11.090 |
[Rank 0 | Local Rank 0] 2026-09-03 08:00:23,540 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-03 08:00:23] iteration 48/ 100
| consumed samples: 384 | elapsed time per iteration (ms): 3131.0 | learning rate: 6.378187E-06 | global batch size: 8 | loss: 1.492616
E+00 | grad norm: 9.734 |
[Rank 0 | Local Rank 0] 2026-09-03 08:00:26,610 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-03 08:00:26] iteration 49/ 100
| consumed samples: 392 | elapsed time per iteration (ms): 3067.9 | learning rate: 6.209609E-06 | global batch size: 8 | loss: 1.219772
E+00 | grad norm: 10.486 |
[Rank 0 | Local Rank 0] 2026-09-03 08:00:34,189 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-03 08:00:34] iteration 50/ 100
| consumed samples: 400 | elapsed time per iteration (ms): 7575.6 | learning rate: 6.039558E-06 | global batch size: 8 | loss: 1.207329
E+00 | grad norm: 9.274 |
[Rank 0 | Local Rank 0] 2026-09-03 08:00:37,834 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-03 08:00:37] iteration 51/ 100
| consumed samples: 408 | elapsed time per iteration (ms): 3643.6 | learning rate: 5.868241E-06 | global batch size: 8 | loss: 1.372535
E+00 | grad norm: 9.554 |
[Rank 0 | Local Rank 0] 2026-09-03 08:00:40,918 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-03 08:00:40] iteration 52/ 100
| consumed samples: 416 | elapsed time per iteration (ms): 3082.4 | learning rate: 5.695866E-06 | global batch size: 8 | loss: 1.537471
E+00 | grad norm: 9.649 |
[Rank 0 | Local Rank 0] 2026-09-03 08:00:43,971 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-03 08:00:43] iteration 53/ 100
| consumed samples: 424 | elapsed time per iteration (ms): 3051.7 | learning rate: 5.522642E-06 | global batch size: 8 | loss: 1.302097
E+00 | grad norm: 10.081 |
[Rank 0 | Local Rank 0] 2026-09-03 08:00:46,998 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-03 08:00:46] iteration 54/ 100
| consumed samples: 432 | elapsed time per iteration (ms): 3025.5 | learning rate: 5.348782E-06 | global batch size: 8 | loss: 1.302673
E+00 | grad norm: 10.239 |
[Rank 0 | Local Rank 0] 2026-09-03 08:00:50,041 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-03 08:00:50] iteration 55/ 100
| consumed samples: 440 | elapsed time per iteration (ms): 3041.0 | learning rate: 5.174497E-06 | global batch size: 8 | loss: 1.234624
E+00 | grad norm: 11.013 |
[Rank 0 | Local Rank 0] 2026-09-03 08:00:53,099 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-03 08:00:53] iteration 56/ 100
| consumed samples: 448 | elapsed time per iteration (ms): 3056.0 | learning rate: 5.000000E-06 | global batch size: 8 | loss: 1.430351
E+00 | grad norm: 9.324 |
[Rank 0 | Local Rank 0] 2026-09-03 08:00:56,130 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-03 08:00:56] iteration 57/ 100
| consumed samples: 456 | elapsed time per iteration (ms): 3028.9 | learning rate: 4.825503E-06 | global batch size: 8 | loss: 1.389697
E+00 | grad norm: 9.510 |
[Rank 0 | Local Rank 0] 2026-09-03 08:00:59,176 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-03 08:00:59] iteration 58/ 100
| consumed samples: 464 | elapsed time per iteration (ms): 3044.9 | learning rate: 4.651218E-06 | global batch size: 8 | loss: 1.368783
E+00 | grad norm: 8.696 |
[Rank 0 | Local Rank 0] 2026-09-03 08:01:02,956 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-03 08:01:02] iteration 59/ 100
| consumed samples: 472 | elapsed time per iteration (ms): 3777.7 | learning rate: 4.477358E-06 | global batch size: 8 | loss: 1.178461
E+00 | grad norm: 9.908 |
[Rank 0 | Local Rank 0] 2026-09-03 08:01:05,992 INFO [mindspeed_mm.fsdpt.train.train_engine:48] => [2026-09-03 08:01:05] iteration 60/ 100
| consumed samples: 480 | elapsed time per iteration (ms): 3033.9 | learning rate: 4.304134E-06 | global batch size: 8 | loss: 1.261029
E+00 | grad norm: 10.948 |

[Rank 0 | Local Rank 0] 2026-09-03 08:01:09,016 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 08:01:09] iteration 61/ 100
| consumed samples: 488 | elapsed time per iteration (ms): 3022.2 | learning rate: 4.131759E-06 | global batch size: 8 | loss: 1.263459
E+00 | grad norm: 9.487 |
[Rank 0 | Local Rank 0] 2026-09-03 08:01:12,049 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 08:01:12] iteration 62/ 100
| consumed samples: 496 | elapsed time per iteration (ms): 3031.5 | learning rate: 3.960442E-06 | global batch size: 8 | loss: 1.184981
E+00 | grad norm: 11.182 |
[Rank 0 | Local Rank 0] 2026-09-03 08:01:15,191 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 08:01:15] iteration 63/ 100
| consumed samples: 504 | elapsed time per iteration (ms): 3141.2 | learning rate: 3.790391E-06 | global batch size: 8 | loss: 1.437771
E+00 | grad norm: 10.467 |
[Rank 0 | Local Rank 0] 2026-09-03 08:01:18,400 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 08:01:18] iteration 64/ 100
| consumed samples: 512 | elapsed time per iteration (ms): 3207.0 | learning rate: 3.621813E-06 | global batch size: 8 | loss: 1.324935
E+00 | grad norm: 10.262 |
[Rank 0 | Local Rank 0] 2026-09-03 08:01:21,570 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 08:01:21] iteration 65/ 100
| consumed samples: 520 | elapsed time per iteration (ms): 3164.8 | learning rate: 3.454915E-06 | global batch size: 8 | loss: 1.293230
E+00 | grad norm: 10.401 |
[Rank 0 | Local Rank 0] 2026-09-03 08:01:24,642 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 08:01:24] iteration 66/ 100
| consumed samples: 528 | elapsed time per iteration (ms): 3070.9 | learning rate: 3.289899E-06 | global batch size: 8 | loss: 1.349204
E+00 | grad norm: 9.911 |
[Rank 0 | Local Rank 0] 2026-09-03 08:01:27,724 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 08:01:27] iteration 67/ 100
| consumed samples: 536 | elapsed time per iteration (ms): 3079.0 | learning rate: 3.126967E-06 | global batch size: 8 | loss: 1.227242
E+00 | grad norm: 9.959 |
[Rank 0 | Local Rank 0] 2026-09-03 08:01:31,593 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 08:01:31] iteration 68/ 100
| consumed samples: 544 | elapsed time per iteration (ms): 3866.8 | learning rate: 2.966317E-06 | global batch size: 8 | loss: 1.503712
E+00 | grad norm: 11.535 |
[Rank 0 | Local Rank 0] 2026-09-03 08:01:34,754 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 08:01:34] iteration 69/ 100
| consumed samples: 552 | elapsed time per iteration (ms): 3159.1 | learning rate: 2.808144E-06 | global batch size: 8 | loss: 1.445690
E+00 | grad norm: 10.165 |
[Rank 0 | Local Rank 0] 2026-09-03 08:01:37,792 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 08:01:37] iteration 70/ 100
| consumed samples: 560 | elapsed time per iteration (ms): 3036.4 | learning rate: 2.652642E-06 | global batch size: 8 | loss: 1.275503
E+00 | grad norm: 10.565 |
[Rank 0 | Local Rank 0] 2026-09-03 08:01:40,837 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 08:01:40] iteration 71/ 100
| consumed samples: 568 | elapsed time per iteration (ms): 3044.3 | learning rate: 2.500000E-06 | global batch size: 8 | loss: 1.293976
E+00 | grad norm: 10.792 |
[Rank 0 | Local Rank 0] 2026-09-03 08:01:43,887 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 08:01:43] iteration 72/ 100
| consumed samples: 576 | elapsed time per iteration (ms): 3048.2 | learning rate: 2.350404E-06 | global batch size: 8 | loss: 1.404052
E+00 | grad norm: 10.184 |
[Rank 0 | Local Rank 0] 2026-09-03 08:01:46,907 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 08:01:46] iteration 73/ 100
| consumed samples: 584 | elapsed time per iteration (ms): 3017.5 | learning rate: 2.204035E-06 | global batch size: 8 | loss: 1.360970
E+00 | grad norm: 9.668 |
[Rank 0 | Local Rank 0] 2026-09-03 08:01:49,902 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 08:01:49] iteration 74/ 100
| consumed samples: 592 | elapsed time per iteration (ms): 2993.2 | learning rate: 2.061074E-06 | global batch size: 8 | loss: 1.286685
E+00 | grad norm: 10.644 |
[Rank 0 | Local Rank 0] 2026-09-03 08:01:52,965 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 08:01:52] iteration 75/ 100
| consumed samples: 600 | elapsed time per iteration (ms): 3060.6 | learning rate: 1.921693E-06 | global batch size: 8 | loss: 1.367728
E+00 | grad norm: 10.969 |
[Rank 0 | Local Rank 0] 2026-09-03 08:01:56,044 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 08:01:56] iteration 76/ 100
| consumed samples: 608 | elapsed time per iteration (ms): 3077.9 | learning rate: 1.786062E-06 | global batch size: 8 | loss: 1.345968
E+00 | grad norm: 9.341 |
[Rank 0 | Local Rank 0] 2026-09-03 08:01:59,503 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 08:01:59] iteration 77/ 100
| consumed samples: 616 | elapsed time per iteration (ms): 3457.2 | learning rate: 1.654347E-06 | global batch size: 8 | loss: 1.402554
E+00 | grad norm: 9.939 |
[Rank 0 | Local Rank 0] 2026-09-03 08:02:02,571 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 08:02:02] iteration 78/ 100
| consumed samples: 624 | elapsed time per iteration (ms): 3066.5 | learning rate: 1.526708E-06 | global batch size: 8 | loss: 1.396679
E+00 | grad norm: 9.987 |
[Rank 0 | Local Rank 0] 2026-09-03 08:02:05,731 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 08:02:05] iteration 79/ 100
| consumed samples: 632 | elapsed time per iteration (ms): 3158.5 | learning rate: 1.403301E-06 | global batch size: 8 | loss: 1.428060
E+00 | grad norm: 9.750 |
[Rank 0 | Local Rank 0] 2026-09-03 08:02:08,858 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 08:02:08] iteration 80/ 100
| consumed samples: 640 | elapsed time per iteration (ms): 3125.1 | learning rate: 1.284276E-06 | global batch size: 8 | loss: 1.249187
E+00 | grad norm: 8.882 |
[Rank 0 | Local Rank 0] 2026-09-03 08:02:11,974 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 08:02:11] iteration 81/ 100
| consumed samples: 648 | elapsed time per iteration (ms): 3114.1 | learning rate: 1.169778E-06 | global batch size: 8 | loss: 1.287505
E+00 | grad norm: 10.903 |
[Rank 0 | Local Rank 0] 2026-09-03 08:02:15,040 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 08:02:15] iteration 82/ 100
| consumed samples: 656 | elapsed time per iteration (ms): 3064.8 | learning rate: 1.059946E-06 | global batch size: 8 | loss: 1.222947
E+00 | grad norm: 10.375 |
[Rank 0 | Local Rank 0] 2026-09-03 08:02:18,053 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 08:02:18] iteration 83/ 100
| consumed samples: 664 | elapsed time per iteration (ms): 3011.9 | learning rate: 9.549150E-07 | global batch size: 8 | loss: 1.257640
E+00 | grad norm: 9.477 |
[Rank 0 | Local Rank 0] 2026-09-03 08:02:21,080 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 08:02:21] iteration 84/ 100
| consumed samples: 672 | elapsed time per iteration (ms): 3025.8 | learning rate: 8.548121E-07 | global batch size: 8 | loss: 1.182749
E+00 | grad norm: 9.401 |
[Rank 0 | Local Rank 0] 2026-09-03 08:02:24,465 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 08:02:24] iteration 85/ 100
| consumed samples: 680 | elapsed time per iteration (ms): 3383.1 | learning rate: 7.597595E-07 | global batch size: 8 | loss: 1.213120
E+00 | grad norm: 11.609 |
[Rank 0 | Local Rank 0] 2026-09-03 08:02:27,487 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 08:02:27] iteration 86/ 100
| consumed samples: 688 | elapsed time per iteration (ms): 3019.4 | learning rate: 6.698730E-07 | global batch size: 8 | loss: 1.352807
E+00 | grad norm: 11.076 |
[Rank 0 | Local Rank 0] 2026-09-03 08:02:30,592 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 08:02:30] iteration 87/ 100
| consumed samples: 696 | elapsed time per iteration (ms): 3102.9 | learning rate: 5.852620E-07 | global batch size: 8 | loss: 1.316924
E+00 | grad norm: 9.116 |
[Rank 0 | Local Rank 0] 2026-09-03 08:02:33,660 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 08:02:33] iteration 88/ 100
| consumed samples: 704 | elapsed time per iteration (ms): 3067.3 | learning rate: 5.060298E-07 | global batch size: 8 | loss: 1.176548
E+00 | grad norm: 10.029 |
[Rank 0 | Local Rank 0] 2026-09-03 08:02:36,674 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 08:02:36] iteration 89/ 100
| consumed samples: 712 | elapsed time per iteration (ms): 3011.9 | learning rate: 4.322727E-07 | global batch size: 8 | loss: 1.068325
E+00 | grad norm: 10.111 |
[Rank 0 | Local Rank 0] 2026-09-03 08:02:39,645 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 08:02:39] iteration 90/ 100
| consumed samples: 720 | elapsed time per iteration (ms): 2968.7 | learning rate: 3.640807E-07 | global batch size: 8 | loss: 1.242069
E+00 | grad norm: 9.891 |
[Rank 0 | Local Rank 0] 2026-09-03 08:02:42,737 INFO [mindspeed_mm.fsdp.train.train_engine:48] => [2026-09-03 08:02:42] iteration 91/ 100
| consumed samples: 728 | elapsed time per iteration (ms): 3091.0 | learning rate: 3.015369E-07 | global batch size: 8 | loss: 1.354719
E+00 | grad norm: 8.363 |

```
[Rank 0 | Local Rank 0] 2026-09-03 08:02:45,662 INFO [mindspeed_mm.fsd.train.train_engine:48] => [2026-09-03 08:02:45] iteration 92/ 100
| consumed samples: 736 | elapsed time per iteration (ms): 2922.4 | learning rate: 2.447174E-07 | global batch size: 8 | loss: 1.288227
E+00 | grad norm: 10.419 |
[Rank 0 | Local Rank 0] 2026-09-03 08:02:48,690 INFO [mindspeed_mm.fsd.train.train_engine:48] => [2026-09-03 08:02:48] iteration 93/ 100
| consumed samples: 744 | elapsed time per iteration (ms): 3027.2 | learning rate: 1.936915E-07 | global batch size: 8 | loss: 1.160878
E+00 | grad norm: 10.213 |
[Rank 0 | Local Rank 0] 2026-09-03 08:02:52,137 INFO [mindspeed_mm.fsd.train.train_engine:48] => [2026-09-03 08:02:52] iteration 94/ 100
| consumed samples: 752 | elapsed time per iteration (ms): 3444.9 | learning rate: 1.485214E-07 | global batch size: 8 | loss: 1.374482
E+00 | grad norm: 9.899 |
[Rank 0 | Local Rank 0] 2026-09-03 08:02:55,126 INFO [mindspeed_mm.fsd.train.train_engine:48] => [2026-09-03 08:02:55] iteration 95/ 100
| consumed samples: 760 | elapsed time per iteration (ms): 2987.3 | learning rate: 1.092620E-07 | global batch size: 8 | loss: 1.310308
E+00 | grad norm: 10.077 |
[Rank 0 | Local Rank 0] 2026-09-03 08:02:58,182 INFO [mindspeed_mm.fsd.train.train_engine:48] => [2026-09-03 08:02:58] iteration 96/ 100
| consumed samples: 768 | elapsed time per iteration (ms): 3055.1 | learning rate: 7.596123E-08 | global batch size: 8 | loss: 1.284240
E+00 | grad norm: 9.441 |
[Rank 0 | Local Rank 0] 2026-09-03 08:03:01,244 INFO [mindspeed_mm.fsd.train.train_engine:48] => [2026-09-03 08:03:01] iteration 97/ 100
| consumed samples: 776 | elapsed time per iteration (ms): 3060.1 | learning rate: 4.865966E-08 | global batch size: 8 | loss: 1.362904
E+00 | grad norm: 10.474 |
[Rank 0 | Local Rank 0] 2026-09-03 08:03:04,259 INFO [mindspeed_mm.fsd.train.train_engine:48] => [2026-09-03 08:03:04] iteration 98/ 100
| consumed samples: 784 | elapsed time per iteration (ms): 3013.4 | learning rate: 2.739052E-08 | global batch size: 8 | loss: 1.361278
E+00 | grad norm: 11.917 |
[Rank 0 | Local Rank 0] 2026-09-03 08:03:07,259 INFO [mindspeed_mm.fsd.train.train_engine:48] => [2026-09-03 08:03:07] iteration 99/ 100
| consumed samples: 792 | elapsed time per iteration (ms): 2998.7 | learning rate: 1.217975E-08 | global batch size: 8 | loss: 1.280154
E+00 | grad norm: 9.923 |
[Rank 0 | Local Rank 0] 2026-09-03 08:03:10,166 INFO [mindspeed_mm.fsd.train.train_engine:48] => [2026-09-03 08:03:10] iteration 100/ 100
| consumed samples: 800 | elapsed time per iteration (ms): 2904.8 | learning rate: 3.045865E-09 | global batch size: 8 | loss: 1.423446
E+00 | grad norm: 9.433 |
/root/mmmm_26/lib/python3.10/site-packages/torch/distributed/checkpoint/filesystem.py:592: UserWarning: Detected an existing checkpoint in save_path/iter_0000100/.metadata, overwriting since self.overwrite=True. Past version 2.5 of PyTorch, `overwrite` will default to False. Set this variable to True to maintain this functionality or False to raise when an existing checkpoint is found.
  warnings.warn(
/root/mmmm_26/lib/python3.10/site-packages/torch_npu/distributed/distributed_c10d.py:110: UserWarning: HCCL doesn't support gather at the moment. Implemented with allgather instead.
  warnings.warn("HCCL doesn't support gather at the moment. Implemented with allgather instead.")
[Rank 0 | Local Rank 0] 2026-09-03 08:03:15,533 INFO [mindspeed_mm.fsd.checkpoint.dcp_checkpoint:48] => Saved checkpoint to ./save_path/iter_0000100
```

20260724_官方基线日志_triton精度日志

文件: 20260724_官方基线日志_triton精度日志.log | 行数: 713 | 大小: 95.5 KB | ★ 官方 7/24 参考基线日志 (精度/性能对比基准)
采集说明: 完整日志 (从程序启动开始, 未经裁剪)

```
W0724 16:01:34.540000 880073 site-packages/torch/distributed/run.py:766]
W0724 16:01:34.540000 880073 site-packages/torch/distributed/run.py:766] *****
W0724 16:01:34.540000 880073 site-packages/torch/distributed/run.py:766] Setting OMP_NUM_THREADS environment variable for each process to be 1 in de
fault, to avoid your system being overloaded, please further tune the variable for optimal performance in your application as needed.
W0724 16:01:34.540000 880073 site-packages/torch/distributed/run.py:766] *****
[INFO] Starting configuration loading...
[INFO] Loaded configuration file: examples/qwen3_5/qwen3_5_4B_config.yaml
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdps.data.data_utils.func_utils.mod
el_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.model_name_or_path (type: <class 'str'>, value: /home/model/Qwe
n3.5-0.8B-hf)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdps.data.data_utils.func_utils.mod
el_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.use_fast_tokenizer (type: <class 'bool'>, value: True)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdps.data.data_utils.func_utils.mod
el_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.split_special_tokens (type: <class 'bool'>, value: False)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdps.data.data_utils.func_utils.mod
el_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.image_max_pixels (type: <class 'int'>, value: 262144)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdps.data.data_utils.func_utils.mod
el_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.image_min_pixels (type: <class 'int'>, value: 1024)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdps.data.data_utils.func_utils.mod
el_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.video_max_pixels (type: <class 'int'>, value: 16384)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdps.data.data_utils.func_utils.mod
el_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.video_min_pixels (type: <class 'int'>, value: 0)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdps.data.data_utils.func_utils.mod
el_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.video_fps (type: <class 'float'>, value: 2.0)
[WARNING] Field 'preprocess_parameters' exists but is not a BaseArguments subclass: typing.Optional[mindspeed_mm.fsdps.data.data_utils.func_utils.mod
el_args.ProcessorArguments]
[INFO] Automatically registered YAML field: data.dataset_param.preprocess_parameters.video_maxlen (type: <class 'int'>, value: 64)
[INFO] Automatically registered YAML field: model.gdn_implementation (type: <class 'str'>, value: triton)
[INFO] Automatically registered YAML field: model.causal_conv1d_implementation (type: <class 'str'>, value: triton)
[INFO] Automatically registered YAML field: model.skip_gdn_recompute (type: <class 'bool'>, value: True)
[INFO] Automatically registered YAML field: model.skip_flash_attn_recompute (type: <class 'bool'>, value: True)
[INFO] Automatically registered YAML field: features.recompute (type: <class 'bool'>, value: True)
=====
BaseArguments Configuration Manager Summary
=====

Configuration class: Arguments
Configuration file: examples/qwen3_5/qwen3_5_4B_config.yaml
Allow YAML extensions: True
Allow CLI override: True
Strict CLI validation: True
Allow registering YAML fields: True

Defined field count: 203
Dynamic field count: 14
Total field count: 217

===== Configuration Details =====
parallel:
  data_parallel_size: 2
  fully_shard_parallel_size: 2
  fsdp:
    ignored_modules: []
    apply_modules: ['model.visual', 'model.visual.blocks.*', 'model.language_model.embed_tokens', 'model.language_model', 'model.language_model.la
yers.*'], 'lm_head', 'mtp']
    param_dtype: bf16
    reduce_dtype: fp32
    output_dtype: None
    cast_forward_inputs: True
    reshard_after_forward: True
    num_to_forward_prefetch: 1
    num_to_backward_prefetch: 1
    pregather: False
    hook_modules: None
    cpu_offload: False
  tensor_parallel_size: 1
  tp_plan:
    colwise_parallel: []
    rowwise_parallel: []
    sequence_parallel: []
  ring_attention_size: 1
  ulysses_parallel_size: 1
  expert_parallel_size: 1
  expert_fully_shard_parallel_size: 2
  ep_plan:
    apply_modules: []
    use_npu_fused_ops: True
```

```

    dispatcher: alltoall
    apply_efsd_p_modules: []
recompute: False
recompute_plan:
    apply_modules: []
    use_reentrant: False
local_rank: 0
global_rank: 0
world_size: 2
model:
    model_id: qwen3_5
    model_name_or_path: /home/model/Qwen3.5-0.8B-hf
    trust_remote_code: True
    attn_implementation: flash_attention_2
    freeze: ['model.visual']
    mtp_num_layers: 0
    mtp_loss_scaling_factor: 0.1
    gdn_implementation: triton
    causal_conv1d_implementation: triton
    skip_gdn_recompute: True
    skip_flash_attn_recompute: True
data:
    dataset_param:
        dataset_type: huggingface
        basic_parameters:
            cache_dir: ./cache_dir/
            template: qwen3_vl_nothink
            chat_template: None
            enable_thinking: False
            dataset_dir: /home/dataset/llava_coco/coco/
            dataset: ['data/s00977099/output_llava_coco_data.json']
            cutoff_len: 1024
            train_on_prompt: False
            mask_history: False
            streaming: False
            overwrite_cache: False
            preprocessing_batch_size: 1000
            preprocessing_num_workers: 16
            max_samples: None
            tool_format: None
            val_dataset: None
            val_max_samples: None
            val_rate: None
            packing: None
            neat_packing: False
            preprocess_on_fly: True
            async_preprocess: False
            async_preprocess_buffer_size: None
            stage: sft
        preprocess_parameters:
            model_name_or_path: /home/model/Qwen3.5-0.8B-hf
            use_fast_tokenizer: True
            split_special_tokens: False
            image_resolution: 512
            video_resolution: 128
            image_max_pixels: 262144
            image_min_pixels: 1024
            video_max_pixels: 16384
            video_min_pixels: 0
            video_fps: 2.0
            video_maxlen: 64
            image_do_pan_and_scan: False
            crop_to_patches: False
            use_audio_in_video: False
            audio_sampling_rate: 16000
            trust_remote_code: False
    attr:
        ranking: False
        system: None
        images: images
        videos: None
        audios: None
        prompt: None
        query: None
        response: None
        history: None
        messages: messages
        tools: None
        role_tag: role
        content_tag: content
        user_tag: user
        assistant_tag: assistant
        observation_tag: observation
        function_tag: function_call
        system_tag: system
        chosen: None
        rejected: None
        formatting: sharegpt
val_dataset_param: None
dataloading_param:
    dataloader_mode: sampler
    sampler_type: BaseRandomBatchSampler
    shuffle: False

```

```

drop_last: True
pin_memory: True
collate_param:
  collator_id: None
  model_name: qwen3v1
  ignore_pad_token_for_loss: True
  pad_to_multiple_of: 8
num_workers: 8
enable_preload: False
training:
  profile:
    enable: False
    start_step: 1
    end_step: 2
    save_path: ./profiling
    record_shapes: True
    with_memory: True
    with_stack: True
    ranks: [0]
    _train_steps: -1
    local_rank: 0
    global_rank: 0
    world_size: 2
    profile_this_rank: False
lora:
  enable: False
  rank: 8
  alpha: 16
  target_modules: ['q_proj', 'k_proj', 'v_proj']
  dropout: 0.0
  init_lora_weights: True
  pretrained_lora_path: None
  lora_target_modules_support: None
lr: 1e-05
lr_min: 0.0
lr_start: 0.0
weight_decay: 0.0
no_decay_modules: []
no_decay_params: []
optimizer: adamw
matched_adamw_rms: 0.2
muon_momentum: 0.95
ns_steps: 5
adam_fused: True
adam_beta1: 0.9
adam_beta2: 0.999
adam_eps: 1e-08
clip_grad: 0.0
clip_grad_foreach: True
micro_batch_size: 4
gradient_accumulation_steps: 1
lr_warmup_ratio: 0.1
lr_decay_style: cosine
lr_decay_ratio: 1.0
enable_full_determinism: False
ckpt_manager: dcp
save_async: False
load_checkpoint_path: None
save_steps: 0
seed: 42
max_steps: None
init_model_with_meta_device: True
train_iters: 100
train_epochs: None
load: /home/model/Qwen3.5-0.8B-dcp
load_format: auto
load_strict: False
load_rank0_and_broadcast: False
no_load_optim: True
no_load_rng: True
save: ./save_path
no_save_optim: True
no_save_rng: True
log_interval: 1
save_interval: 100
val_interval: 0
val_micro_batch_size: None
use_deter_comp: False
allow_hf32: None
plugin: ['mindspeed_mm/fsdp/models/qwen3_5', 'mindspeed_mm/fsdp/data/datasets/huggingface']
lr_scaling_plan: None
quantization_plan: QuantizeConfig(recipe_name=None, apply_modules=[], ignored_modules=[], quant_converters=[], extra_args={}, enable_fsdp_low_precision_all_gather=True, fsdp_low_precision_all_gather_mode='on-demand')
_train_steps: -1
local_rank: 0
global_rank: 0
world_size: 2
global_batch_size: 8
tools:
  profile:
    enable: False
    profile_type: static
    ranks: [0]

```

```

static_param:
  level: level1
  with_stack: False
  with_memory: False
  record_shapes: False
  with_cpu: True
  save_path: ./profiling
  start_step: 10
  end_step: 11
  data_simplification: False
  aic_metrics_type: PipeUtilization
  analyse_flag: True
memory_profile:
  enable: False
  start_step: 1
  end_step: 2
  save_path: ./memory_snapshot
  dump_ranks: [0]
  stacks: all
  max_entries: None
  mem_info: False
features:
  recompute_plan:
    apply_modules: ['model.visual.blocks.*', 'model.language_model.layers.*']
    use_reentrant: False
  loss_cfg:
    loss_type: default
    router_aux_loss_coef: 0.0
    router_aux_loss_offload: False
  enable_chunk_loss: True
  enable_dynamic_chunk_loss: False
  chunkloss_plan:
    apply_module: lm_head
    chunk_size: 1024
    total_chunk_size: 4096
  enable_activation_offload: True
  activation_offload_plan:
    apply_modules: ['model.visual.blocks.*', 'model.language_model.layers.*']
  enable_chunk_mbs: False
  chunkmbs_plan:
    apply_modules: ['model.language_model.layers.*']
    chunk_mbs: 2
    batch_dim: 0
    chunk_arg_indexes: [0]
    chunk_kwarg_names: ['position_embeddings', 'position_ids', 'rope_deltas', 'attention_mask']
  enable_ep_balance: False
  ep_balance_plan:
    max_dup_experts_num: 2
  recompute: True
data.dataset_param.preprocess_parameters.model_name_or_path: /home/model/Qwen3.5-0.8B-hf
data.dataset_param.preprocess_parameters.use_fast_tokenizer: True
data.dataset_param.preprocess_parameters.split_special_tokens: False
data.dataset_param.preprocess_parameters.image_max_pixels: 262144
data.dataset_param.preprocess_parameters.image_min_pixels: 1024
data.dataset_param.preprocess_parameters.video_max_pixels: 16384
data.dataset_param.preprocess_parameters.video_min_pixels: 0
data.dataset_param.preprocess_parameters.video_fps: 2.0
data.dataset_param.preprocess_parameters.video_maxlen: 64

```

===== Dynamic Fields =====

```

data.dataset_param.preprocess_parameters.image_max_pixels
data.dataset_param.preprocess_parameters.image_min_pixels
data.dataset_param.preprocess_parameters.model_name_or_path
data.dataset_param.preprocess_parameters.split_special_tokens
data.dataset_param.preprocess_parameters.use_fast_tokenizer
data.dataset_param.preprocess_parameters.video_fps
data.dataset_param.preprocess_parameters.video_max_pixels
data.dataset_param.preprocess_parameters.video_maxlen
data.dataset_param.preprocess_parameters.video_min_pixels
features.recompute
model.causal_conv1d_implementation
model.gdn_implementation
model.skip_flash_attn_recompute
model.skip_gdn_recompute

```

=====

Successfully imported: mindspeed_mm.fsdpy.models.qwen3_5

Successfully imported: mindspeed_mm.fsdpy.models.qwen3_5

[Rank 1 | Local Rank 1] 2026-07-24 16:01:50,595 INFO [torch_npy.env:15] => get env TRITON_CACHE_DIR = ./triton_cache

[Rank 0 | Local Rank 0] 2026-07-24 16:01:50,597 INFO [torch_npy.env:15] => get env TRITON_CACHE_DIR = ./triton_cache

Successfully imported: mindspeed_mm.fsdpy.data.datasets.huggingface

Successfully imported: mindspeed_mm.fsdpy.data.datasets.huggingface

[Rank 1 | Local Rank 1] 2026-07-24 16:01:51,293 INFO [torch_npy.env:15] => get env USE_TF = FALSE

[Rank 1 | Local Rank 1] 2026-07-24 16:01:51,293 INFO [datasets:114] => Disabling Tensorflow because USE_TORCH is set

[Rank 0 | Local Rank 0] 2026-07-24 16:01:51,307 INFO [torch_npy.env:15] => get env USE_TF = FALSE

[Rank 0 | Local Rank 0] 2026-07-24 16:01:51,307 INFO [datasets:114] => Disabling Tensorflow because USE_TORCH is set

[Rank 1 | Local Rank 1] 2026-07-24 16:01:51,427 INFO [torch_npy.env:15] => get env MASTER_ADDR = localhost

[Rank 1 | Local Rank 1] 2026-07-24 16:01:51,427 INFO [torch_npy.env:15] => get env MASTER_PORT = 6000

[Rank 1 | Local Rank 1] 2026-07-24 16:01:51,427 INFO [torch_npy.env:15] => get env TORCHELASTIC_USE_AGENT_STORE = True

[Rank 0 | Local Rank 0] 2026-07-24 16:01:51,435 INFO [torch_npy.env:15] => get env MASTER_ADDR = localhost

[Rank 0 | Local Rank 0] 2026-07-24 16:01:51,435 INFO [torch_npy.env:15] => get env MASTER_PORT = 6000

[Rank 0 | Local Rank 0] 2026-07-24 16:01:51,435 INFO [torch_npy.env:15] => get env TORCHELASTIC_USE_AGENT_STORE = True

[Rank 0 | Local Rank 0] 2026-07-24 16:01:51,441 INFO [mindspeed_mm.fsdpy.distributed.parallel_state:87] => Parallel state initialized:

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

```
[Rank 0 | Local Rank 0] 2026-07-24 16:01:59,728 WARNING [mindspeed_mm.fsdps.data.data_loader.data_collator:177] => model_name is deprecated and will be removed in a future version. Please use collator_id instead. Available collator_ids: llm_pretrain, qwen3asr, qwen3omni, qwen3vl, step3_vl.
[Rank 1 | Local Rank 1] 2026-07-24 16:02:00,664 INFO [httpx:1025] => HTTP Request: HEAD https://s3.amazonaws.com/datasets.huggingface.co/datasets/datasets/json/json.py "HTTP/1.1 200 OK"
[Rank 1 | Local Rank 1] 2026-07-24 16:02:00,867 WARNING [mindspeed_mm.fsdps.data.data_loader.data_collator:177] => model_name is deprecated and will be removed in a future version. Please use collator_id instead. Available collator_ids: llm_pretrain, qwen3asr, qwen3omni, qwen3vl, step3_vl.
2026-07-24 16:02:02,313 - template - 730 - INFO - Replace eos token: <|im_end|>
[Rank 0 | Local Rank 0] 2026-07-24 16:02:02,313 INFO [mindspeed_mm.fsdps.data.data_utils.func_utils.template:730] => Replace eos token: <|im_end|>
2026-07-24 16:02:03,497 - template - 730 - INFO - Replace eos token: <|im_end|>
[Rank 1 | Local Rank 1] 2026-07-24 16:02:03,497 INFO [mindspeed_mm.fsdps.data.data_utils.func_utils.template:730] => Replace eos token: <|im_end|>
/root/miniconda3/envs/szm_mindspeedmm/lib/python3.11/site-packages/torch_npu/distributed/distributed_c10d.py:158: UserWarning: HCCL doesn't support gather at the moment. Implemented with allgather instead.
  warnings.warn("HCCL doesn't support gather at the moment. Implemented with allgather instead.")
[rank1]:[W724 16:02:03.479879561 TensorFactories.cpp:340] Warning: Cannot create tensor with internal format while allow_internal_format=False, tensor will be created with base format. (function operator())
[Rank 0 | Local Rank 0] 2026-07-24 16:02:04,314 INFO [mindspeed_mm.fsdps.checkpoint.dcp_checkpoint:47] => Loaded checkpoint from /home/model/Qwen3.5-0.8B-dcp/release
[2026-07-24 16:02:04] [WARNING] [880220] profiler.py: Profiler won't be using warmup, this can skew profiler results
/data/s00977099/MindSpeed-MM/mindspeed_mm/fsdps/features/memory/async_offload.py:25: UserWarning: TypedStorage is deprecated. It will be removed in the future and UntypedStorage will be the only storage class. This should only matter to you if you are using storages directly. To access UntypedStorage directly, use tensor.untyped_storage() instead of tensor.storage()
  if tensor.storage().size() <= 0:
/data/s00977099/MindSpeed-MM/mindspeed_mm/fsdps/features/memory/async_offload.py:25: UserWarning: TypedStorage is deprecated. It will be removed in the future and UntypedStorage will be the only storage class. This should only matter to you if you are using storages directly. To access UntypedStorage directly, use tensor.untyped_storage() instead of tensor.storage()
  if tensor.storage().size() <= 0:
[rank0]:[W724 16:02:05.331867563 TensorFactories.cpp:340] Warning: Cannot create tensor with internal format while allow_internal_format=False, tensor will be created with base format. (function operator())
[Rank 0 | Local Rank 0] 2026-07-24 16:04:03,724 INFO [mindspeed_mm.fsdps.train.train_engine:47] => [2026-07-24 16:04:03] iteration 1/ 100
| consumed samples: 8 | elapsed time per iteration (ms): 119405.2 | learning rate: 0.000000E+00 | global batch size: 8 | loss: 1.9246E+00 | grad norm: 134.266 |
[Rank 0 | Local Rank 0] 2026-07-24 16:04:06,228 INFO [mindspeed_mm.fsdps.train.train_engine:47] => [2026-07-24 16:04:06] iteration 2/ 100
| consumed samples: 16 | elapsed time per iteration (ms): 2503.1 | learning rate: 1.000000E-06 | global batch size: 8 | loss: 1.791238E+00 | grad norm: 112.753 |
[Rank 0] (after 2 iterations) memory (MB) | allocated: 5970.84716796875 | max allocated: 16178.6923828125 | reserved: 18680.0 | max reserved: 18680.0
[Rank 0 | Local Rank 0] 2026-07-24 16:04:15,712 INFO [mindspeed_mm.fsdps.train.train_engine:47] => [2026-07-24 16:04:15] iteration 3/ 100
| consumed samples: 24 | elapsed time per iteration (ms): 9483.1 | learning rate: 2.000000E-06 | global batch size: 8 | loss: 1.784901E+00 | grad norm: 88.392 |
[Rank 0 | Local Rank 0] 2026-07-24 16:04:25,515 INFO [mindspeed_mm.fsdps.train.train_engine:47] => [2026-07-24 16:04:25] iteration 4/ 100
| consumed samples: 32 | elapsed time per iteration (ms): 9801.8 | learning rate: 3.000000E-06 | global batch size: 8 | loss: 1.506463E+00 | grad norm: 58.012 |
[Rank 0 | Local Rank 0] 2026-07-24 16:04:27,549 INFO [mindspeed_mm.fsdps.train.train_engine:47] => [2026-07-24 16:04:27] iteration 5/ 100
| consumed samples: 40 | elapsed time per iteration (ms): 2032.8 | learning rate: 4.000000E-06 | global batch size: 8 | loss: 1.408003E+00 | grad norm: 23.029 |
[Rank 0 | Local Rank 0] 2026-07-24 16:04:36,900 INFO [mindspeed_mm.fsdps.train.train_engine:47] => [2026-07-24 16:04:36] iteration 6/ 100
| consumed samples: 48 | elapsed time per iteration (ms): 9350.9 | learning rate: 5.000000E-06 | global batch size: 8 | loss: 1.499018E+00 | grad norm: 15.763 |
[Rank 0 | Local Rank 0] 2026-07-24 16:04:38,928 INFO [mindspeed_mm.fsdps.train.train_engine:47] => [2026-07-24 16:04:38] iteration 7/ 100
| consumed samples: 56 | elapsed time per iteration (ms): 2026.6 | learning rate: 6.000000E-06 | global batch size: 8 | loss: 1.395148E+00 | grad norm: 15.688 |
[Rank 0 | Local Rank 0] 2026-07-24 16:04:48,297 INFO [mindspeed_mm.fsdps.train.train_engine:47] => [2026-07-24 16:04:48] iteration 8/ 100
| consumed samples: 64 | elapsed time per iteration (ms): 9368.2 | learning rate: 7.000000E-06 | global batch size: 8 | loss: 1.465237E+00 | grad norm: 12.628 |
[Rank 0 | Local Rank 0] 2026-07-24 16:04:50,375 INFO [mindspeed_mm.fsdps.train.train_engine:47] => [2026-07-24 16:04:50] iteration 9/ 100
| consumed samples: 72 | elapsed time per iteration (ms): 2077.3 | learning rate: 8.000000E-06 | global batch size: 8 | loss: 1.268659E+00 | grad norm: 12.821 |
[Rank 0 | Local Rank 0] 2026-07-24 16:04:59,869 INFO [mindspeed_mm.fsdps.train.train_engine:47] => [2026-07-24 16:04:59] iteration 10/ 100
| consumed samples: 80 | elapsed time per iteration (ms): 9492.5 | learning rate: 9.000000E-06 | global batch size: 8 | loss: 1.405683E+00 | grad norm: 10.895 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:01,883 INFO [mindspeed_mm.fsdps.train.train_engine:47] => [2026-07-24 16:05:01] iteration 11/ 100
| consumed samples: 88 | elapsed time per iteration (ms): 2012.7 | learning rate: 1.000000E-05 | global batch size: 8 | loss: 1.432606E+00 | grad norm: 12.031 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:11,451 INFO [mindspeed_mm.fsdps.train.train_engine:47] => [2026-07-24 16:05:11] iteration 12/ 100
| consumed samples: 96 | elapsed time per iteration (ms): 9567.7 | learning rate: 9.96954E-06 | global batch size: 8 | loss: 1.401814E+00 | grad norm: 11.461 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:11,891 INFO [mindspeed_mm.fsdps.train.train_engine:47] => [2026-07-24 16:05:11] iteration 13/ 100
| consumed samples: 104 | elapsed time per iteration (ms): 439.1 | learning rate: 9.987820E-06 | global batch size: 8 | loss: 1.465265E+00 | grad norm: 10.048 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:12,344 INFO [mindspeed_mm.fsdps.train.train_engine:47] => [2026-07-24 16:05:12] iteration 14/ 100
| consumed samples: 112 | elapsed time per iteration (ms): 451.7 | learning rate: 9.972609E-06 | global batch size: 8 | loss: 1.471768E+00 | grad norm: 11.265 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:21,851 INFO [mindspeed_mm.fsdps.train.train_engine:47] => [2026-07-24 16:05:21] iteration 15/ 100
| consumed samples: 120 | elapsed time per iteration (ms): 9504.5 | learning rate: 9.951340E-06 | global batch size: 8 | loss: 1.418221E+00 | grad norm: 9.823 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:22,261 INFO [mindspeed_mm.fsdps.train.train_engine:47] => [2026-07-24 16:05:22] iteration 16/ 100
| consumed samples: 128 | elapsed time per iteration (ms): 409.9 | learning rate: 9.924039E-06 | global batch size: 8 | loss: 1.432037E+00 | grad norm: 11.301 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:22,681 INFO [mindspeed_mm.fsdps.train.train_engine:47] => [2026-07-24 16:05:22] iteration 17/ 100
| consumed samples: 136 | elapsed time per iteration (ms): 419.4 | learning rate: 9.890738E-06 | global batch size: 8 | loss: 1.411621E+00 | grad norm: 10.323 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:23,109 INFO [mindspeed_mm.fsdps.train.train_engine:47] => [2026-07-24 16:05:23] iteration 18/ 100
| consumed samples: 144 | elapsed time per iteration (ms): 426.8 | learning rate: 9.851479E-06 | global batch size: 8 | loss: 1.380019E+00 | grad norm: 10.323 |
```

+00 | grad norm: 11.871 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:23,533 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:05:23] iteration 19/ 100
| consumed samples: 152 | elapsed time per iteration (ms): 423.2 | learning rate: 9.806308E-06 | global batch size: 8 | loss: 1.462066E
+00 | grad norm: 10.028 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:23,961 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:05:23] iteration 20/ 100
| consumed samples: 160 | elapsed time per iteration (ms): 427.5 | learning rate: 9.755283E-06 | global batch size: 8 | loss: 1.350745E
+00 | grad norm: 10.383 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:24,361 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:05:24] iteration 21/ 100
| consumed samples: 168 | elapsed time per iteration (ms): 398.9 | learning rate: 9.698463E-06 | global batch size: 8 | loss: 1.255295E
+00 | grad norm: 10.741 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:24,790 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:05:24] iteration 22/ 100
| consumed samples: 176 | elapsed time per iteration (ms): 428.9 | learning rate: 9.635919E-06 | global batch size: 8 | loss: 1.431265E
+00 | grad norm: 9.815 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:25,190 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:05:25] iteration 23/ 100
| consumed samples: 184 | elapsed time per iteration (ms): 399.5 | learning rate: 9.567727E-06 | global batch size: 8 | loss: 1.479332E
+00 | grad norm: 11.547 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:25,616 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:05:25] iteration 24/ 100
| consumed samples: 192 | elapsed time per iteration (ms): 425.2 | learning rate: 9.493970E-06 | global batch size: 8 | loss: 1.481910E
+00 | grad norm: 9.780 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:26,032 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:05:26] iteration 25/ 100
| consumed samples: 200 | elapsed time per iteration (ms): 414.9 | learning rate: 9.414738E-06 | global batch size: 8 | loss: 1.421660E
+00 | grad norm: 11.045 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:26,421 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:05:26] iteration 26/ 100
| consumed samples: 208 | elapsed time per iteration (ms): 388.9 | learning rate: 9.330127E-06 | global batch size: 8 | loss: 1.313326E
+00 | grad norm: 11.947 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:26,911 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:05:26] iteration 27/ 100
| consumed samples: 216 | elapsed time per iteration (ms): 489.2 | learning rate: 9.240240E-06 | global batch size: 8 | loss: 1.478197E
+00 | grad norm: 10.455 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:27,338 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:05:27] iteration 28/ 100
| consumed samples: 224 | elapsed time per iteration (ms): 425.7 | learning rate: 9.145188E-06 | global batch size: 8 | loss: 1.352563E
+00 | grad norm: 10.383 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:27,779 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:05:27] iteration 29/ 100
| consumed samples: 232 | elapsed time per iteration (ms): 440.4 | learning rate: 9.045085E-06 | global batch size: 8 | loss: 1.452959E
+00 | grad norm: 10.197 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:28,193 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:05:28] iteration 30/ 100
| consumed samples: 240 | elapsed time per iteration (ms): 413.4 | learning rate: 8.940054E-06 | global batch size: 8 | loss: 1.285910E
+00 | grad norm: 9.376 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:30,251 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:05:30] iteration 31/ 100
| consumed samples: 248 | elapsed time per iteration (ms): 2057.2 | learning rate: 8.830222E-06 | global batch size: 8 | loss: 1.320338E
+00 | grad norm: 9.492 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:30,681 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:05:30] iteration 32/ 100
| consumed samples: 256 | elapsed time per iteration (ms): 429.3 | learning rate: 8.715724E-06 | global batch size: 8 | loss: 1.436788E
+00 | grad norm: 9.369 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:31,119 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:05:31] iteration 33/ 100
| consumed samples: 264 | elapsed time per iteration (ms): 436.9 | learning rate: 8.596699E-06 | global batch size: 8 | loss: 1.277287E
+00 | grad norm: 9.742 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:31,548 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:05:31] iteration 34/ 100
| consumed samples: 272 | elapsed time per iteration (ms): 428.1 | learning rate: 8.473292E-06 | global batch size: 8 | loss: 1.293409E
+00 | grad norm: 10.337 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:33,605 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:05:33] iteration 35/ 100
| consumed samples: 280 | elapsed time per iteration (ms): 2056.4 | learning rate: 8.345653E-06 | global batch size: 8 | loss: 1.410657E
+00 | grad norm: 9.264 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:34,305 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:05:34] iteration 36/ 100
| consumed samples: 288 | elapsed time per iteration (ms): 699.3 | learning rate: 8.213938E-06 | global batch size: 8 | loss: 1.407040E
+00 | grad norm: 9.936 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:34,740 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:05:34] iteration 37/ 100
| consumed samples: 296 | elapsed time per iteration (ms): 434.7 | learning rate: 8.078307E-06 | global batch size: 8 | loss: 1.293624E
+00 | grad norm: 11.824 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:35,176 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:05:35] iteration 38/ 100
| consumed samples: 304 | elapsed time per iteration (ms): 435.3 | learning rate: 7.938926E-06 | global batch size: 8 | loss: 1.471243E
+00 | grad norm: 10.320 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:35,612 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:05:35] iteration 39/ 100
| consumed samples: 312 | elapsed time per iteration (ms): 434.3 | learning rate: 7.795965E-06 | global batch size: 8 | loss: 1.354094E
+00 | grad norm: 9.545 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:36,045 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:05:36] iteration 40/ 100
| consumed samples: 320 | elapsed time per iteration (ms): 433.1 | learning rate: 7.649596E-06 | global batch size: 8 | loss: 1.515606E
+00 | grad norm: 8.956 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:36,482 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:05:36] iteration 41/ 100
| consumed samples: 328 | elapsed time per iteration (ms): 435.3 | learning rate: 7.500000E-06 | global batch size: 8 | loss: 1.403632E
+00 | grad norm: 10.493 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:36,916 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:05:36] iteration 42/ 100
| consumed samples: 336 | elapsed time per iteration (ms): 432.8 | learning rate: 7.347358E-06 | global batch size: 8 | loss: 1.288146E
+00 | grad norm: 9.205 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:37,316 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:05:37] iteration 43/ 100
| consumed samples: 344 | elapsed time per iteration (ms): 399.9 | learning rate: 7.191856E-06 | global batch size: 8 | loss: 1.322229E
+00 | grad norm: 10.296 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:37,756 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:05:37] iteration 44/ 100
| consumed samples: 352 | elapsed time per iteration (ms): 439.2 | learning rate: 7.033683E-06 | global batch size: 8 | loss: 1.357560E
+00 | grad norm: 11.093 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:38,179 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:05:38] iteration 45/ 100
| consumed samples: 360 | elapsed time per iteration (ms): 422.4 | learning rate: 6.873033E-06 | global batch size: 8 | loss: 1.290732E
+00 | grad norm: 9.466 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:38,583 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:05:38] iteration 46/ 100
| consumed samples: 368 | elapsed time per iteration (ms): 403.1 | learning rate: 6.710101E-06 | global batch size: 8 | loss: 1.331221E
+00 | grad norm: 10.476 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:40,596 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:05:40] iteration 47/ 100
| consumed samples: 376 | elapsed time per iteration (ms): 2011.8 | learning rate: 6.545085E-06 | global batch size: 8 | loss: 1.391710E
+00 | grad norm: 10.833 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:41,020 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:05:41] iteration 48/ 100
| consumed samples: 384 | elapsed time per iteration (ms): 424.0 | learning rate: 6.378187E-06 | global batch size: 8 | loss: 1.470920E
+00 | grad norm: 9.603 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:41,447 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:05:41] iteration 49/ 100
| consumed samples: 392 | elapsed time per iteration (ms): 425.8 | learning rate: 6.209609E-06 | global batch size: 8 | loss: 1.273075E

+00 | grad norm: 10.072 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:50,987 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:05:50] iteration 50/ 100
| consumed samples: 400 | elapsed time per iteration (ms): 9539.1 | learning rate: 6.039558E-06 | global batch size: 8 | loss: 1.285009
+00 | grad norm: 9.127 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:51,417 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:05:51] iteration 51/ 100
| consumed samples: 408 | elapsed time per iteration (ms): 429.0 | learning rate: 5.868241E-06 | global batch size: 8 | loss: 1.361099E
+00 | grad norm: 9.245 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:51,855 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:05:51] iteration 52/ 100
| consumed samples: 416 | elapsed time per iteration (ms): 437.5 | learning rate: 5.695866E-06 | global batch size: 8 | loss: 1.557325E
+00 | grad norm: 9.374 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:52,278 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:05:52] iteration 53/ 100
| consumed samples: 424 | elapsed time per iteration (ms): 421.9 | learning rate: 5.522642E-06 | global batch size: 8 | loss: 1.304100E
+00 | grad norm: 10.079 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:52,710 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:05:52] iteration 54/ 100
| consumed samples: 432 | elapsed time per iteration (ms): 431.3 | learning rate: 5.348782E-06 | global batch size: 8 | loss: 1.278860E
+00 | grad norm: 9.576 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:53,122 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:05:53] iteration 55/ 100
| consumed samples: 440 | elapsed time per iteration (ms): 411.7 | learning rate: 5.174497E-06 | global batch size: 8 | loss: 1.300684E
+00 | grad norm: 10.596 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:53,566 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:05:53] iteration 56/ 100
| consumed samples: 448 | elapsed time per iteration (ms): 443.5 | learning rate: 5.000000E-06 | global batch size: 8 | loss: 1.402378E
+00 | grad norm: 9.496 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:54,021 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:05:54] iteration 57/ 100
| consumed samples: 456 | elapsed time per iteration (ms): 454.2 | learning rate: 4.825503E-06 | global batch size: 8 | loss: 1.392132E
+00 | grad norm: 9.649 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:54,448 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:05:54] iteration 58/ 100
| consumed samples: 464 | elapsed time per iteration (ms): 426.6 | learning rate: 4.651218E-06 | global batch size: 8 | loss: 1.377736E
+00 | grad norm: 8.852 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:54,849 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:05:54] iteration 59/ 100
| consumed samples: 472 | elapsed time per iteration (ms): 400.2 | learning rate: 4.477358E-06 | global batch size: 8 | loss: 1.213646E
+00 | grad norm: 10.248 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:55,241 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:05:55] iteration 60/ 100
| consumed samples: 480 | elapsed time per iteration (ms): 390.9 | learning rate: 4.304134E-06 | global batch size: 8 | loss: 1.269621E
+00 | grad norm: 10.533 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:55,651 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:05:55] iteration 61/ 100
| consumed samples: 488 | elapsed time per iteration (ms): 409.1 | learning rate: 4.131759E-06 | global batch size: 8 | loss: 1.300824E
+00 | grad norm: 9.549 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:56,084 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:05:56] iteration 62/ 100
| consumed samples: 496 | elapsed time per iteration (ms): 432.1 | learning rate: 3.960442E-06 | global batch size: 8 | loss: 1.175227E
+00 | grad norm: 11.454 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:56,498 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:05:56] iteration 63/ 100
| consumed samples: 504 | elapsed time per iteration (ms): 413.4 | learning rate: 3.790391E-06 | global batch size: 8 | loss: 1.406934E
+00 | grad norm: 9.713 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:56,931 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:05:56] iteration 64/ 100
| consumed samples: 512 | elapsed time per iteration (ms): 432.4 | learning rate: 3.621813E-06 | global batch size: 8 | loss: 1.334252E
+00 | grad norm: 10.411 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:57,336 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:05:57] iteration 65/ 100
| consumed samples: 520 | elapsed time per iteration (ms): 404.2 | learning rate: 3.454915E-06 | global batch size: 8 | loss: 1.297628E
+00 | grad norm: 10.460 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:57,774 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:05:57] iteration 66/ 100
| consumed samples: 528 | elapsed time per iteration (ms): 436.4 | learning rate: 3.289899E-06 | global batch size: 8 | loss: 1.415693E
+00 | grad norm: 9.580 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:58,207 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:05:58] iteration 67/ 100
| consumed samples: 536 | elapsed time per iteration (ms): 432.7 | learning rate: 3.126967E-06 | global batch size: 8 | loss: 1.293933E
+00 | grad norm: 9.831 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:58,641 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:05:58] iteration 68/ 100
| consumed samples: 544 | elapsed time per iteration (ms): 433.5 | learning rate: 2.966317E-06 | global batch size: 8 | loss: 1.509670E
+00 | grad norm: 10.334 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:59,087 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:05:59] iteration 69/ 100
| consumed samples: 552 | elapsed time per iteration (ms): 445.0 | learning rate: 2.808144E-06 | global batch size: 8 | loss: 1.426578E
+00 | grad norm: 9.381 |
[Rank 0 | Local Rank 0] 2026-07-24 16:05:59,498 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:05:59] iteration 70/ 100
| consumed samples: 560 | elapsed time per iteration (ms): 409.8 | learning rate: 2.652642E-06 | global batch size: 8 | loss: 1.334287E
+00 | grad norm: 10.359 |
[Rank 0 | Local Rank 0] 2026-07-24 16:06:00,189 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:06:00] iteration 71/ 100
| consumed samples: 568 | elapsed time per iteration (ms): 691.2 | learning rate: 2.500000E-06 | global batch size: 8 | loss: 1.276605E
+00 | grad norm: 10.749 |
[Rank 0 | Local Rank 0] 2026-07-24 16:06:00,868 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:06:00] iteration 72/ 100
| consumed samples: 576 | elapsed time per iteration (ms): 677.4 | learning rate: 2.350404E-06 | global batch size: 8 | loss: 1.401531E
+00 | grad norm: 9.748 |
[Rank 0 | Local Rank 0] 2026-07-24 16:06:01,309 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:06:01] iteration 73/ 100
| consumed samples: 584 | elapsed time per iteration (ms): 440.9 | learning rate: 2.204035E-06 | global batch size: 8 | loss: 1.380450E
+00 | grad norm: 9.611 |
[Rank 0 | Local Rank 0] 2026-07-24 16:06:01,751 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:06:01] iteration 74/ 100
| consumed samples: 592 | elapsed time per iteration (ms): 441.2 | learning rate: 2.061074E-06 | global batch size: 8 | loss: 1.328064E
+00 | grad norm: 10.235 |
[Rank 0 | Local Rank 0] 2026-07-24 16:06:02,140 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:06:02] iteration 75/ 100
| consumed samples: 600 | elapsed time per iteration (ms): 388.3 | learning rate: 1.921693E-06 | global batch size: 8 | loss: 1.367047E
+00 | grad norm: 9.576 |
[Rank 0 | Local Rank 0] 2026-07-24 16:06:02,569 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:06:02] iteration 76/ 100
| consumed samples: 608 | elapsed time per iteration (ms): 427.8 | learning rate: 1.786062E-06 | global batch size: 8 | loss: 1.337646E
+00 | grad norm: 9.834 |
[Rank 0 | Local Rank 0] 2026-07-24 16:06:02,992 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:06:02] iteration 77/ 100
| consumed samples: 616 | elapsed time per iteration (ms): 422.0 | learning rate: 1.654347E-06 | global batch size: 8 | loss: 1.401107E
+00 | grad norm: 9.559 |
[Rank 0 | Local Rank 0] 2026-07-24 16:06:03,403 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:06:03] iteration 78/ 100
| consumed samples: 624 | elapsed time per iteration (ms): 410.3 | learning rate: 1.526708E-06 | global batch size: 8 | loss: 1.418421E
+00 | grad norm: 9.876 |
[Rank 0 | Local Rank 0] 2026-07-24 16:06:03,843 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:06:03] iteration 79/ 100
| consumed samples: 632 | elapsed time per iteration (ms): 439.2 | learning rate: 1.403301E-06 | global batch size: 8 | loss: 1.449646E
+00 | grad norm: 9.676 |
[Rank 0 | Local Rank 0] 2026-07-24 16:06:04,281 INFO [mindspeed_mm.fsdsp.train.train_engine:47] => [2026-07-24 16:06:04] iteration 80/ 100
| consumed samples: 640 | elapsed time per iteration (ms): 437.9 | learning rate: 1.284276E-06 | global batch size: 8 | loss: 1.276647E

```
+00 | grad norm: 8.813 |
[Rank 0 | Local Rank 0] 2026-07-24 16:06:04,686 INFO [mindspeed_mm.fsdp.train.train_engine:47] => [2026-07-24 16:06:04] iteration      81/    100
| consumed samples:      648 | elapsed time per iteration (ms): 404.0 | learning rate: 1.169778E-06 | global batch size:      8 | loss: 1.295212E
+00 | grad norm: 10.193 |
[Rank 0 | Local Rank 0] 2026-07-24 16:06:05,133 INFO [mindspeed_mm.fsdp.train.train_engine:47] => [2026-07-24 16:06:05] iteration      82/    100
| consumed samples:      656 | elapsed time per iteration (ms): 446.2 | learning rate: 1.059946E-06 | global batch size:      8 | loss: 1.275704E
+00 | grad norm: 9.978 |
[Rank 0 | Local Rank 0] 2026-07-24 16:06:05,545 INFO [mindspeed_mm.fsdp.train.train_engine:47] => [2026-07-24 16:06:05] iteration      83/    100
| consumed samples:      664 | elapsed time per iteration (ms): 411.0 | learning rate: 9.549150E-07 | global batch size:      8 | loss: 1.263605E
+00 | grad norm: 9.442 |
[Rank 0 | Local Rank 0] 2026-07-24 16:06:05,974 INFO [mindspeed_mm.fsdp.train.train_engine:47] => [2026-07-24 16:06:05] iteration      84/    100
| consumed samples:      672 | elapsed time per iteration (ms): 428.7 | learning rate: 8.548121E-07 | global batch size:      8 | loss: 1.234609E
+00 | grad norm: 9.194 |
[Rank 0 | Local Rank 0] 2026-07-24 16:06:06,499 INFO [mindspeed_mm.fsdp.train.train_engine:47] => [2026-07-24 16:06:06] iteration      85/    100
| consumed samples:      680 | elapsed time per iteration (ms): 523.7 | learning rate: 7.597595E-07 | global batch size:      8 | loss: 1.318295E
+00 | grad norm: 11.694 |
[Rank 0 | Local Rank 0] 2026-07-24 16:06:06,909 INFO [mindspeed_mm.fsdp.train.train_engine:47] => [2026-07-24 16:06:06] iteration      86/    100
| consumed samples:      688 | elapsed time per iteration (ms): 409.5 | learning rate: 6.698730E-07 | global batch size:      8 | loss: 1.360037E
+00 | grad norm: 10.681 |
[Rank 0 | Local Rank 0] 2026-07-24 16:06:07,339 INFO [mindspeed_mm.fsdp.train.train_engine:47] => [2026-07-24 16:06:07] iteration      87/    100
| consumed samples:      696 | elapsed time per iteration (ms): 429.7 | learning rate: 5.852620E-07 | global batch size:      8 | loss: 1.367115E
+00 | grad norm: 8.786 |
[Rank 0 | Local Rank 0] 2026-07-24 16:06:07,779 INFO [mindspeed_mm.fsdp.train.train_engine:47] => [2026-07-24 16:06:07] iteration      88/    100
| consumed samples:      704 | elapsed time per iteration (ms): 438.7 | learning rate: 5.060298E-07 | global batch size:      8 | loss: 1.280388E
+00 | grad norm: 10.103 |
[Rank 0 | Local Rank 0] 2026-07-24 16:06:08,210 INFO [mindspeed_mm.fsdp.train.train_engine:47] => [2026-07-24 16:06:08] iteration      89/    100
| consumed samples:      712 | elapsed time per iteration (ms): 430.1 | learning rate: 4.322727E-07 | global batch size:      8 | loss: 1.110990E
+00 | grad norm: 9.705 |
[Rank 0 | Local Rank 0] 2026-07-24 16:06:08,645 INFO [mindspeed_mm.fsdp.train.train_engine:47] => [2026-07-24 16:06:08] iteration      90/    100
| consumed samples:      720 | elapsed time per iteration (ms): 434.7 | learning rate: 3.640807E-07 | global batch size:      8 | loss: 1.300898E
+00 | grad norm: 9.587 |
[Rank 0 | Local Rank 0] 2026-07-24 16:06:09,111 INFO [mindspeed_mm.fsdp.train.train_engine:47] => [2026-07-24 16:06:09] iteration      91/    100
| consumed samples:      728 | elapsed time per iteration (ms): 465.4 | learning rate: 3.015369E-07 | global batch size:      8 | loss: 1.352559E
+00 | grad norm: 8.153 |
[Rank 0 | Local Rank 0] 2026-07-24 16:06:09,501 INFO [mindspeed_mm.fsdp.train.train_engine:47] => [2026-07-24 16:06:09] iteration      92/    100
| consumed samples:      736 | elapsed time per iteration (ms): 388.9 | learning rate: 2.447174E-07 | global batch size:      8 | loss: 1.290553E
+00 | grad norm: 10.371 |
[Rank 0 | Local Rank 0] 2026-07-24 16:06:09,928 INFO [mindspeed_mm.fsdp.train.train_engine:47] => [2026-07-24 16:06:09] iteration      93/    100
| consumed samples:      744 | elapsed time per iteration (ms): 426.1 | learning rate: 1.936915E-07 | global batch size:      8 | loss: 1.200977E
+00 | grad norm: 10.067 |
[Rank 0 | Local Rank 0] 2026-07-24 16:06:10,375 INFO [mindspeed_mm.fsdp.train.train_engine:47] => [2026-07-24 16:06:10] iteration      94/    100
| consumed samples:      752 | elapsed time per iteration (ms): 446.7 | learning rate: 1.485214E-07 | global batch size:      8 | loss: 1.441367E
+00 | grad norm: 9.688 |
[Rank 0 | Local Rank 0] 2026-07-24 16:06:10,804 INFO [mindspeed_mm.fsdp.train.train_engine:47] => [2026-07-24 16:06:10] iteration      95/    100
| consumed samples:      760 | elapsed time per iteration (ms): 428.3 | learning rate: 1.092620E-07 | global batch size:      8 | loss: 1.313985E
+00 | grad norm: 9.777 |
[Rank 0 | Local Rank 0] 2026-07-24 16:06:12,909 INFO [mindspeed_mm.fsdp.train.train_engine:47] => [2026-07-24 16:06:12] iteration      96/    100
| consumed samples:      768 | elapsed time per iteration (ms): 2103.7 | learning rate: 7.596123E-08 | global batch size:      8 | loss: 1.322307E
+00 | grad norm: 8.616 |
[Rank 0 | Local Rank 0] 2026-07-24 16:06:13,341 INFO [mindspeed_mm.fsdp.train.train_engine:47] => [2026-07-24 16:06:13] iteration      97/    100
| consumed samples:      776 | elapsed time per iteration (ms): 431.3 | learning rate: 4.865966E-08 | global batch size:      8 | loss: 1.329676E
+00 | grad norm: 9.990 |
[Rank 0 | Local Rank 0] 2026-07-24 16:06:13,727 INFO [mindspeed_mm.fsdp.train.train_engine:47] => [2026-07-24 16:06:13] iteration      98/    100
| consumed samples:      784 | elapsed time per iteration (ms): 384.7 | learning rate: 2.739052E-08 | global batch size:      8 | loss: 1.307330E
+00 | grad norm: 11.279 |
[Rank 0 | Local Rank 0] 2026-07-24 16:06:14,164 INFO [mindspeed_mm.fsdp.train.train_engine:47] => [2026-07-24 16:06:14] iteration      99/    100
| consumed samples:      792 | elapsed time per iteration (ms): 437.0 | learning rate: 1.217975E-08 | global batch size:      8 | loss: 1.365586E
+00 | grad norm: 10.154 |
[Rank 0 | Local Rank 0] 2026-07-24 16:06:14,600 INFO [mindspeed_mm.fsdp.train.train_engine:47] => [2026-07-24 16:06:14] iteration     100/    100
| consumed samples:      800 | elapsed time per iteration (ms): 434.8 | learning rate: 3.045865E-09 | global batch size:      8 | loss: 1.448866E
+00 | grad norm: 10.077 |
/root/miniconda3/envs/szm_mindspeedmm/lib/python3.11/site-packages/torch/distributed/checkpoint/filesystem.py:592: UserWarning: Detected an existing
checkpoint in save_path/iter_0000100/.metadata, overwriting since self.overwrite=True. Past version 2.5 of PyTorch, `overwrite` will default to Fals
e. Set this variable to True to maintain this functionality or False to raise when an existing checkpoint is found.
  warnings.warn(
/root/miniconda3/envs/szm_mindspeedmm/lib/python3.11/site-packages/torch/distributed/checkpoint/filesystem.py:592: UserWarning: Detected an existing
checkpoint in save_path/iter_0000100/.metadata, overwriting since self.overwrite=True. Past version 2.5 of PyTorch, `overwrite` will default to Fals
e. Set this variable to True to maintain this functionality or False to raise when an existing checkpoint is found.
  warnings.warn(
/root/miniconda3/envs/szm_mindspeedmm/lib/python3.11/site-packages/torch_npu/distributed/distributed_c10d.py:158: UserWarning: HCCL doesn't support
gather at the moment. Implemented with allgather instead.
  warnings.warn("HCCL doesn't support gather at the moment. Implemented with allgather instead.")
[Rank 0 | Local Rank 0] 2026-07-24 16:06:16,921 INFO [mindspeed_mm.fsdp.checkpoint.dcp_checkpointinter:47] => Saved checkpoint to ./save_path/iter_0000
100
[Rank 0 | Local Rank 0] 2026-07-24 16:06:18,497 INFO [mindspeed_mm.fsdp.checkpoint.dcp_checkpointinter:47] => Saved checkpoint to ./save_path/iter_0000
100
Elapsed Time Per iteration (ms): 2502.0
Average Samples per Second: 3.197
```